

From Zero to Generative



Learning Generative Modelling from scratch

Carolina Cuesta-Lazaro

*Flatiron Institute
Institute for Advanced Studies*

Art: "The art of painting" by Johannes Vermeer

Probabilistic ML has made *high dimensional inference tractable*

$$p(\text{World}|\text{Prompt})$$

["Genie 2: A large-scale foundation model"
Parker-Holder et al (2024)]

Probabilistic ML has made *high dimensional inference tractable*

$$p(\text{World} | \text{Prompt})$$

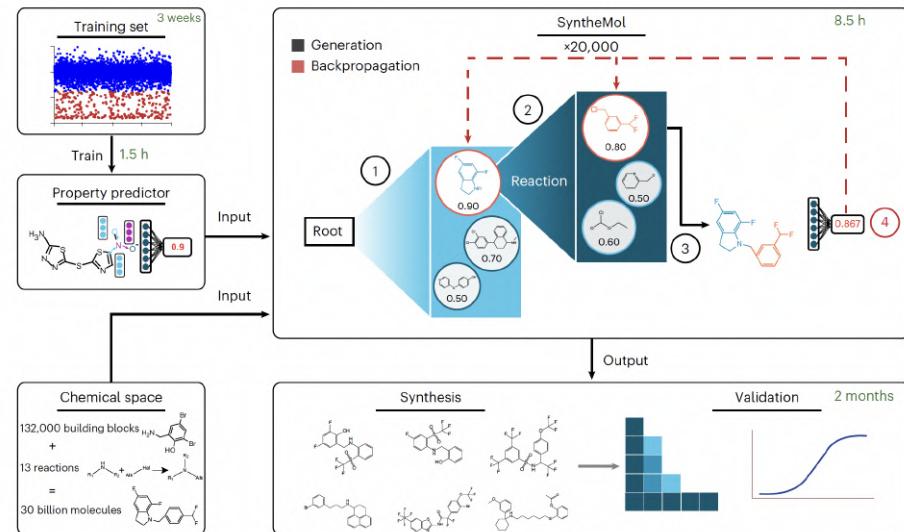
["Genie 2: A large-scale foundation model"
Parker-Holder et al (2024)]

["Genie 3: A new frontier for world models"
Parker-Holder et al (2025)]

Probabilistic ML has made *high dimensional inference tractable*

$p(\text{World}|\text{Prompt})$

$p(\text{Drug}|\text{Properties})$



["Genie 2: A large-scale foundation model"
Parker-Holder et al (2024)]

["Genie 3: A new frontier for world models"
Parker-Holder et al (2025)]

["Generative AI for designing and validating easily synthesizable and structurally novel antibiotics"
Swanson et al]

Parti-350M



A portrait photo of a kangaroo wearing an orange hoodie and blue sunglasses standing on the grass in front of the Sydney Opera House holding a sign on the chest that says Welcome Friends!

<https://parti.research.google>

Parti-350M



Parti-750M



A portrait photo of a kangaroo wearing an orange hoodie and blue sunglasses standing on the grass in front of the Sydney Opera House holding a sign on the chest that says Welcome Friends!

<https://parti.research.google>

Parti-350M



Parti-750M



Parti-3B



A portrait photo of a kangaroo wearing an orange hoodie and blue sunglasses standing on the grass in front of the Sydney Opera House holding a sign on the chest that says Welcome Friends!

<https://parti.research.google>

Parti-350M



Parti-750M



Parti-3B



Parti-20B

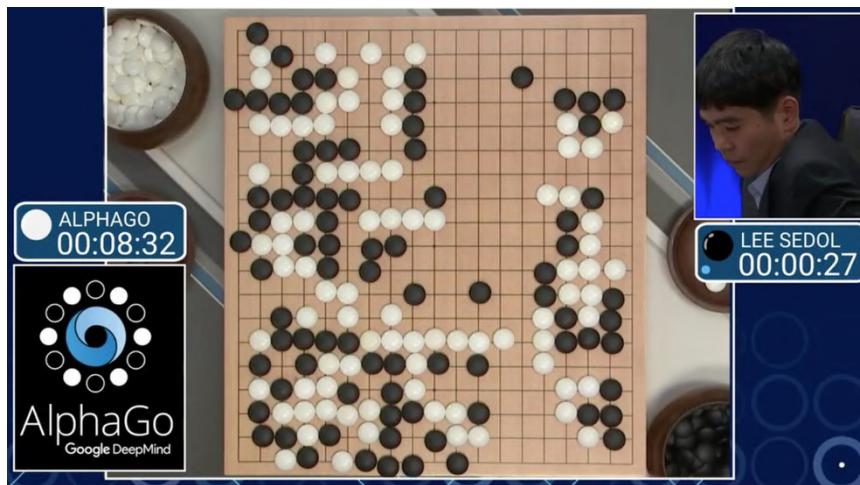


A portrait photo of a kangaroo wearing an orange hoodie and blue sunglasses standing on the grass in front of the Sydney Opera House holding a sign on the chest that says Welcome Friends!

<https://parti.research.google>

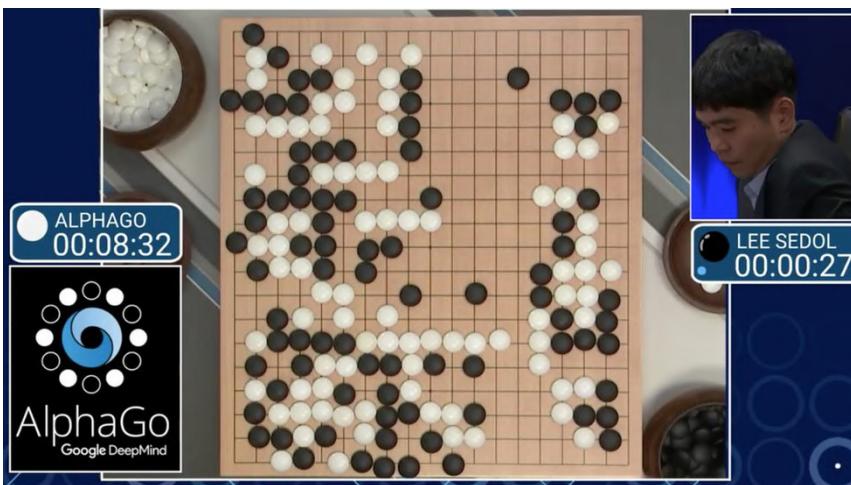
Artificial General Intelligence?

BEFORE



Artificial General Intelligence?

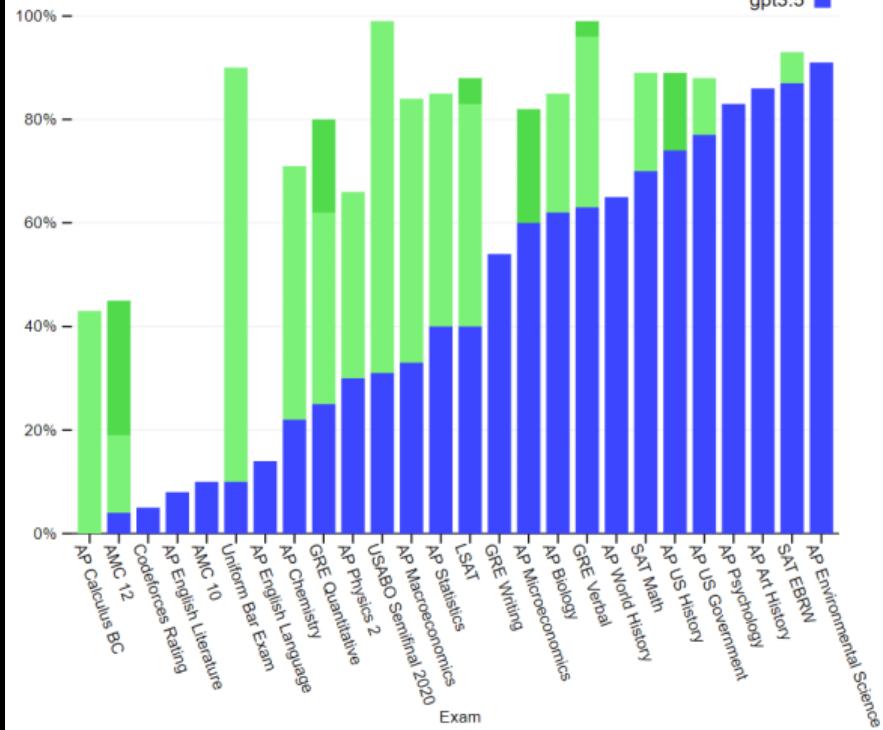
BEFORE



AFTER

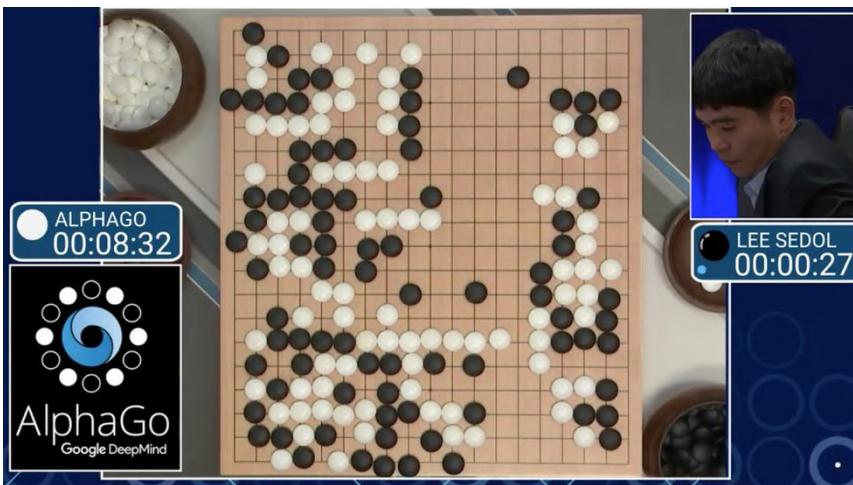
Exam results (ordered by GPT-3.5 performance)

Estimated percentile lower bound (among test takers)

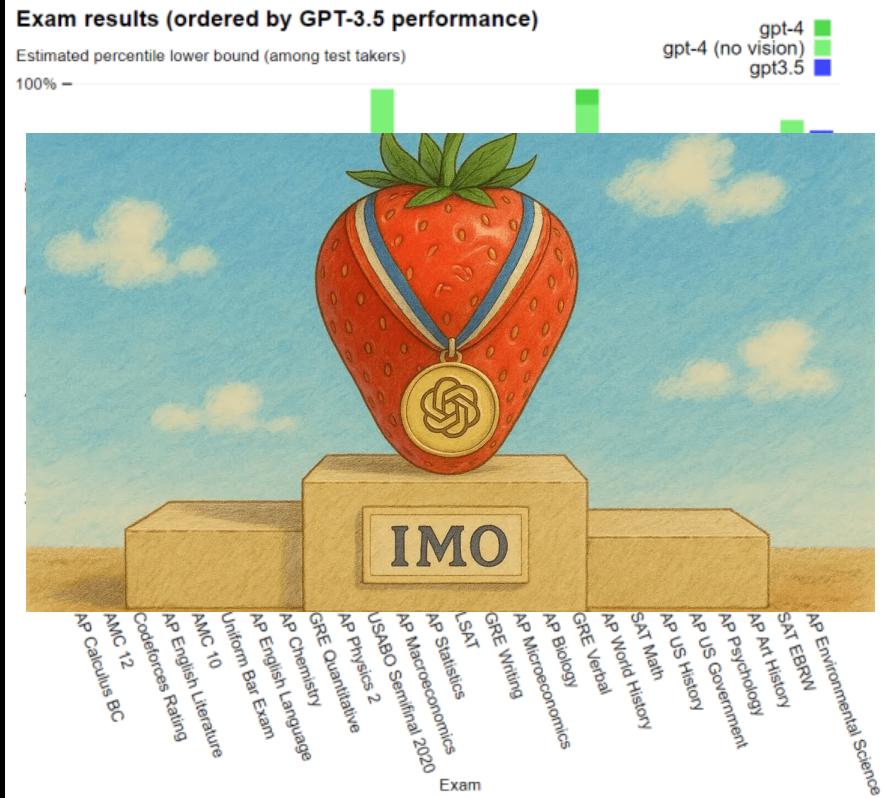


Artificial General Intelligence?

BEFORE

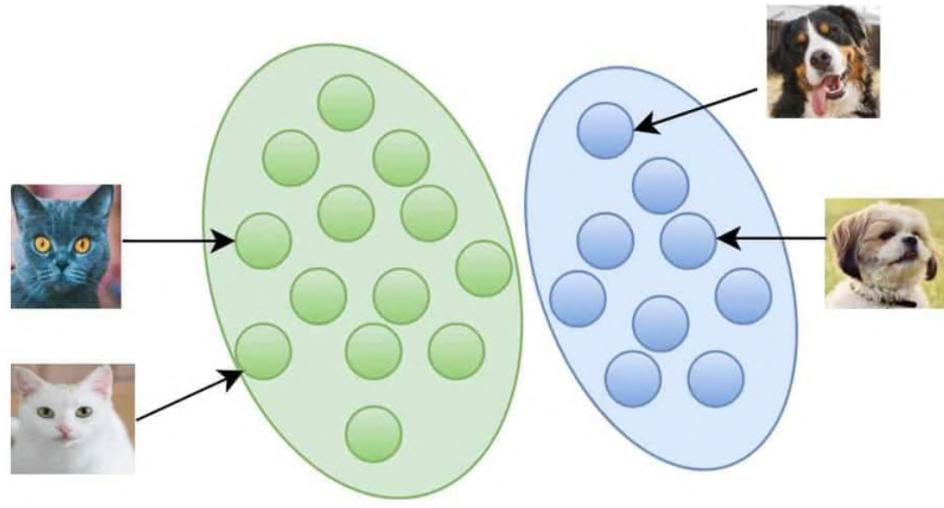


AFTER



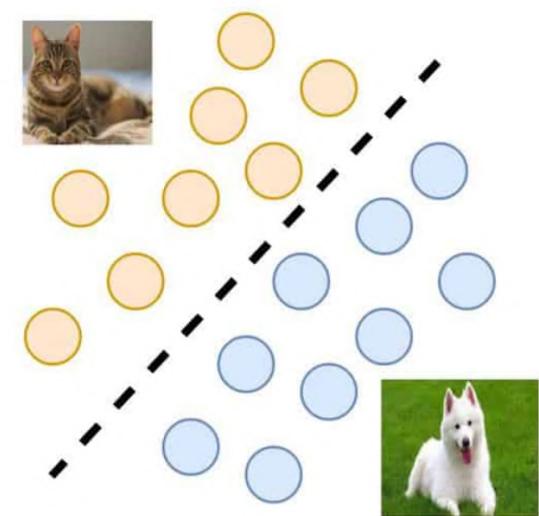
Generation vs Discrimination

$$p(x)$$



Generative

$$p(y|x)$$

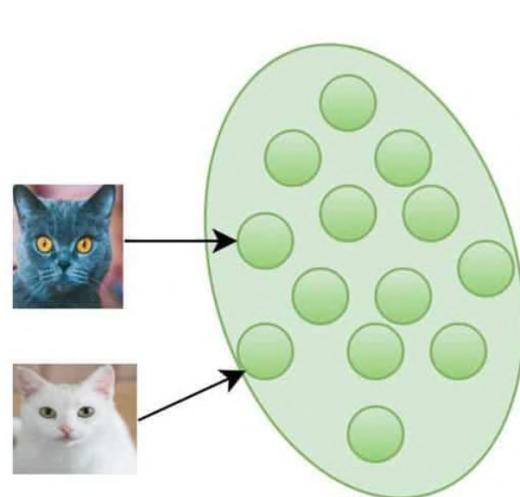


Discriminative

$$p(x|y)$$

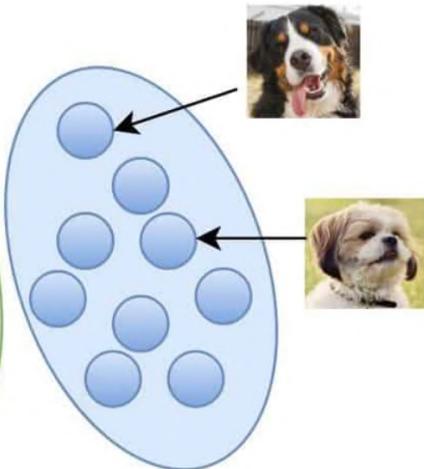
Generation vs Discrimination

$$p(x)$$



Generative

$$p(y|x)$$



Discriminative

$$p(x|y)$$

$$p(x|y) = \frac{p(y|x)p(x)}{p(y)}$$

Generative Models

$$p_{\phi}(x)$$


Generative Models

Data

$$p_{\phi}(\boxed{x})$$



Generative Models



Data

$$p_{\phi}(x)$$

*A PDF that we can
optimize*

Generative Models

*Maximize the likelihood of
the data*



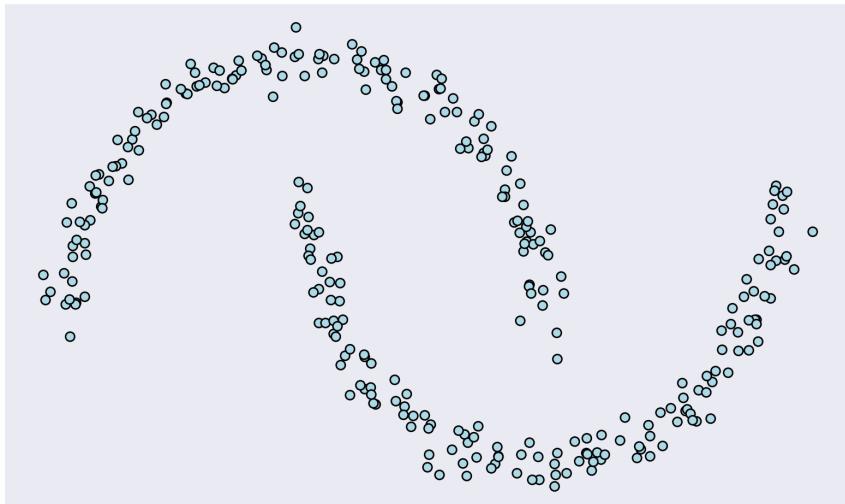
Data

$$p_{\phi}(x)$$

*A PDF that we can
optimize*

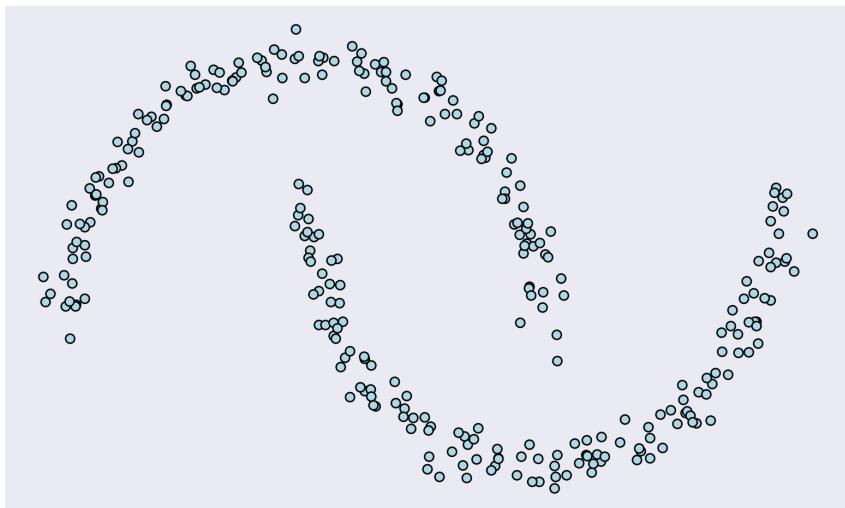
Generative Models 101

Training Samples x_{train}

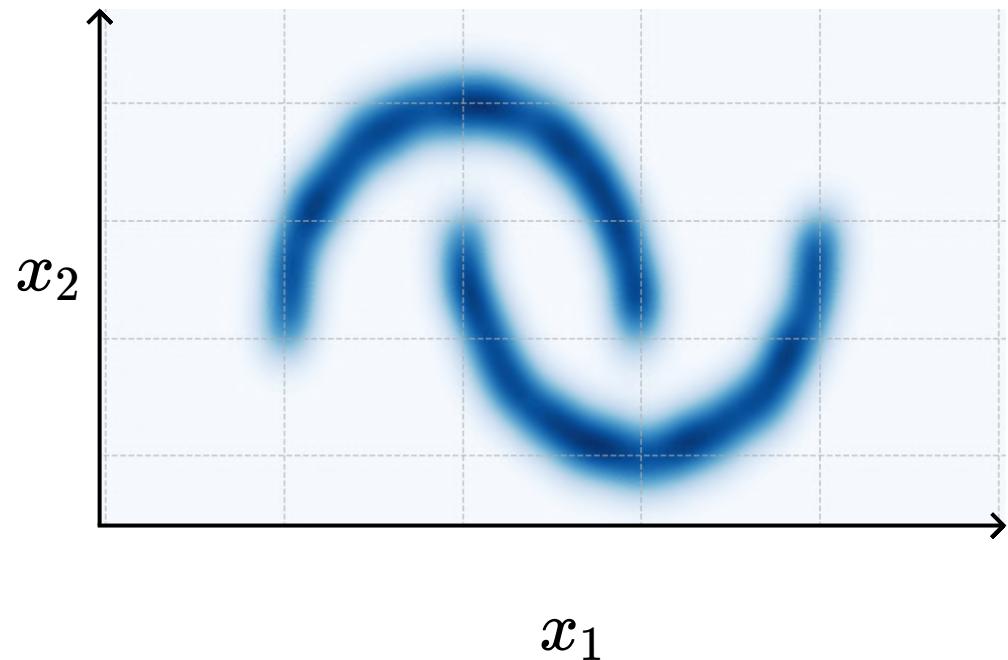


Generative Models 101

Training Samples x_{train}

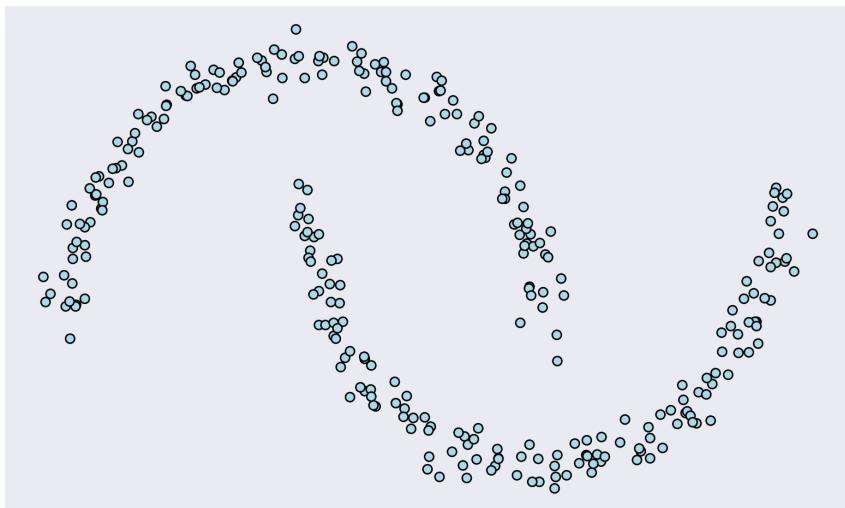


Parametric Model $p_{\phi}(x)$

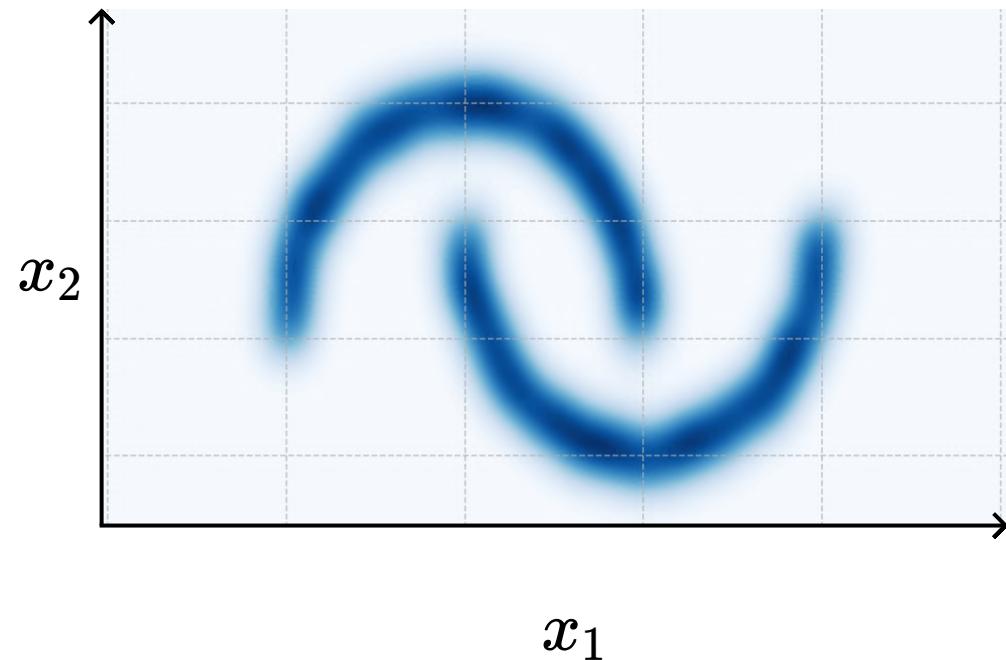


Generative Models 101

Training Samples x_{train}



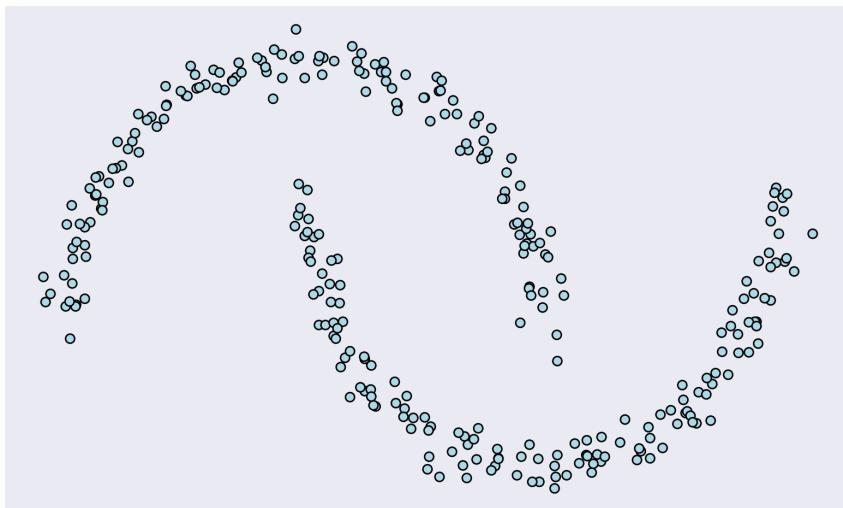
Parametric Model $p_{\phi}(x)$



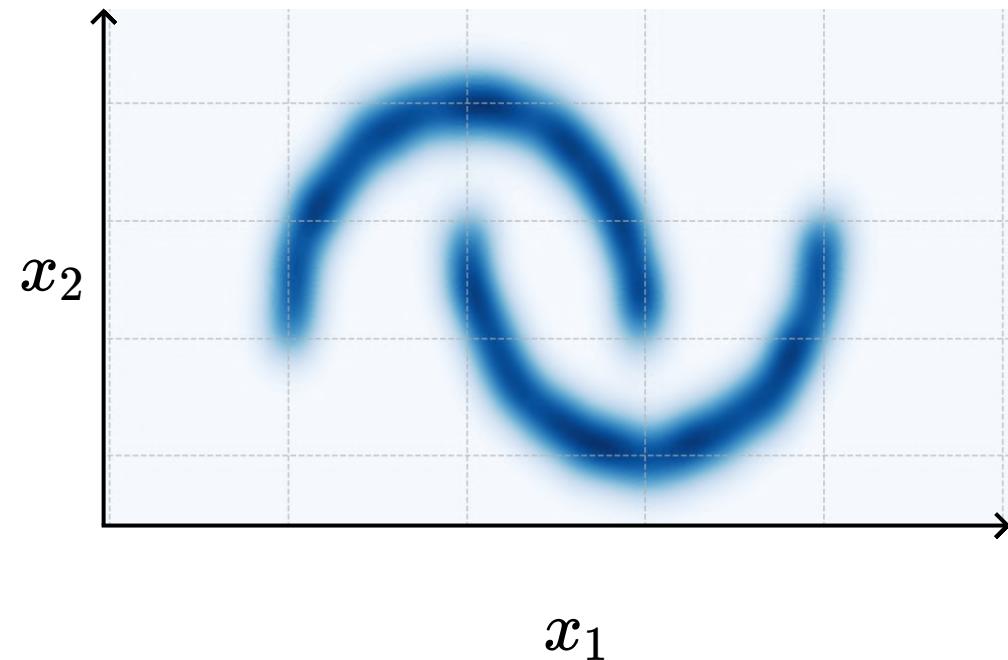
Maximize the likelihood of the training samples

Generative Models 101

Training Samples x_{train}



Parametric Model $p_{\phi}(x)$

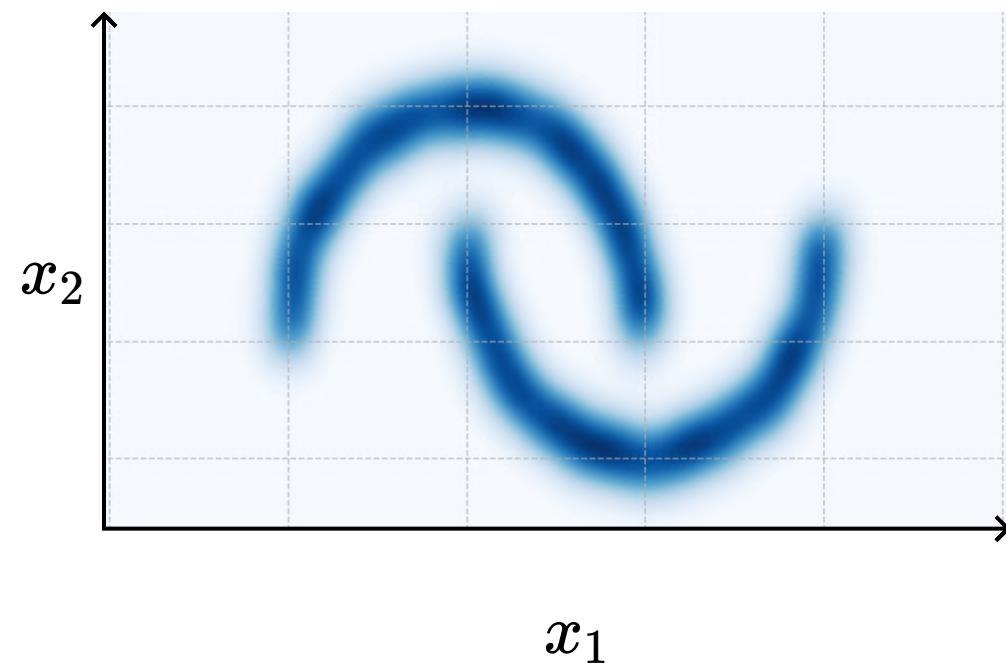


Maximize the likelihood of the training samples

$$\hat{\phi} = \arg \max [\log p_{\phi}(x_{\text{train}})]$$

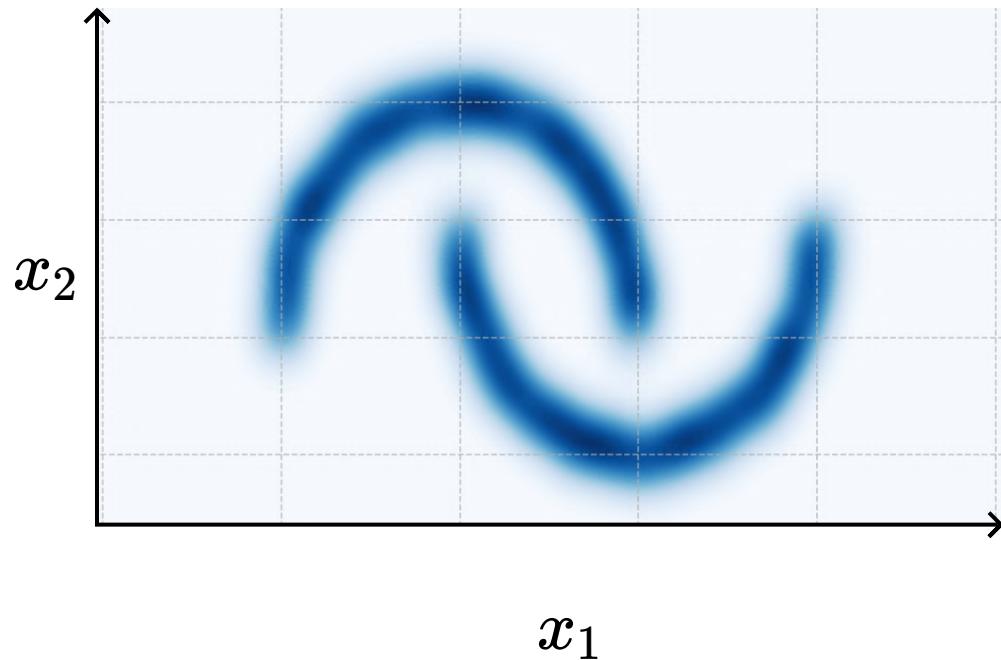
Generative Models: Simulate and Analyze

Trained Model $p_\phi(x)$



Generative Models: Simulate and Analyze

Trained Model $p_\phi(x)$

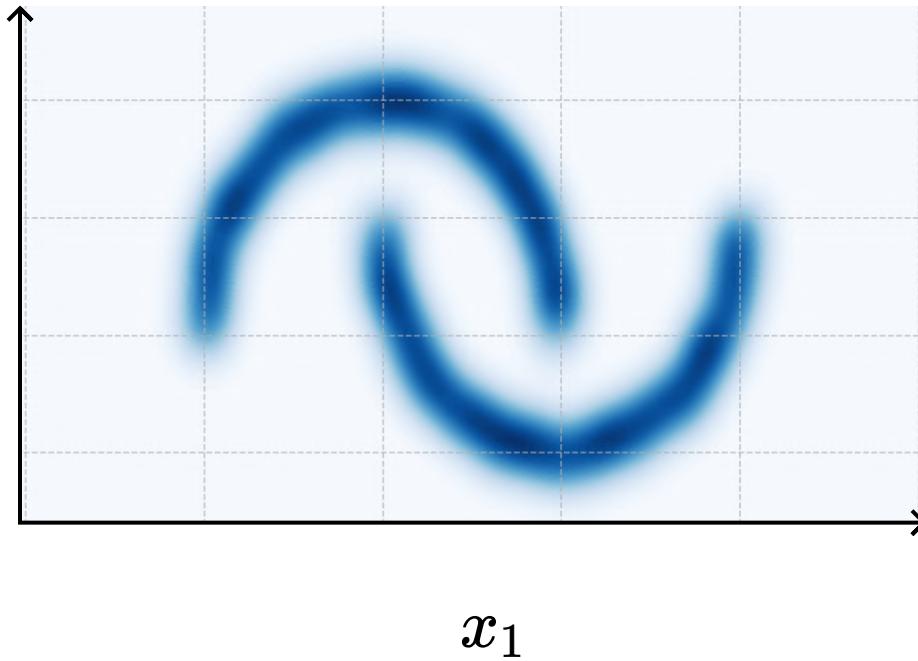


Generate Novel Samples



Generative Models: Simulate and Analyze

Trained Model $p_\phi(x)$



Generate Novel Samples



Evaluate probabilities



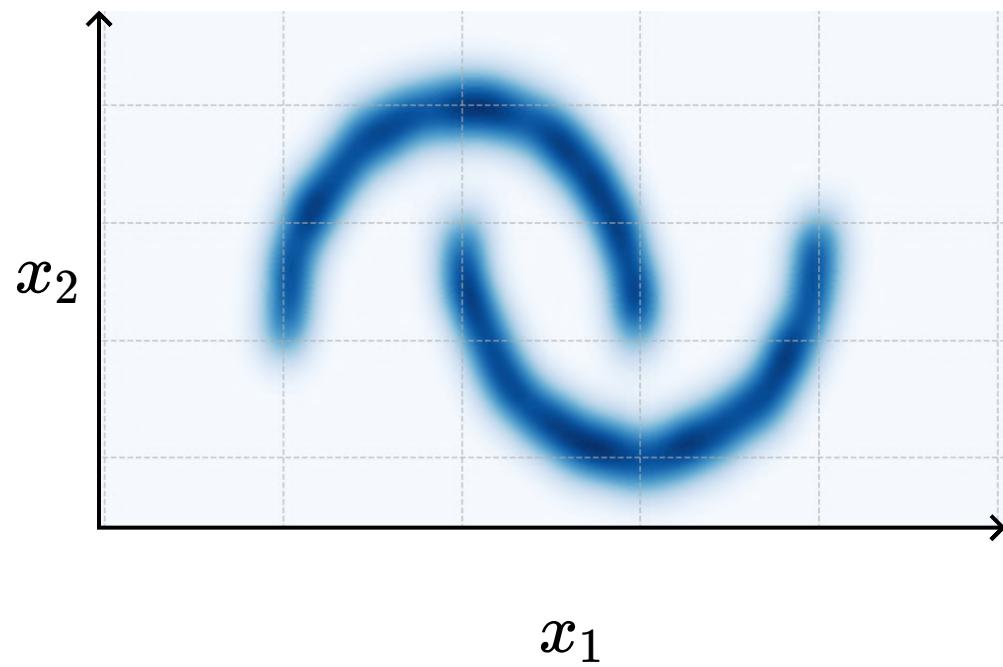
High Probability



Low Probability

Generative Models: Simulate and Analyze

Trained Model $p_\phi(x)$



Generate Novel Samples



✓ Generative Model

✓ Simulator

Evaluate probabilities



High Probability



Low Probability

✓ Generative Model

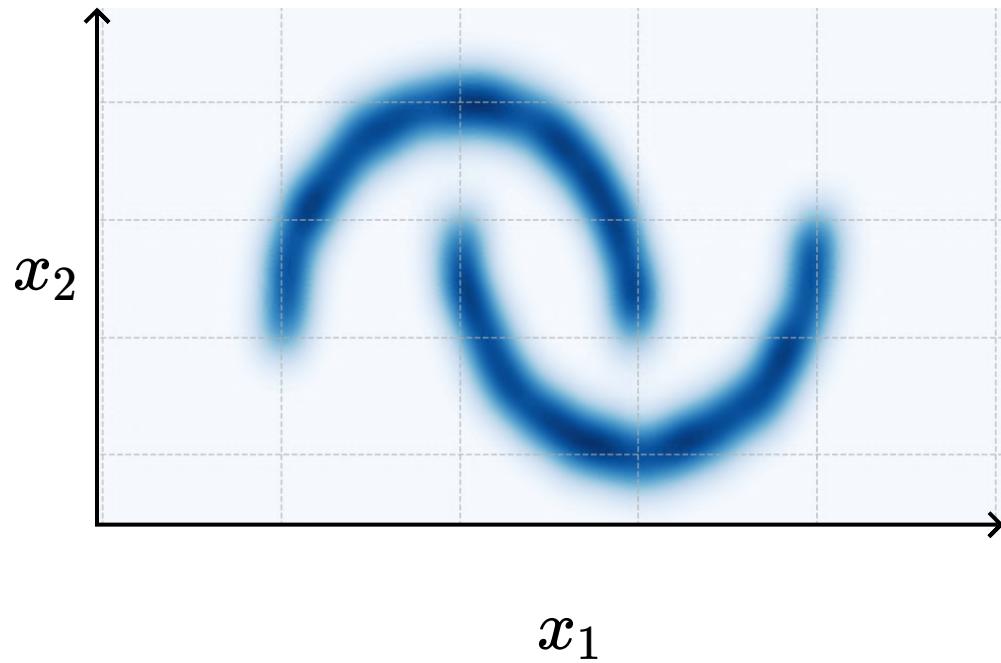
✗ Simulator

Generative Models: Simulate and Analyze

Fast emulators

Generate Novel Samples

Trained Model $p_\phi(x)$



✓ Generative Model

✓ Simulator

Evaluate probabilities



High Probability

Low Probability

✓ Generative Model

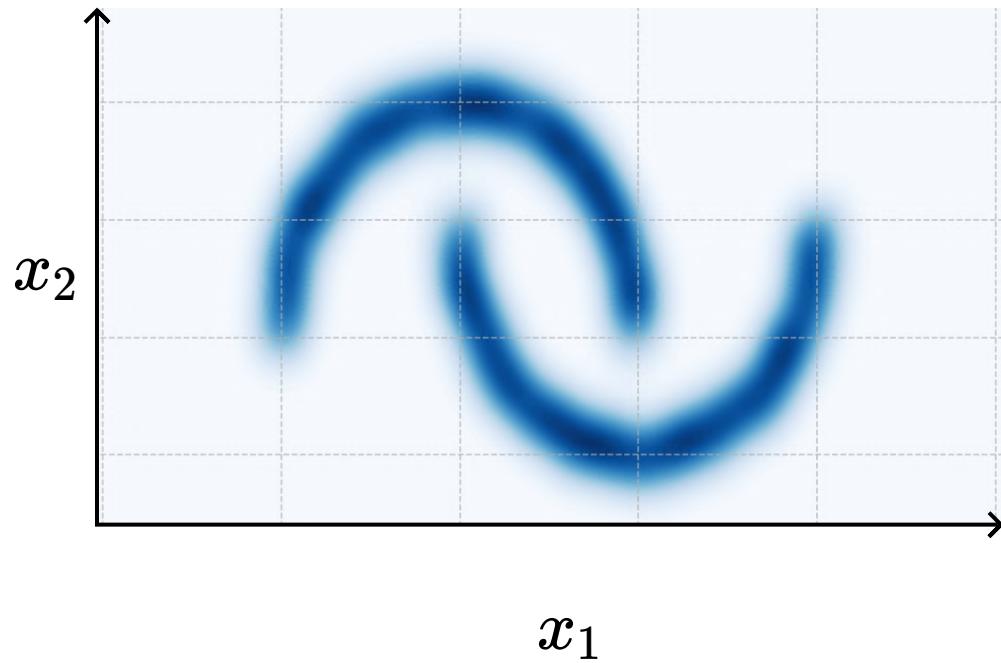
✗ Simulator

Generative Models: Simulate and Analyze

Fast emulators

Generate Novel Samples

Trained Model $p_\phi(x)$



✓ Generative Model

✓ Simulator

Testing Theories

Evaluate probabilities



High Probability

Low Probability

✓ Generative Model

✗ Simulator

Contrastive Learning

A folk music band of anthropomorphic autumn leaves playing bluegrass instrument

BigGAN

VAEs



GANS



2006

2014

2017

2019

2022

2023

Deep Belief
Networks

0 0
1 1
2 2
3 3

Normalising Flows



Diffusion Models



Contrastive Learning

A folk music band of anthropomorphic autumn leaves playing bluegrass instrument

BigGAN

VAEs



GANS



2006

2014

2017

2019

2022

2023

Deep Belief
Networks

0 0
1 1
2 2
3 3

Normalising Flows



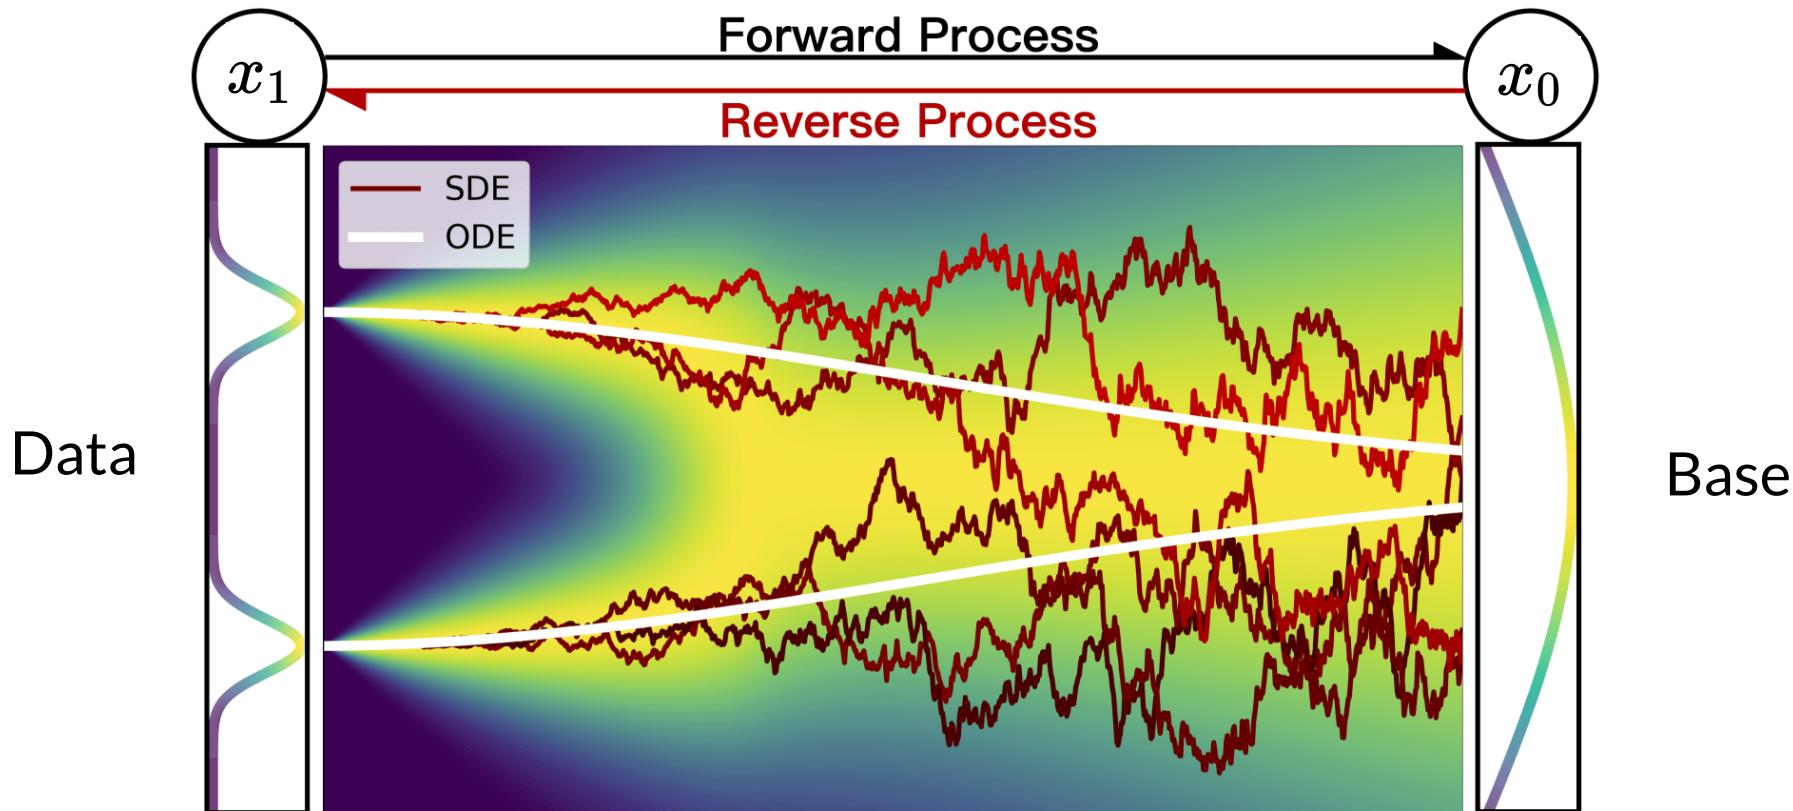
Diffusion Models



Bridging two distributions

*"Creating noise from data is easy;
creating data from noise is generative modeling."*

Yang Song

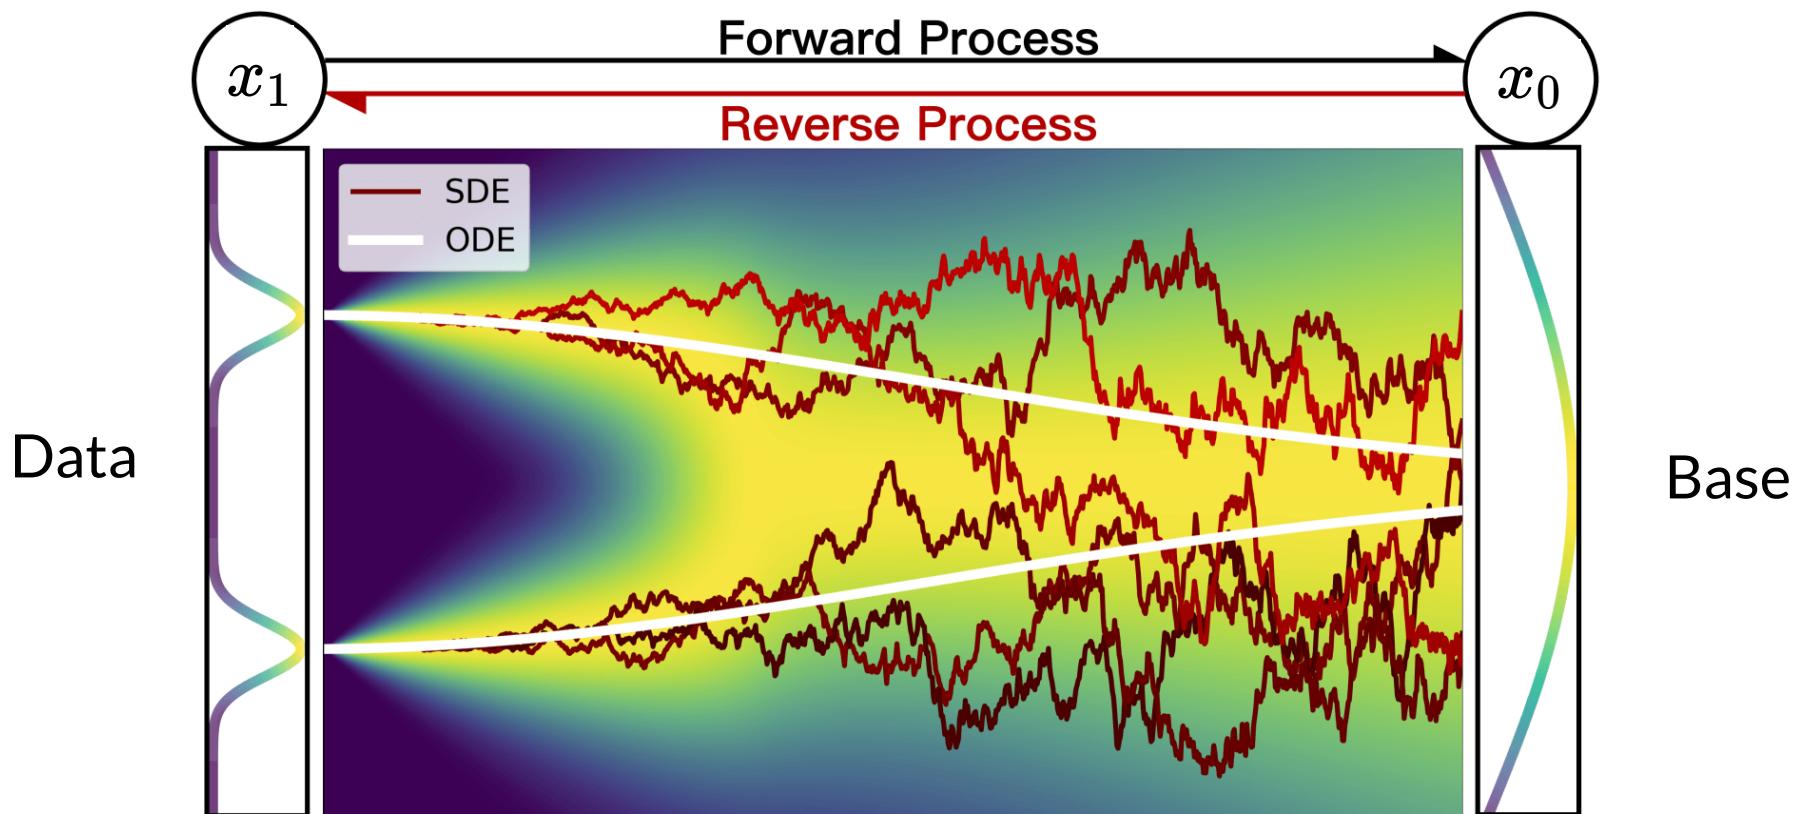


Bridging two distributions

is $p(x_0)$ restricted?

Diffusion: $p(x_0)$ is Gaussian

Normalising flows: $p(x_0)$ can
be evaluated



Bridging two distributions

How is the bridge constrained?

Normalizing flows: Reverse = Forward inverse

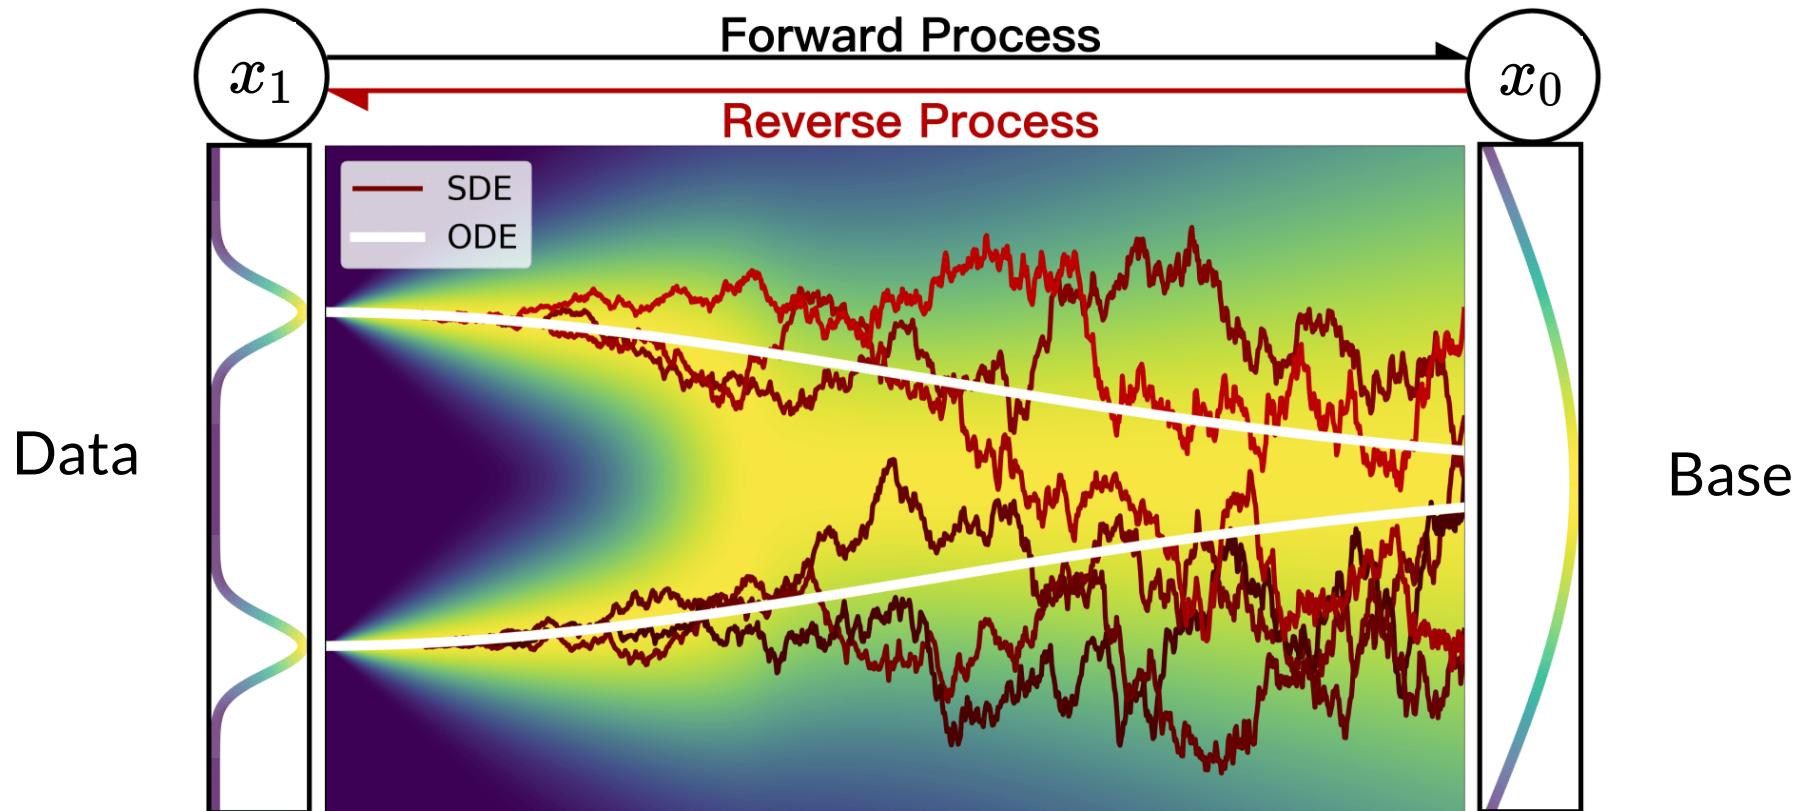
Diffusion: Forward = Gaussian noising

Flow Matching: Forward = Interpolant

is $p(x_0)$ restricted?

Diffusion: $p(x_0)$ is Gaussian

Normalising flows: $p(x_0)$ can
be evaluated



Bridging two distributions

How is the bridge constrained?

Normalizing flows: Reverse = Forward inverse

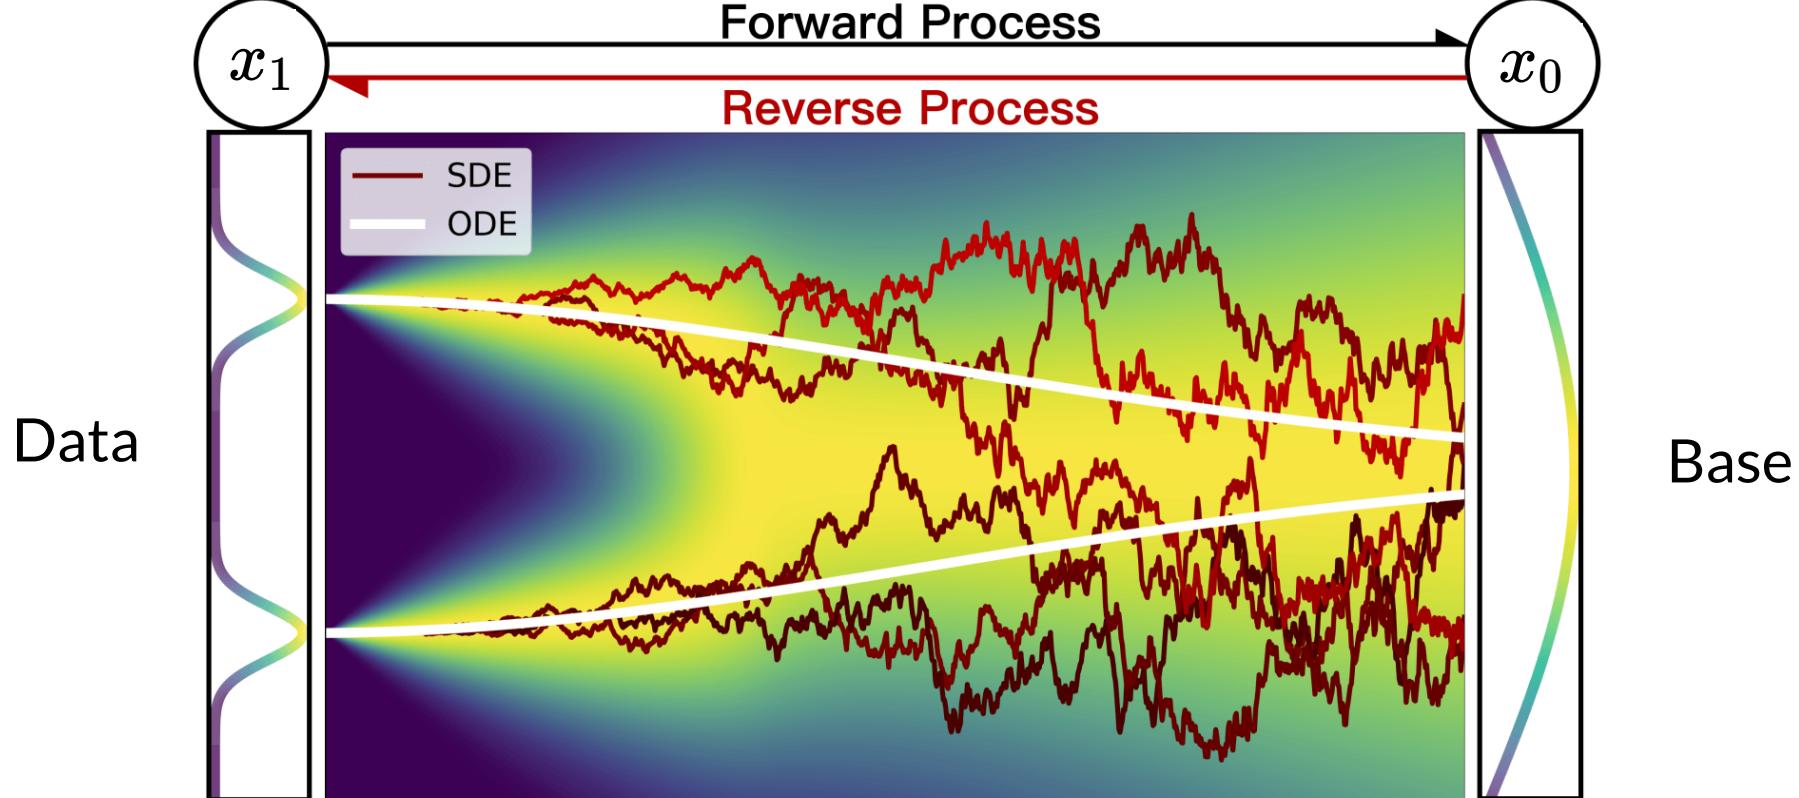
Diffusion: Forward = Gaussian noising

Flow Matching: Forward = Interpolant

is $p(x_0)$ restricted?

Diffusion: $p(x_0)$ is Gaussian

Normalising flows: $p(x_0)$ can
be evaluated



Is bridge stochastic (SDE) or deterministic (ODE)?

Diffusion: Stochastic (SDE)

Normalising flows: Deterministic (ODE)

Bridging two distributions

How is the bridge constrained?

Normalizing flows: Reverse = Forward inverse

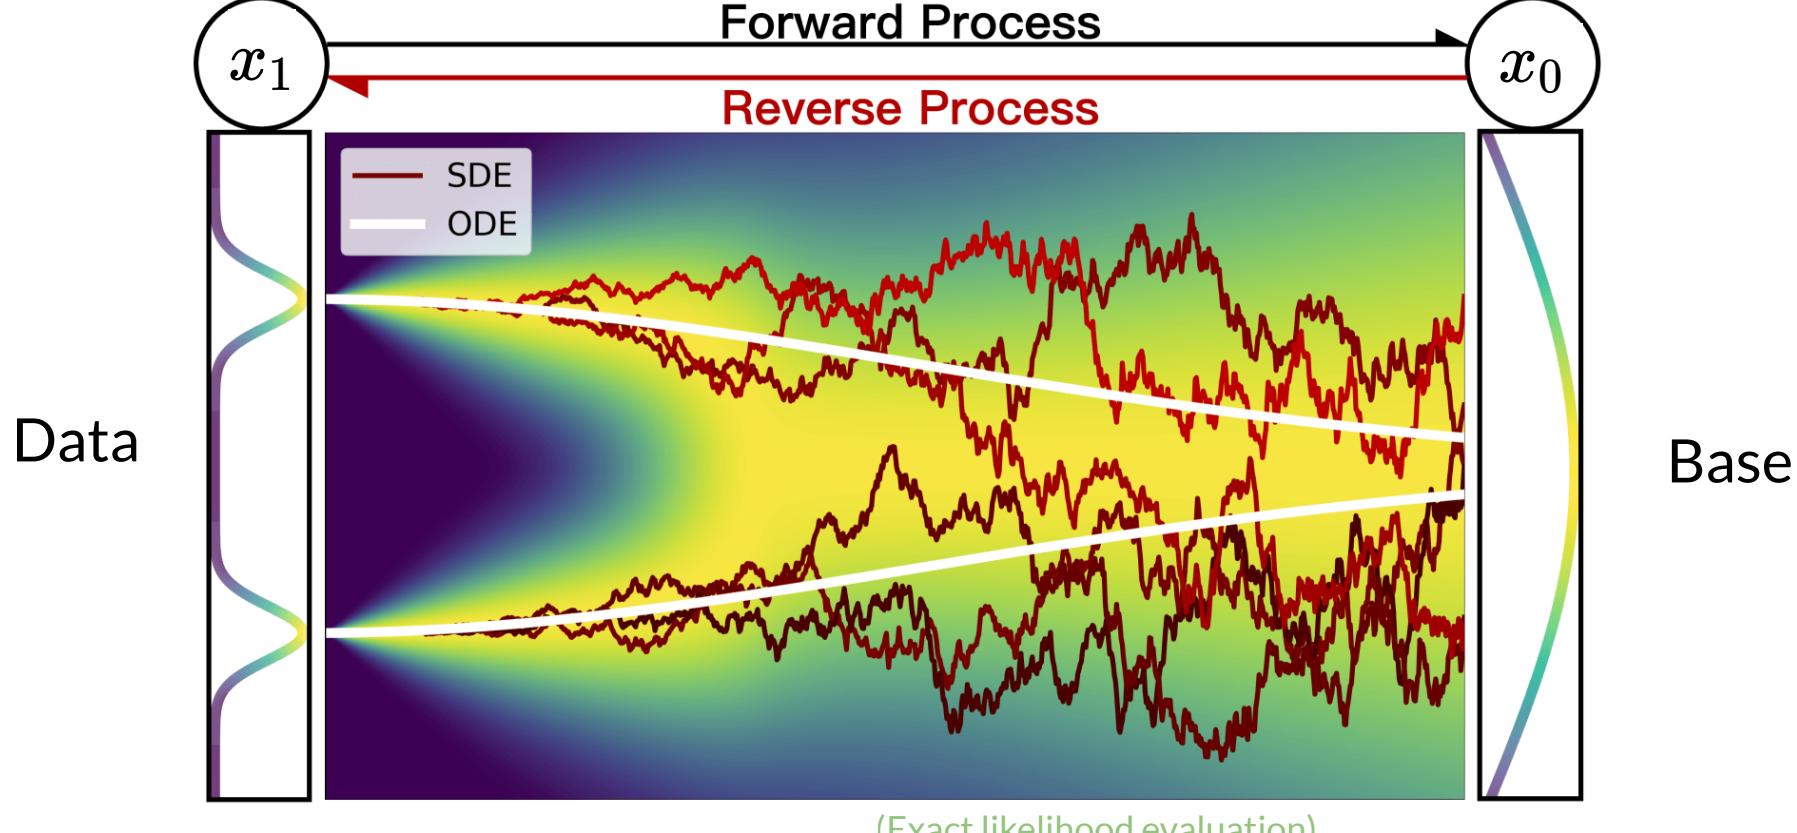
Diffusion: Forward = Gaussian noising

Flow Matching: Forward = Interpolant

is $p(x_0)$ restricted?

Diffusion: $p(x_0)$ is Gaussian

Normalising flows: $p(x_0)$ can
be evaluated



(Exact likelihood evaluation)

Is bridge stochastic (SDE) or deterministic (ODE)?

Diffusion: Stochastic (SDE)

Normalising flows: Deterministic (ODE)

Flow: Change of Variables

$$X \sim \mathcal{N}(\mu, \sigma)$$

Transformation (flow):

$$Y = g(X) = aX + b$$

How is Y distributed?

Flow: Change of Variables

$$X \sim \mathcal{N}(\mu, \sigma)$$

Transformation (flow):

$$Y = g(X) = aX + b$$

How is Y distributed?

$$Y \sim N(a\mu + b, a^2\sigma^2)$$

Flow: Change of Variables

$$X \sim \mathcal{N}(\mu, \sigma)$$

Transformation (flow):

$$Y = g(X) = aX + b$$

How is Y distributed?

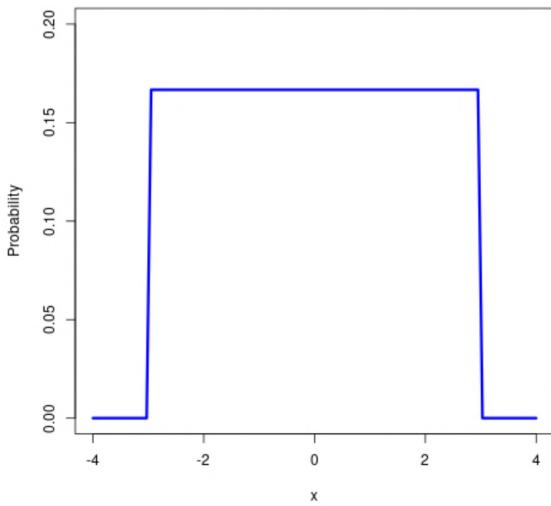
$$Y \sim N(a\mu + b, a^2\sigma^2)$$

$$p_Y(y) = p_X(g^{-1}(y)) \left| \frac{dg^{-1}(y)}{dy} \right|$$

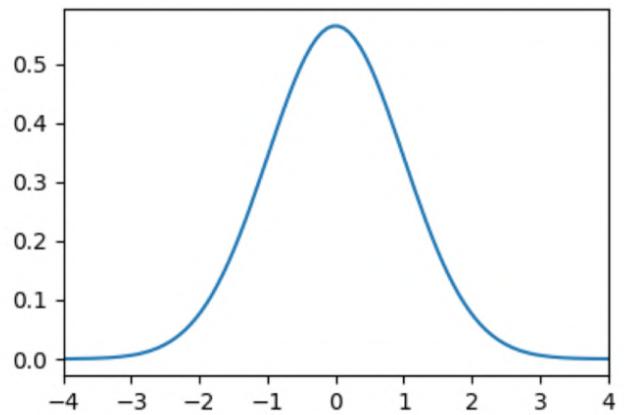
Box-Muller transform

Normalizing flows in 1934

$$\text{Uniform}(0, 1) \rightarrow U_1, U_2$$



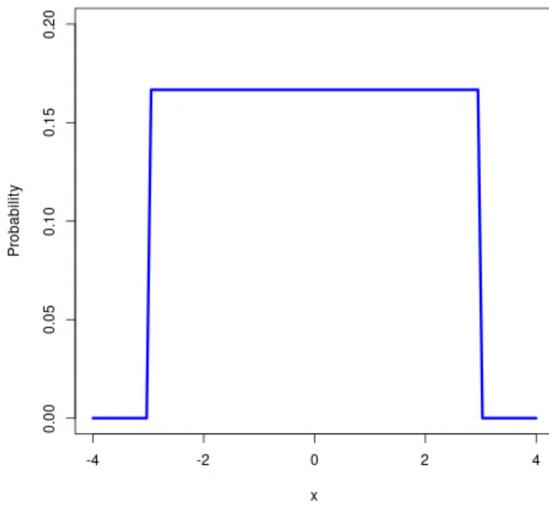
$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$



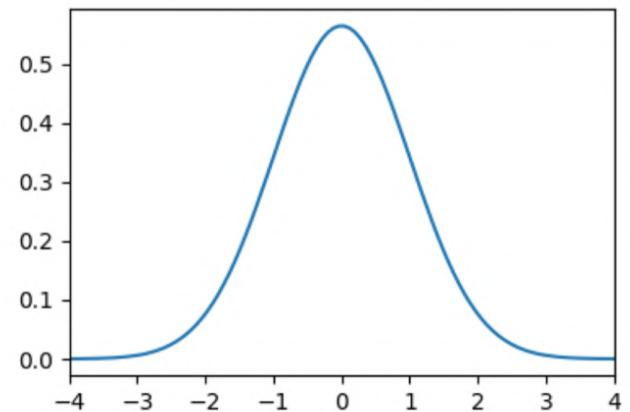
Box-Muller transform

Normalizing flows in 1934

$$\text{Uniform}(0, 1) \rightarrow U_1, U_2$$



$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$



$$Z_0 = \sqrt{-2 \ln U_1} \cos(2\pi U_2)$$

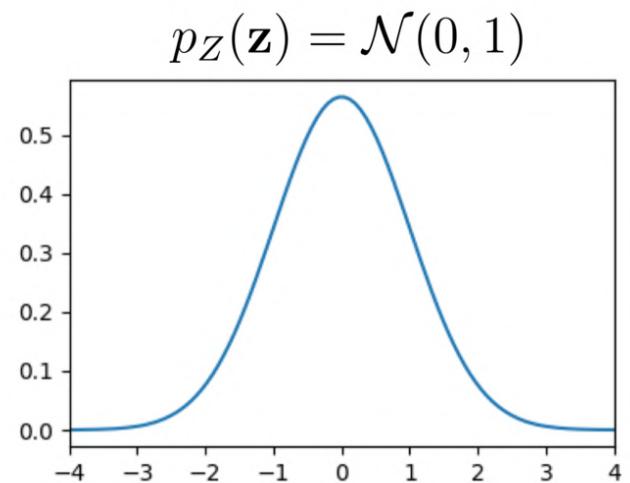
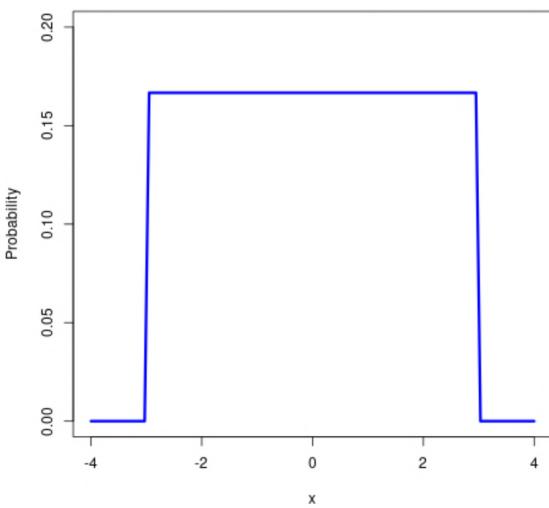
$$Z_1 = \sqrt{-2 \ln U_1} \sin(2\pi U_2)$$

Box-Muller transform

Normalizing flows in 1934

$$\text{Uniform}(0, 1) \rightarrow U_1, U_2$$

$$Z_0, Z_2 \leftarrow N(0, 1)$$

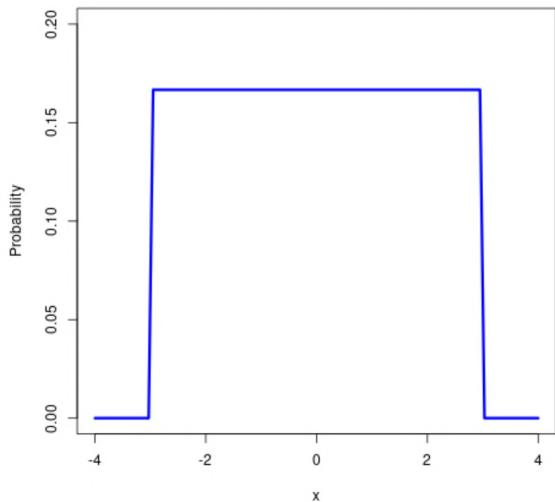


$$Z_0 = \sqrt{-2 \ln U_1} \cos(2\pi U_2)$$

$$Z_1 = \sqrt{-2 \ln U_1} \sin(2\pi U_2)$$

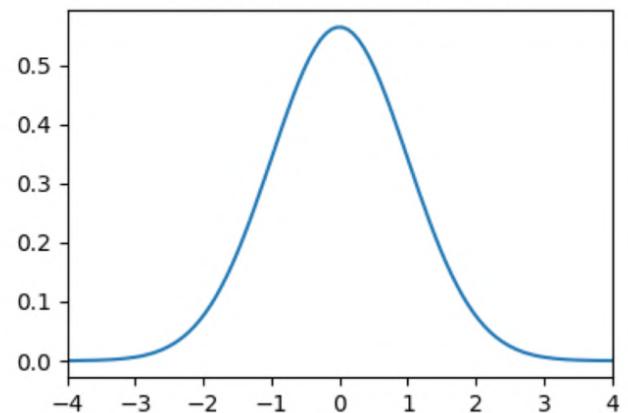
Tutorial 1

$\text{Uniform}(0, 1) \rightarrow U_1, U_2$



$Z_0, Z_2 \leftarrow N(0, 1)$

$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$



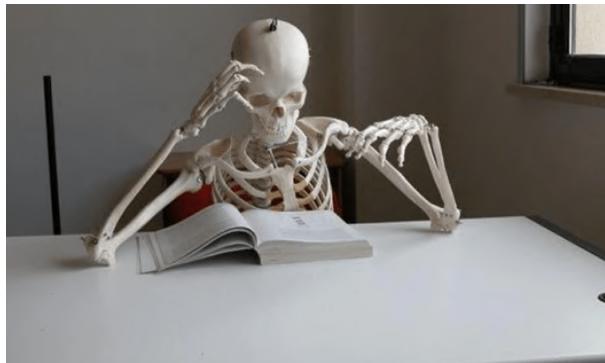
<https://tinyurl.com/heidelbergsschool>

On tutorials



Different people learn differently

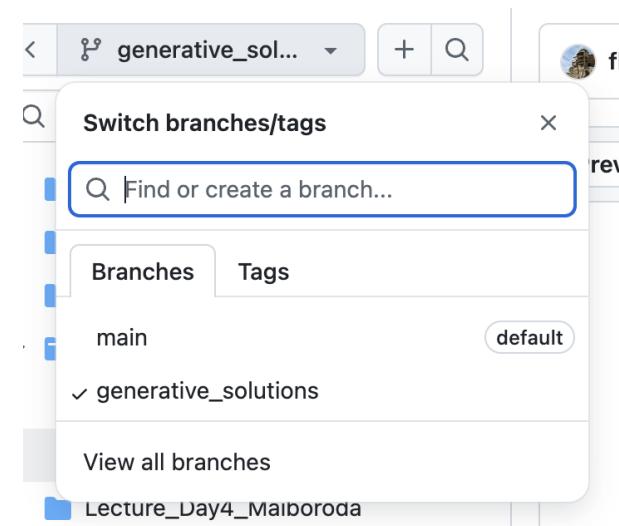
Take the data generator and start from scratch



Fill in the gaps

```
In [ ]: class FlowMatching:  
    def __init__(self, velocity_network, device='cuda'):  
        self.velocity_net = velocity_network  
        self.device = device  
  
    def sample_time(self, batch_size):  
        """Sample random time steps between 0 and 1."""  
        return torch.rand(batch_size, device=self.device)  
  
    def interpolate(self, x0, x1, t):  
        """  
        Linear interpolation between noise (x0) and data (x1).  
        At t=0: return x0 (noise)  
        At t=1: return x1 (data)  
        """  
        # TODO: fill in the interpolant again  
  
        # Your code here  
  
        return  
  
    def compute_target_velocity(self, x0, x1, t):  
        """  
        Compute the target velocity field that the neural network should learn.  
        Think about this: if we have  $x_t = (1-t)x_0 + t x_1$ ,  
        what is  $dx_t/dt$  (the derivative with respect to time)?  
        """  
        # Your code here  
        return # What is  $dx_t/dt$ ?
```

Start from a working solution and play around



On tutorials

Gemini

```
def interpolant(x0, x1, t):
    """
    #TODO:

    Implement the linear interpolation between noise (x0) and data (x1).

    At t=0: we should get pure noise (x0)
    At t=1: we should get real data (x1)
    At t=1: we should get real data (x1)
    At t=0.5: we should get a 50/50 mix

    Args:
        x0: Starting point (noise), shape: (batch_size, ...)
        x1: End point (data), shape: (batch_size, ...)
        t: Time parameter, shape: (batch_size,) or scalar

    Returns:
        x_t: Interpolated samples at time t
    """
    # Your code here
    return
    if t.ndim == 1 and x0.ndim > 1:
        # Reshape t for broadcasting if it's a 1D tensor
        t = t.view(-1, *[1] * (x0.ndim - 1))
    return (1 - t) * x0 + t * x1

fig, axs = plt.subplots(1, 10, figsize=(20, 2))

x1, _ = make_moons(n_samples=500, noise=data_std)
x1 = torch.tensor(x1).float()
x0 = torch.randn_like(x1)
t = torch.tensor(np.linspace(0., 1., 10))
for i, ax in enumerate(axs):
    xt = interpolant(x0, x1, t[i])
    ax.scatter(xt[:, 0], xt[:, 1], s=1, color='k')
    ax.axis('off')
    ax.set_title(f't={t[i]:.2f}')


t=0.00          t=0.11          t=0.22          t=0.33          t=0.44          t=0.55          t=0.66          t=0.77          t=0.88          t=1.00
```

c fill in TODOs

Sure, I can help with that. I'll start by filling in the `interpolant` function.

Accept & Run Accept Cancel

What can I help you build?

Gemini can make mistakes so double-check it and use code with caution. [Learn more](#)



On tutorials



C Please explain this code:

```
<> class FlowMatching:
```

◆ Certainly! The selected code defines the `FlowMatching` class, which implements the core logic for training and sampling from a Flow Matching model. Let's break down its key methods:

- `__init__(self, velocity_network, device='cuda')`: This is the constructor. It initializes the `FlowMatching` object with the provided `velocity network` (which is a neural network that

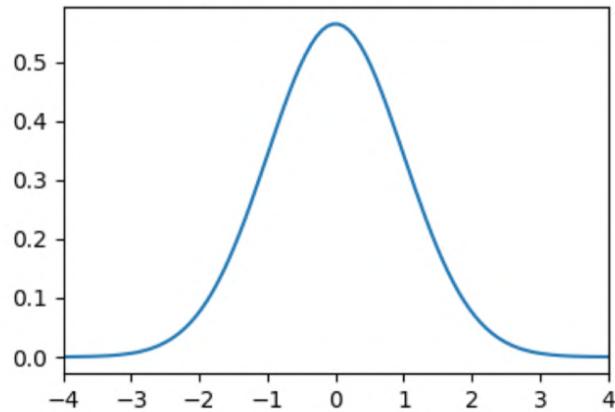
◆ What can I help you build?



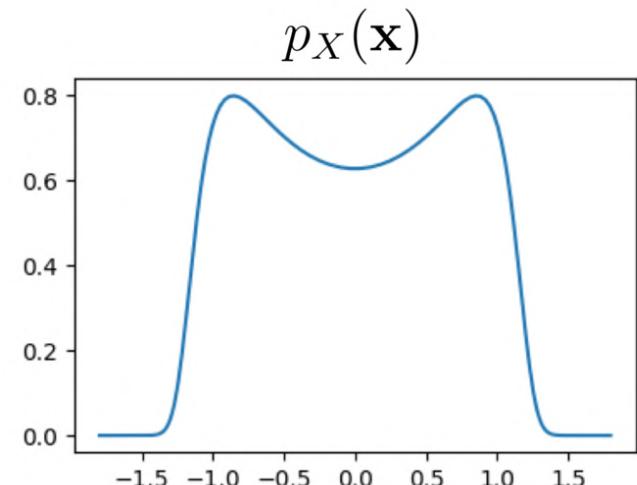
Gemini can make mistakes so double-check it and use code with caution. [Learn more](#)

Normalizing flows

$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$



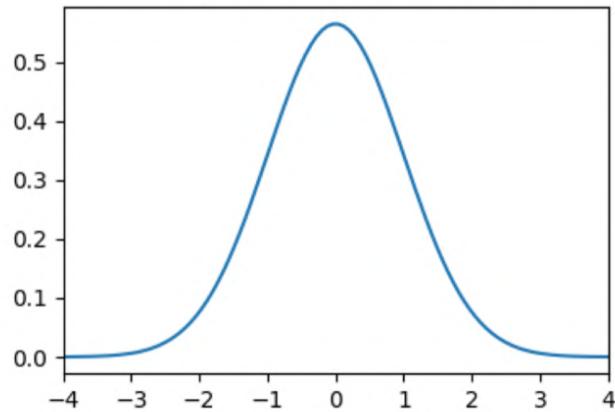
Base distribution



Target distribution

Normalizing flows

$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$

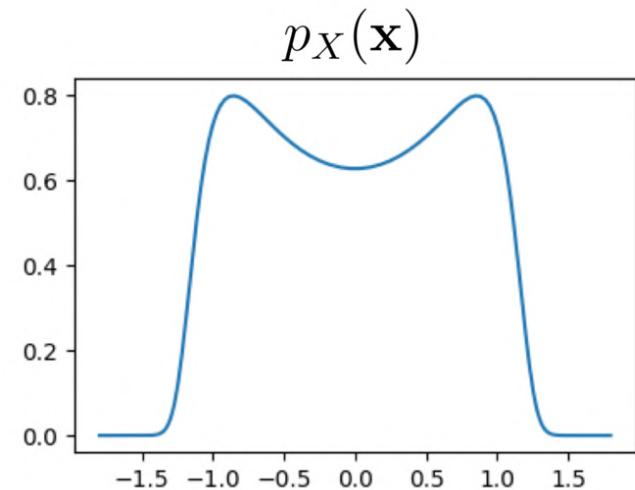
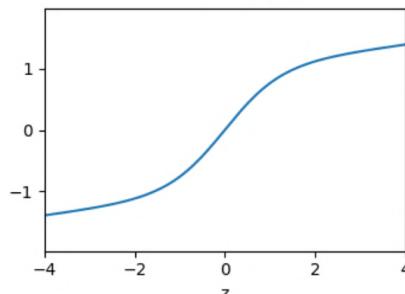


Base distribution

*Invertible
transformation*



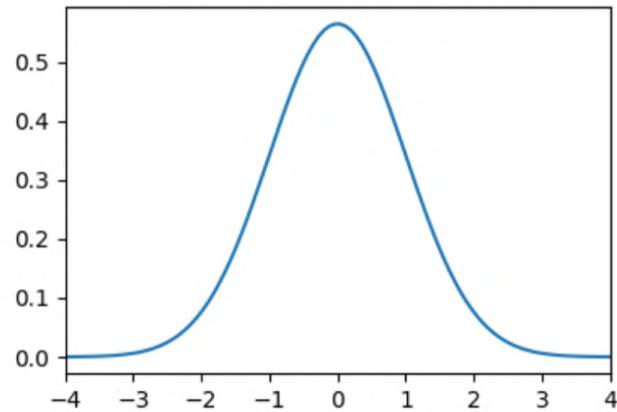
$$\mathbf{x} = g(\mathbf{z})$$



Target distribution

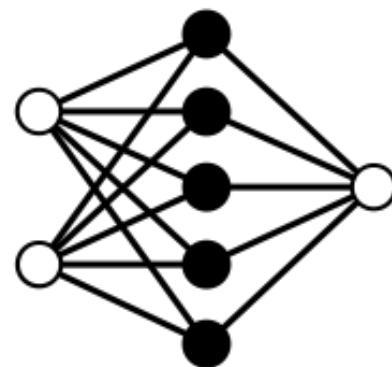
Normalizing flows

$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$

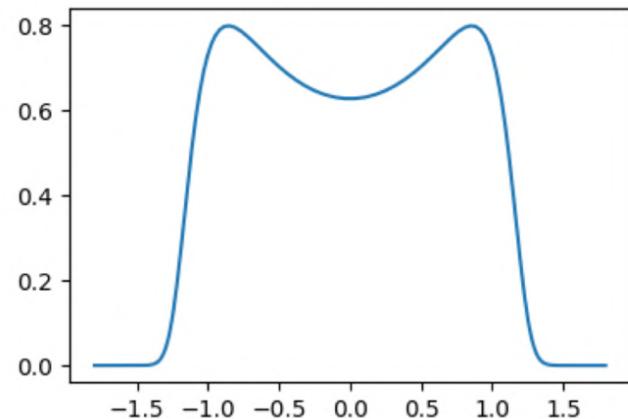


Base distribution

*Invertible
transformation*



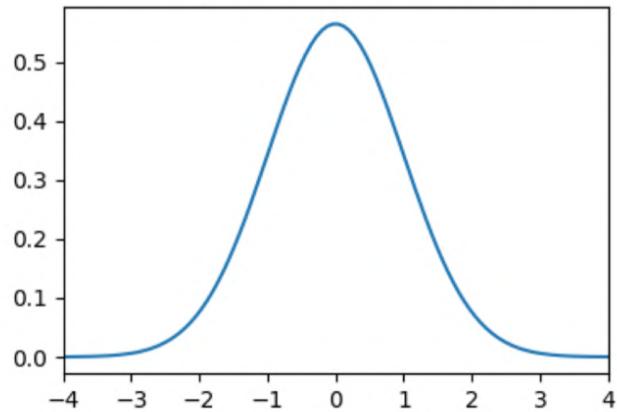
$$p_X(\mathbf{x})$$



Target distribution

Normalizing flows

$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$

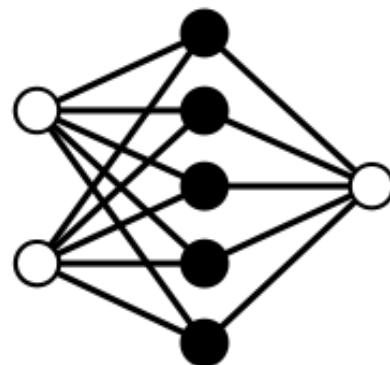


Base distribution

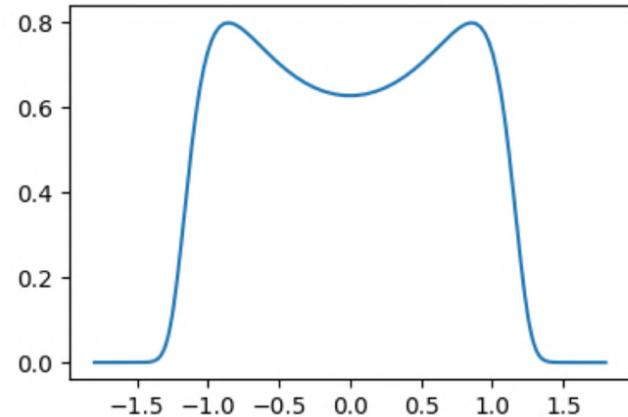


$$z \sim p_Z(z)$$

Invertible transformation



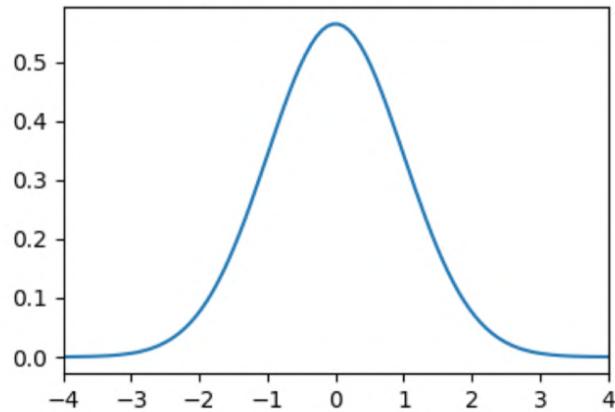
$$p_X(\mathbf{x})$$



Target distribution

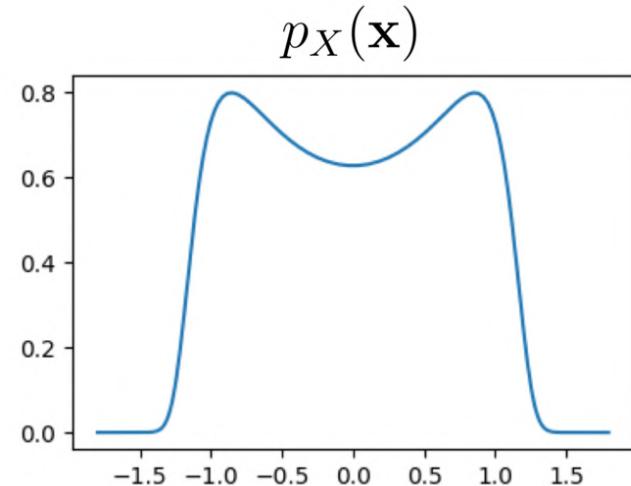
Normalizing flows

$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$



Base distribution

*Invertible
transformation*



Target distribution



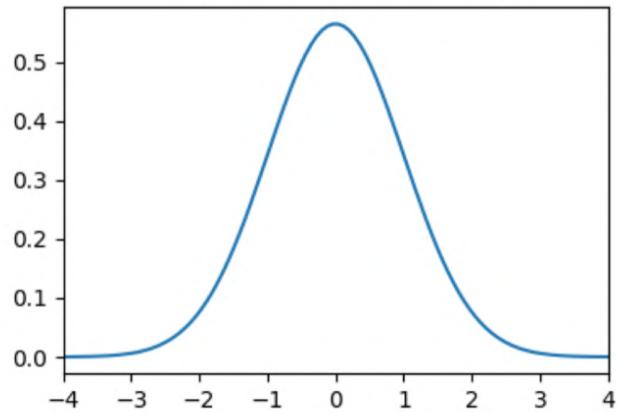
$$z \sim p_Z(z)$$



$$p_Z(z)$$

Normalizing flows

$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$



Base distribution

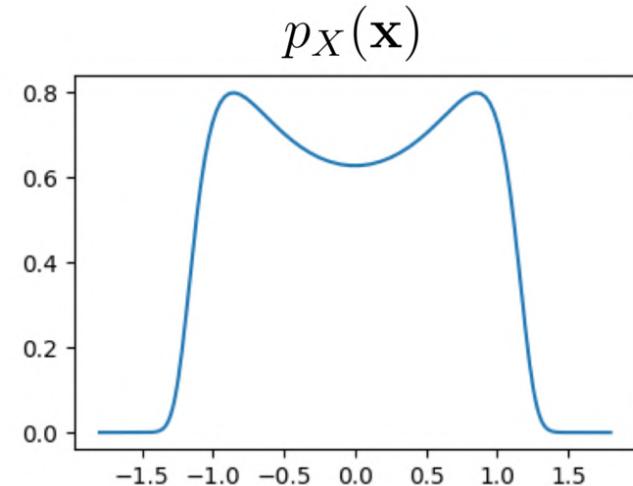
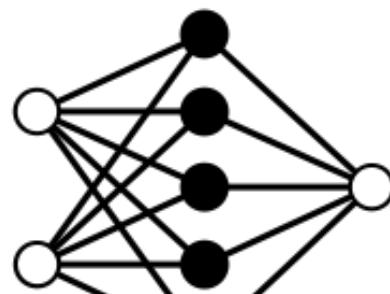


$$z \sim p_Z(z)$$



$$p_Z(z)$$

Invertible transformation



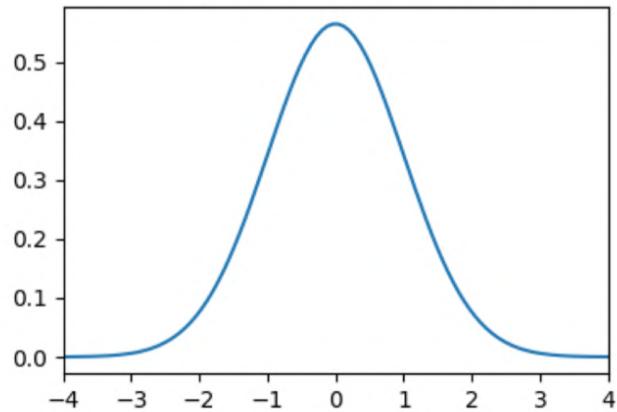
Target distribution



$$Z \sim \mathcal{N}(0, 1) \rightarrow g(z) \rightarrow X$$

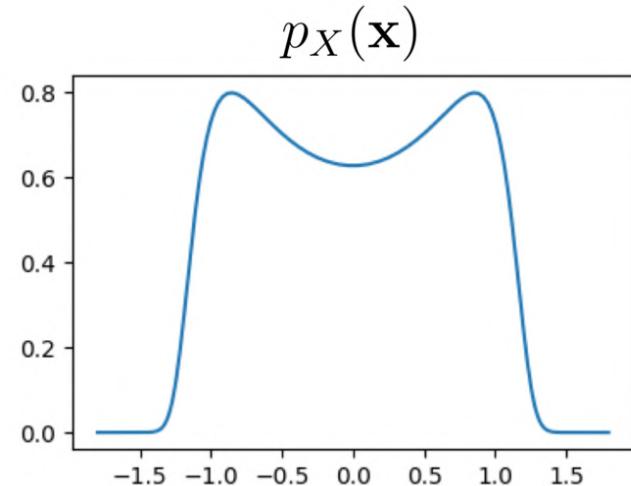
Normalizing flows

$$p_Z(\mathbf{z}) = \mathcal{N}(0, 1)$$



Base distribution

Invertible transformation



Target distribution



$$z \sim p_Z(z)$$



$$p_Z(z)$$



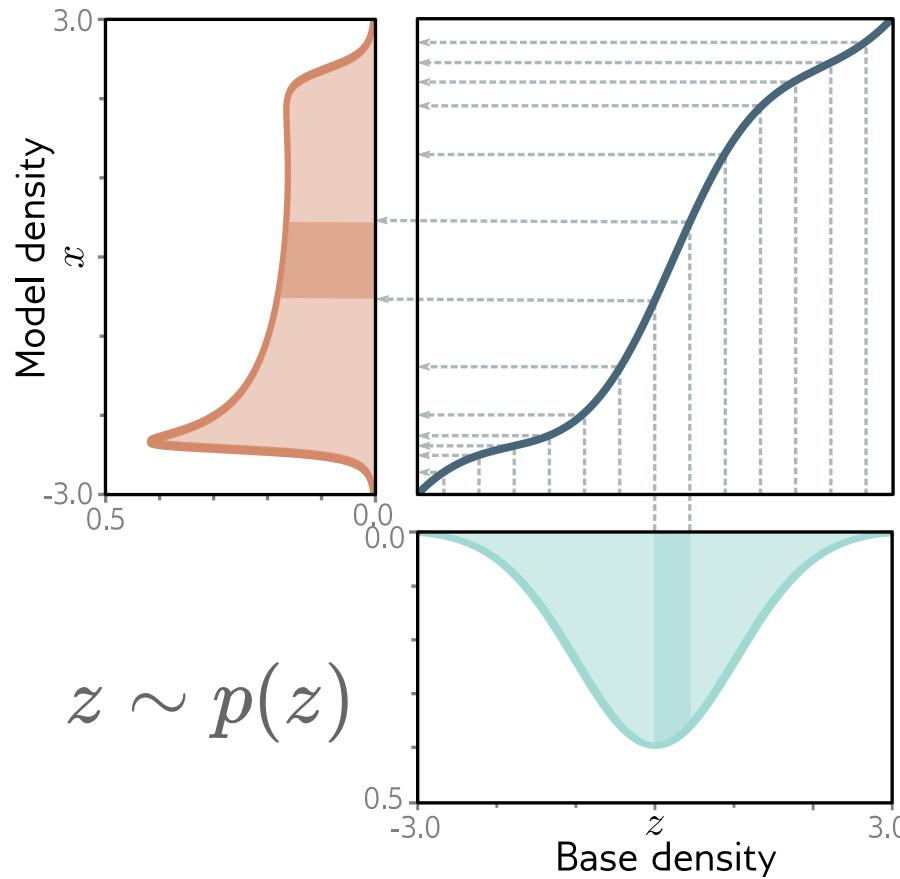
$$Z \sim \mathcal{N}(0, 1) \rightarrow g(z) \rightarrow X$$



$$p_X(x) = p_Z(z) \left| \frac{dz}{dx} \right|$$

Normalizing flows

$$x \sim p(x)$$

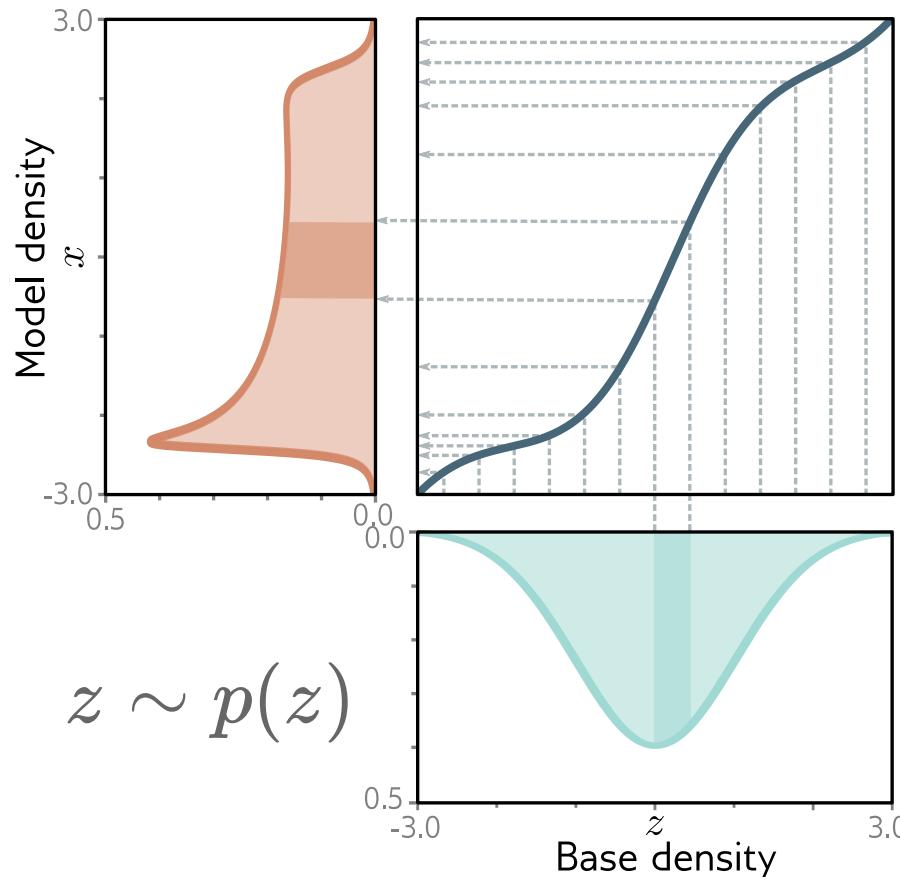


Bijective
 $x = f(z)$

[Image Credit:
"Understanding Deep
Learning" Simon J.D.
Prince]

Normalizing flows

$$x \sim p(x)$$



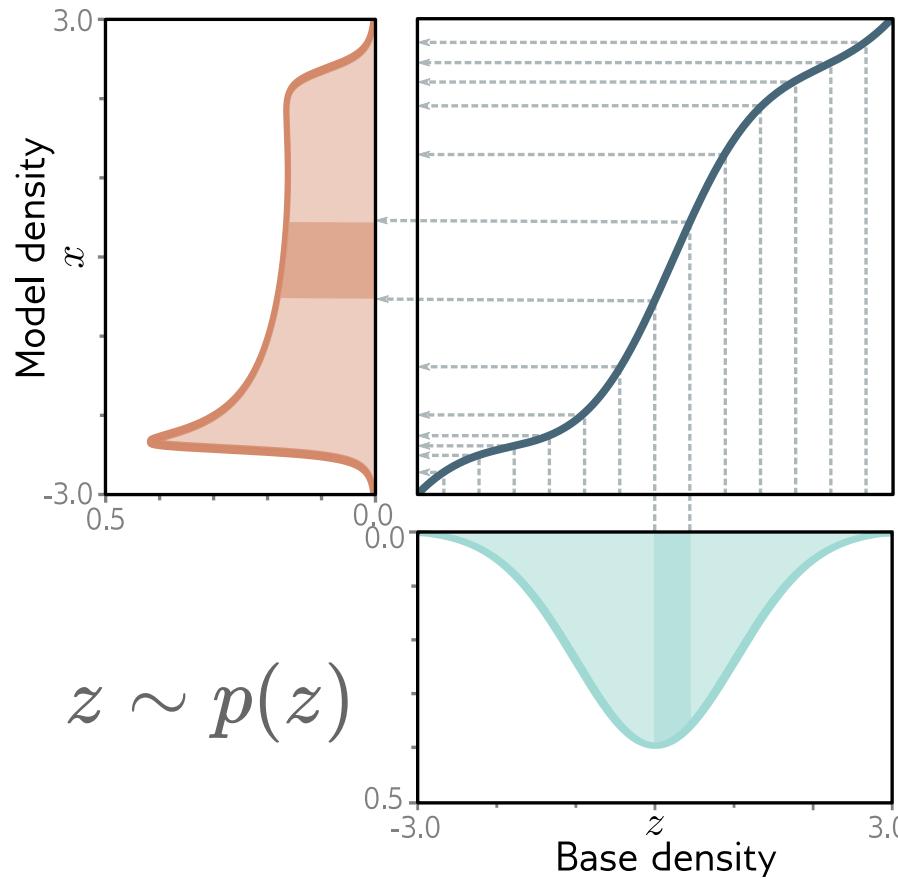
Bijective
 $x = f(z)$ *Sample*

[Image Credit:
"Understanding Deep
Learning" Simon J.D.
Prince]

Normalizing flows

Bijective
 $x = f(z)$ *Sample*

$$x \sim p(x)$$



$$z \sim p(z)$$

[Image Credit:
"Understanding Deep
Learning" Simon J.D.
Prince]

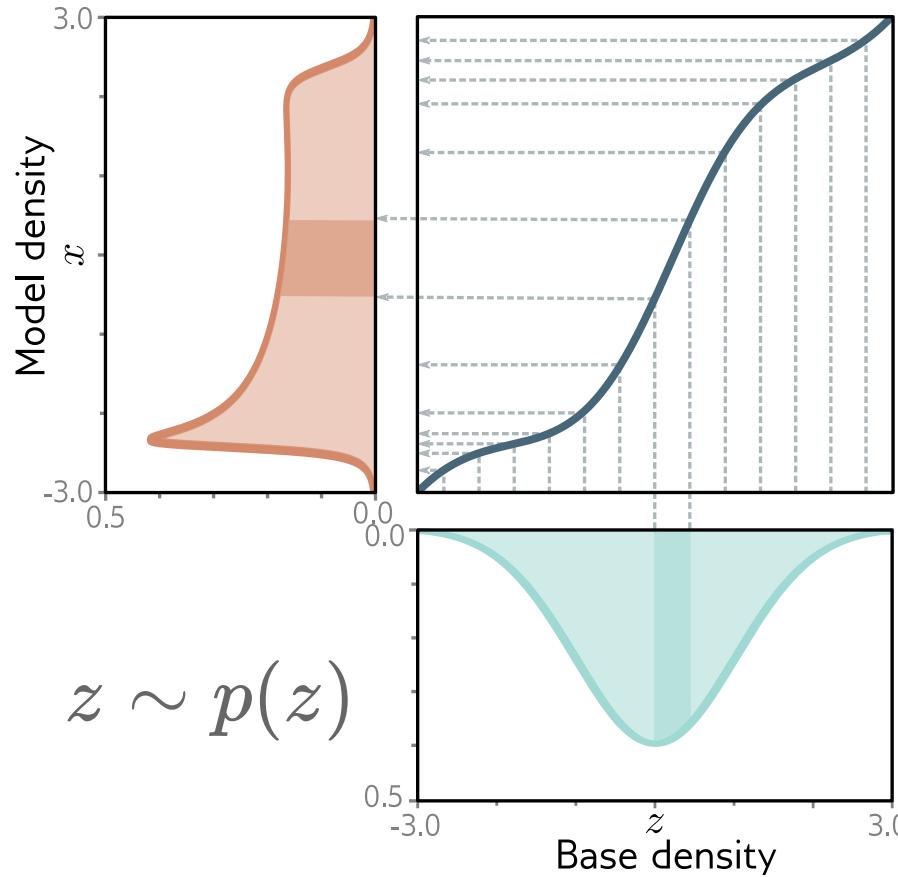
Probability mass conserved locally

$$p(x) = p(z = f^{-1}(x)) |\det J_{f^{-1}}(x)|$$

Normalizing flows

Bijective
 $x = f(z)$  *Sample*

$$x \sim p(x)$$



$$z \sim p(z)$$

Probability mass conserved locally

$$p(x) = p(z = f^{-1}(x)) |\det J_{f^{-1}}(x)|$$
  *Evaluate probabilities*

[Image Credit:
"Understanding Deep
Learning" Simon J.D.
Prince]

$$z_0 \sim p(z)$$

$$z_k = f_k(z_{k-1})$$

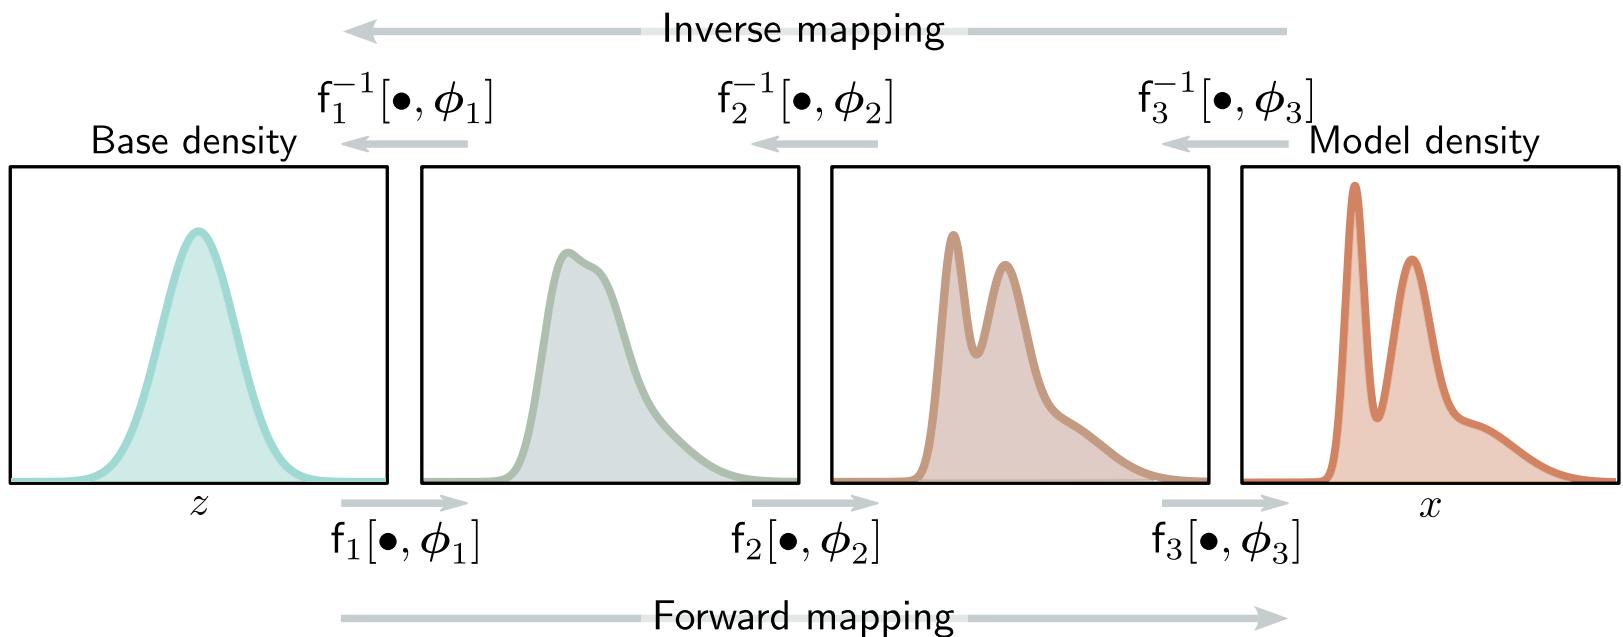
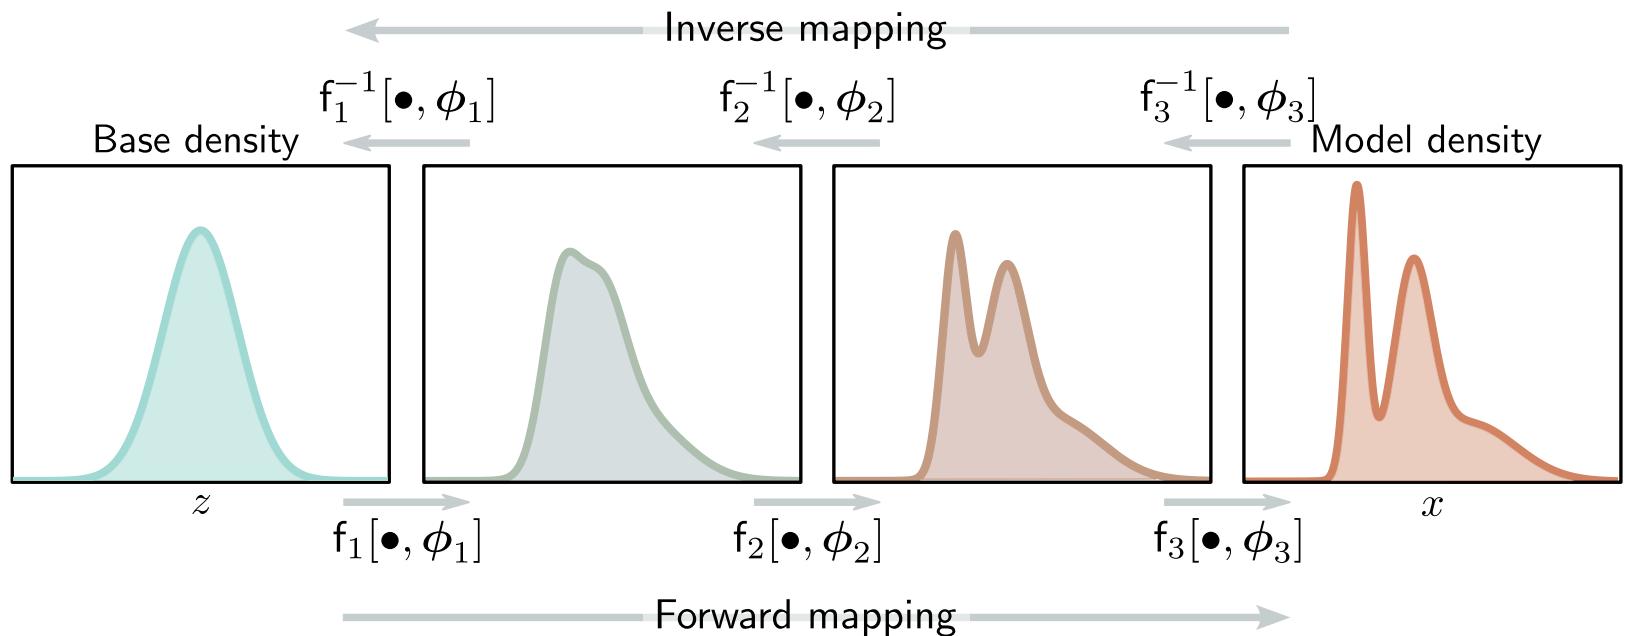


Image Credit: "Understanding Deep Learning" Simon J.D. Prince

$$z_0 \sim p(z)$$

$$z_k = f_k(z_{k-1})$$

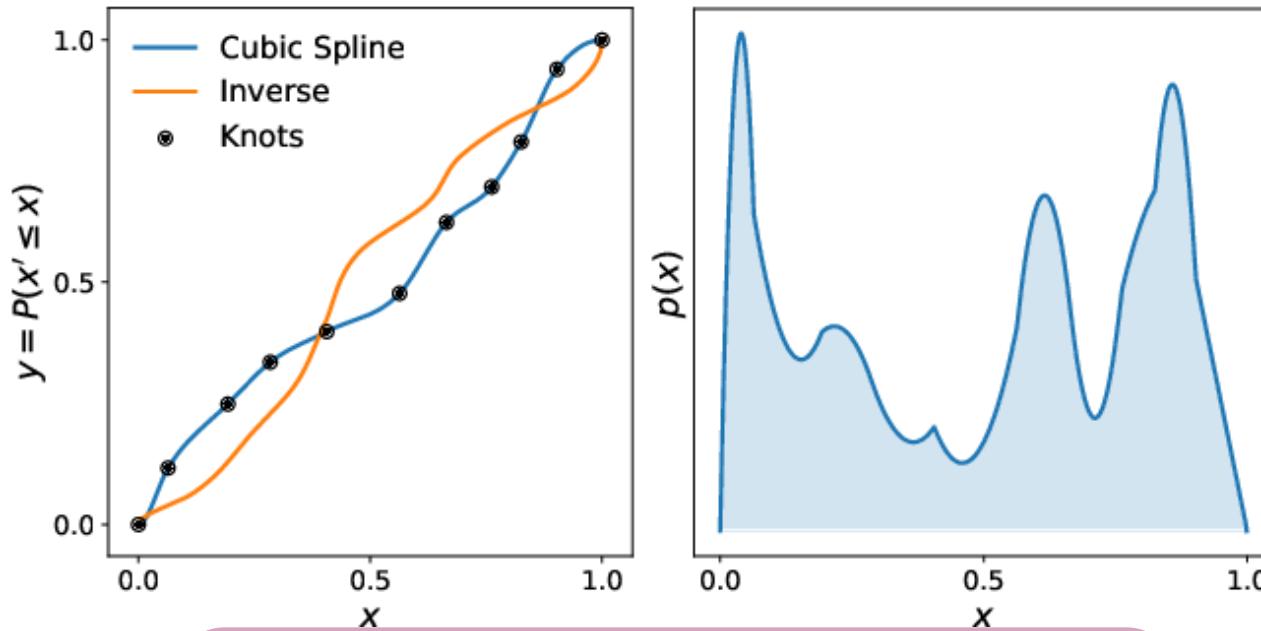


$$\log p(x) = \log p(z_0) - \sum_{k=1}^K \log |\det J_{f_k}(z_{k-1})|$$

Image Credit: "Understanding Deep Learning" Simon J.D. Prince

Invertible functions aren't that common!

Splines



Issues NFs: Lack of flexibility

- Invertible functions
- Tractable Jacobians

Masked Autoregressive Flows

$$p(x) = \prod_i p(x_i | x_{1:i-1})$$

Masked Autoregressive Flows

$$p(x) = \prod_i p(x_i | x_{1:i-1})$$

$$p(x_i | x_{1:i-1}) = \mathcal{N}(x_i | \mu_i, (\exp \alpha_i)^2)$$

Masked Autoregressive Flows

$$p(x) = \prod_i p(x_i | x_{1:i-1}) \quad \mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$
$$p(x_i | x_{1:i-1}) = \mathcal{N}(x_i | \mu_i, (\exp \alpha_i)^2)$$

Neural Network

Masked Autoregressive Flows

Neural Network

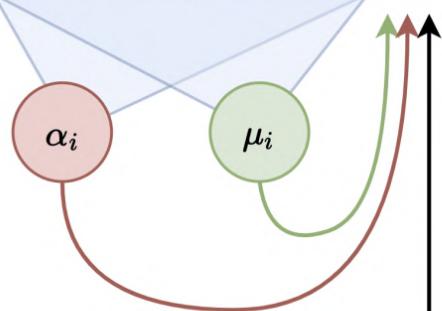
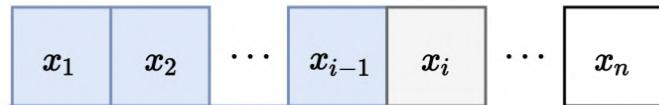
$$p(x) = \prod_i p(x_i | x_{1:i-1})$$

$$\mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$

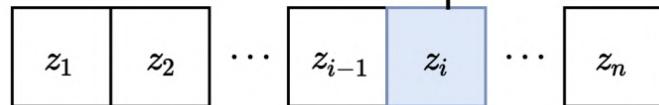
$$p(x_i | x_{1:i-1}) = \mathcal{N}(x_i | \mu_i, (\exp \alpha_i)^2)$$

✓ *Sample*

Transformed distribution



Base distribution



$$x_i = z_i \exp(\alpha_i) + \mu_i$$

Masked Autoregressive Flows

$$p(x) = \prod_i p(x_i | x_{1:i-1})$$

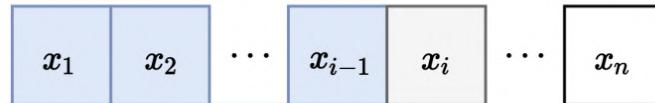
Neural Network

$$\mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$

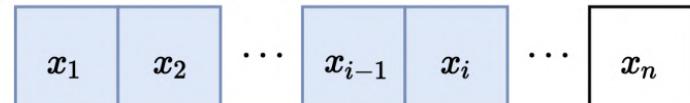
$$p(x_i | x_{1:i-1}) = \mathcal{N}(x_i | \mu_i, (\exp \alpha_i)^2)$$

✓ *Sample*

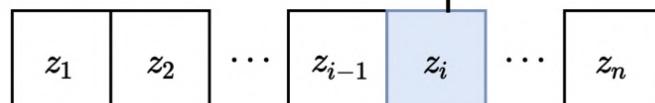
Transformed distribution



✓ *Evaluate probabilities*



Base distribution

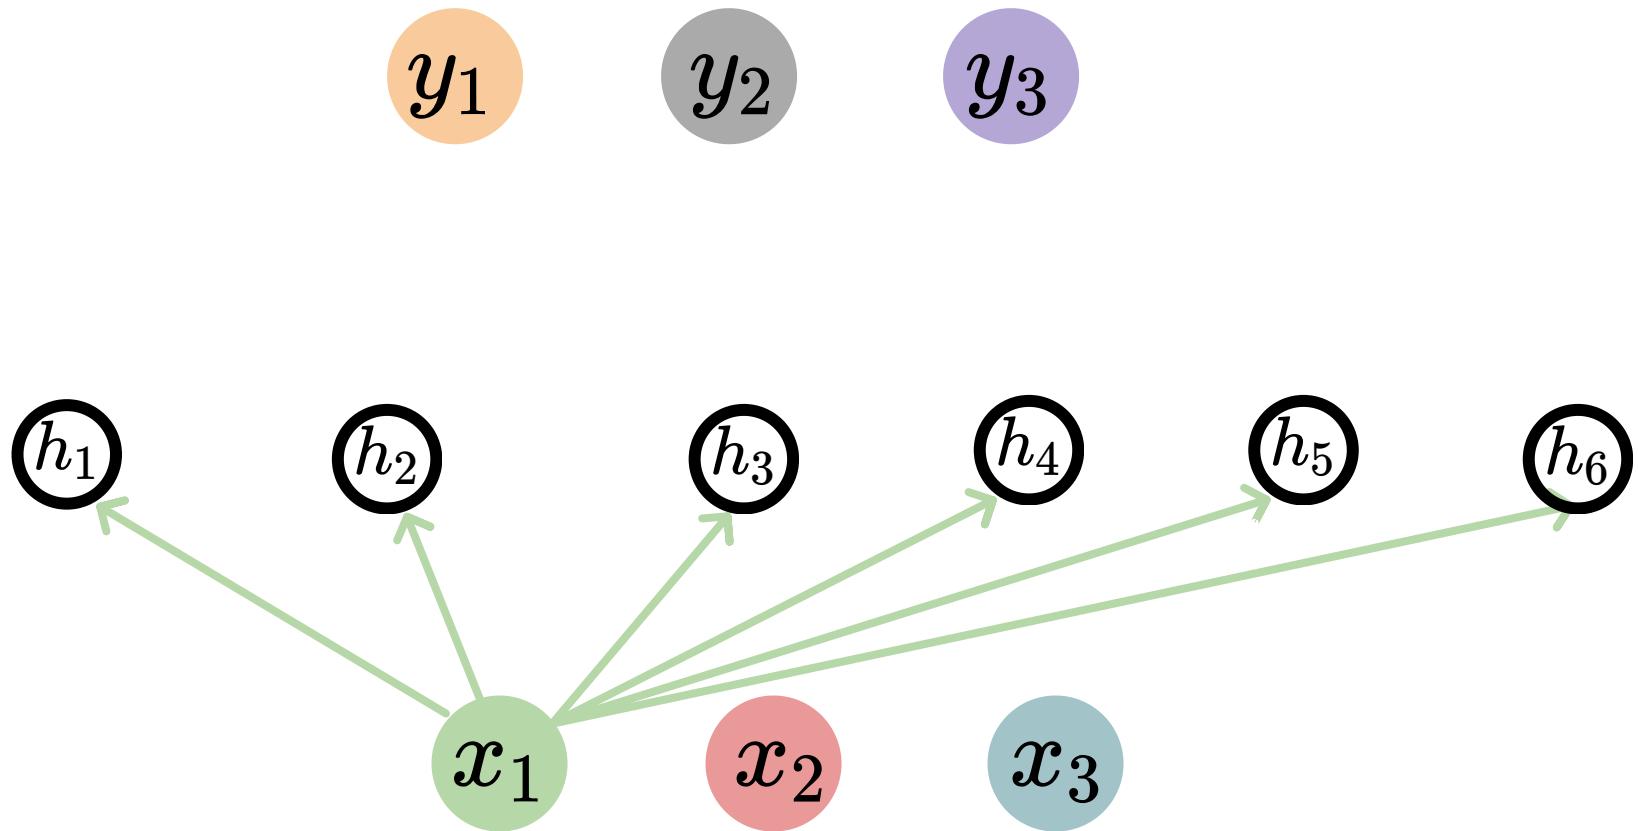


$$x_i = z_i \exp(\alpha_i) + \mu_i$$

$$z_i = (x_i - \mu_i) \exp(-\alpha_i)$$

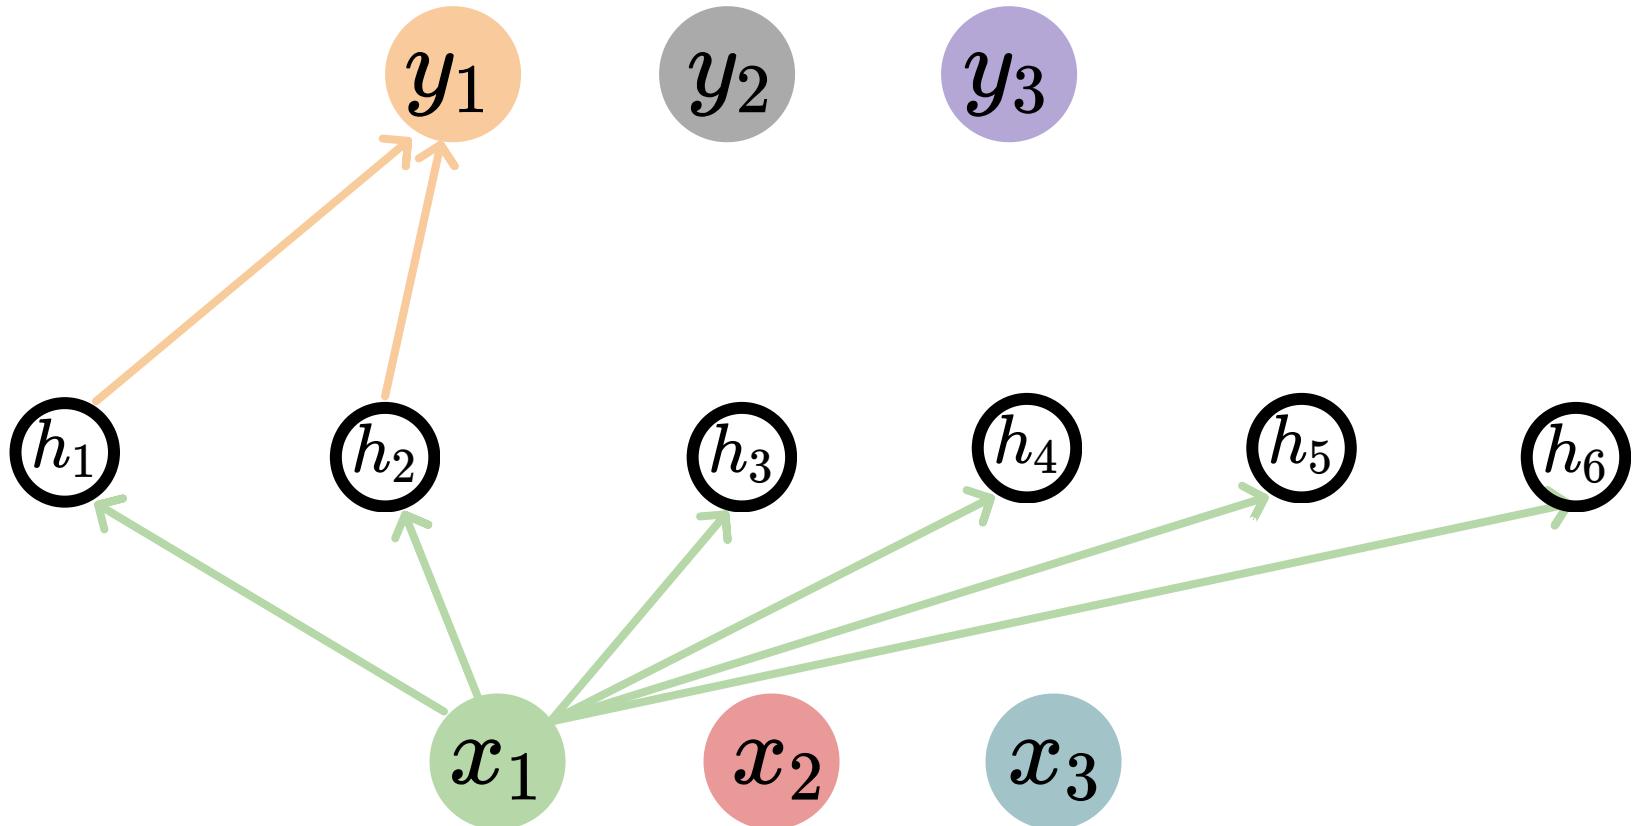
Masked MLP

$$y_i = \mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$



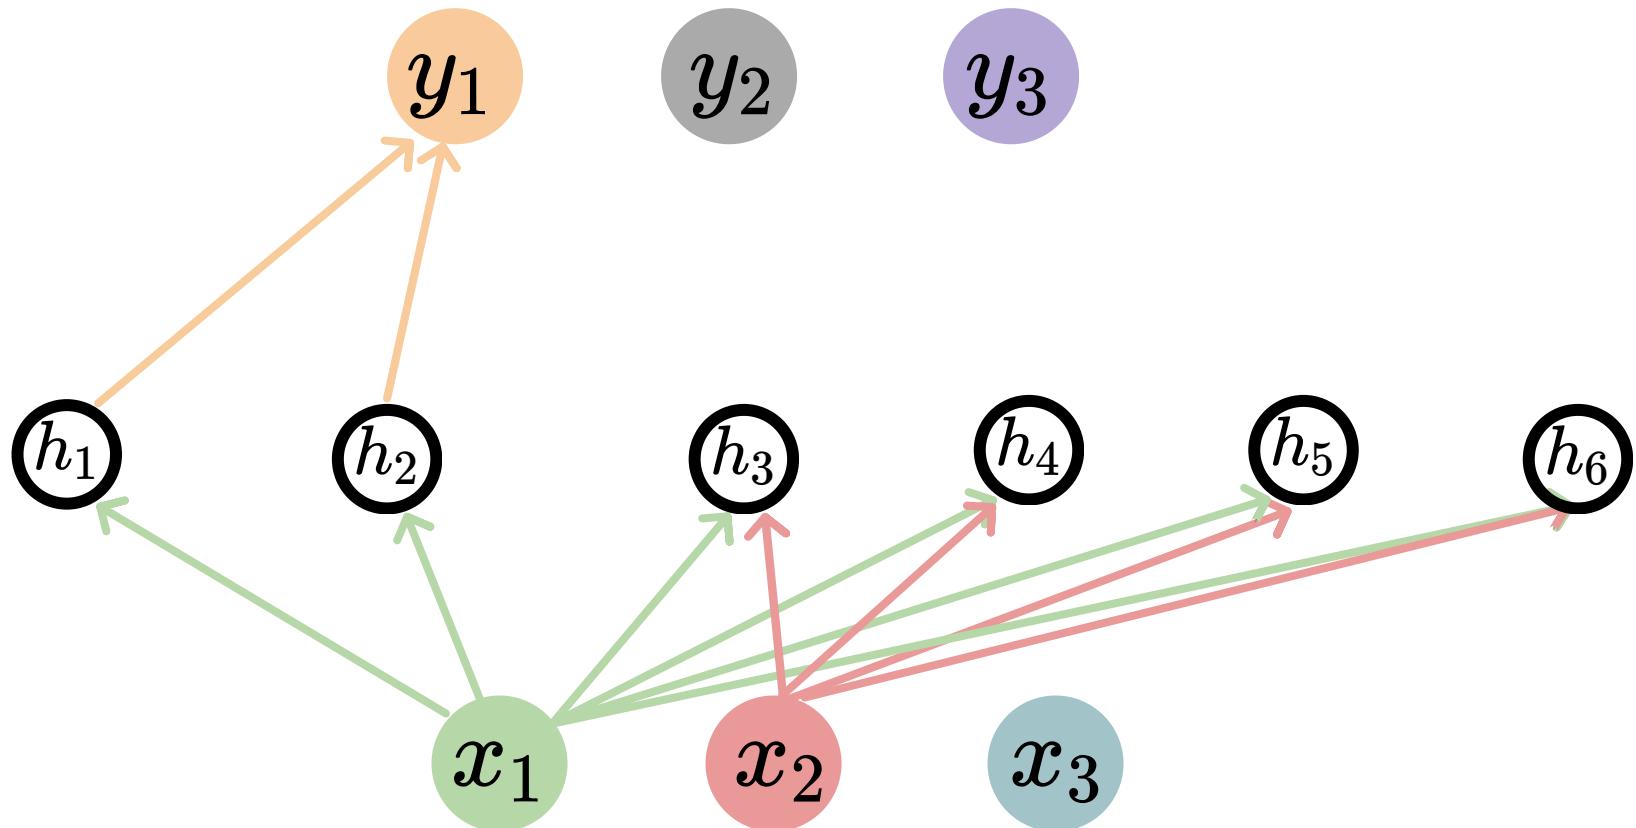
Masked MLP

$$y_i = \mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$



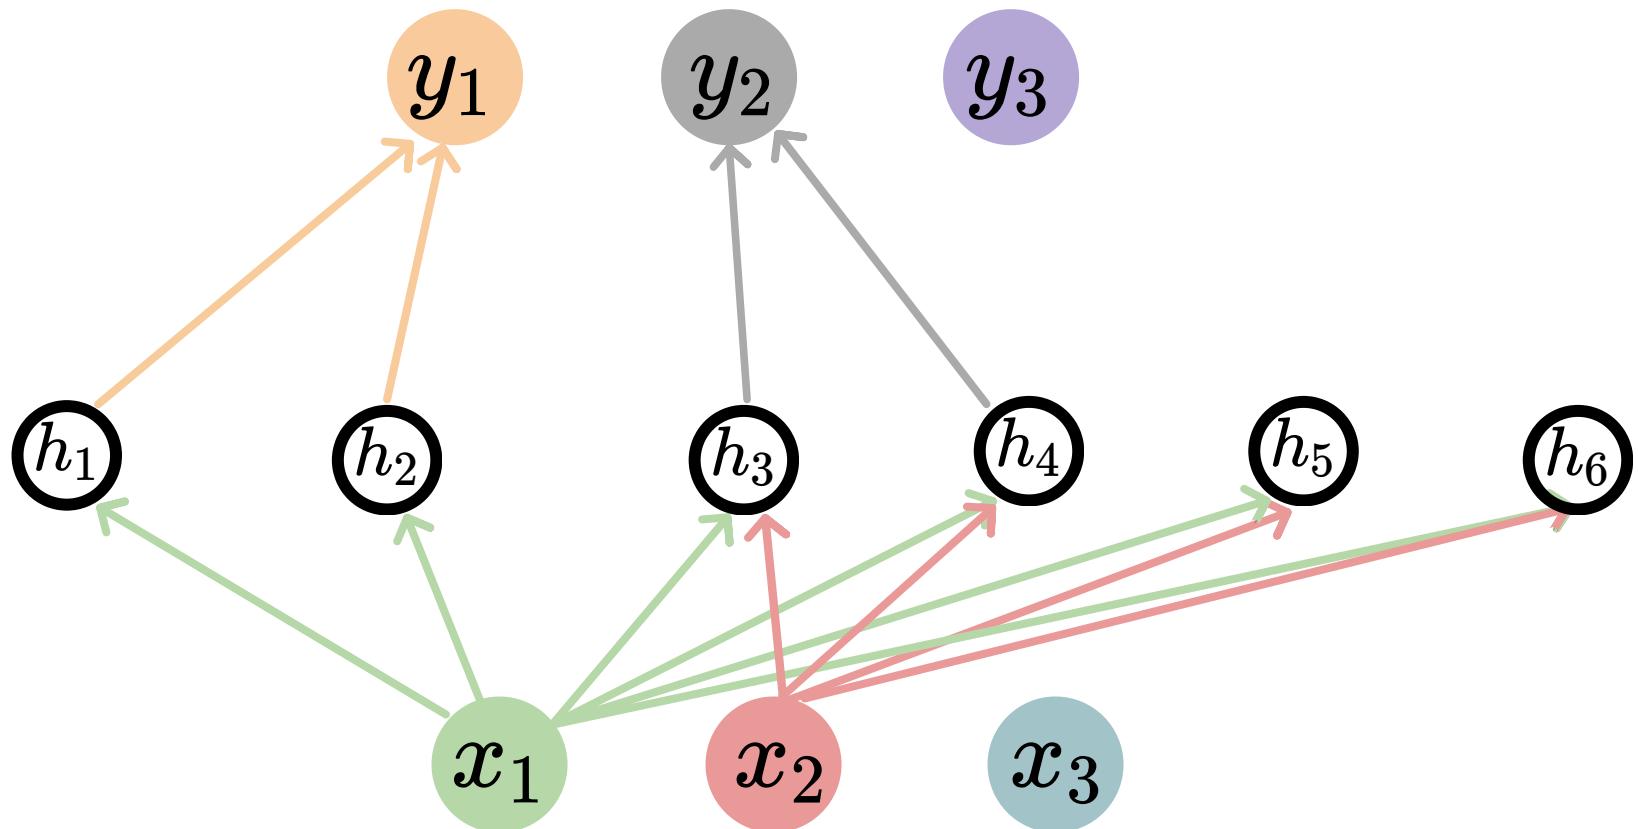
Masked MLP

$$y_i = \mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$



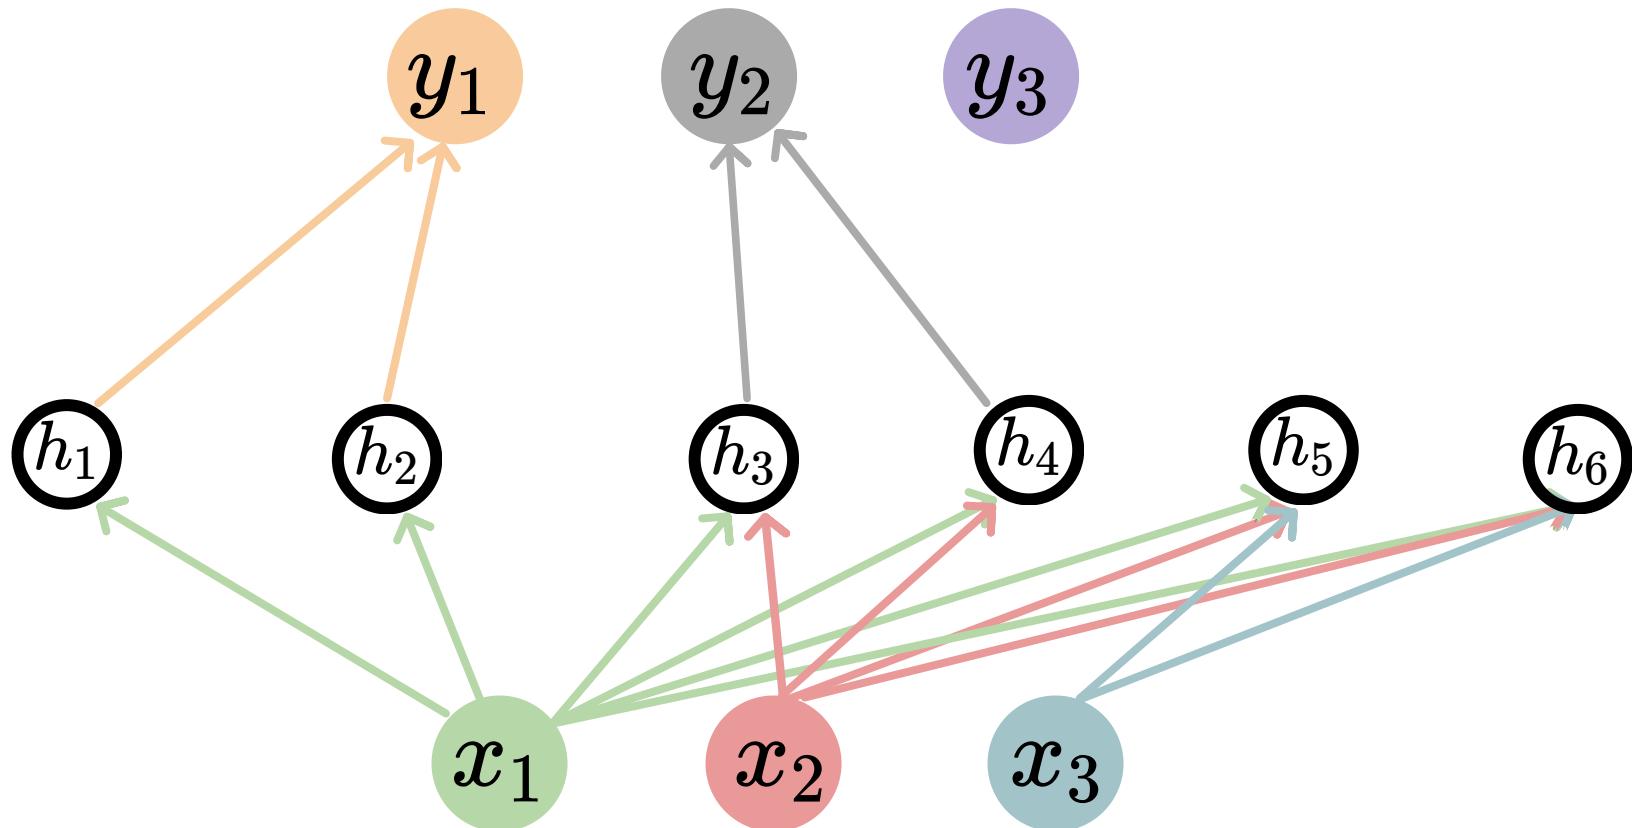
Masked MLP

$$y_i = \mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$



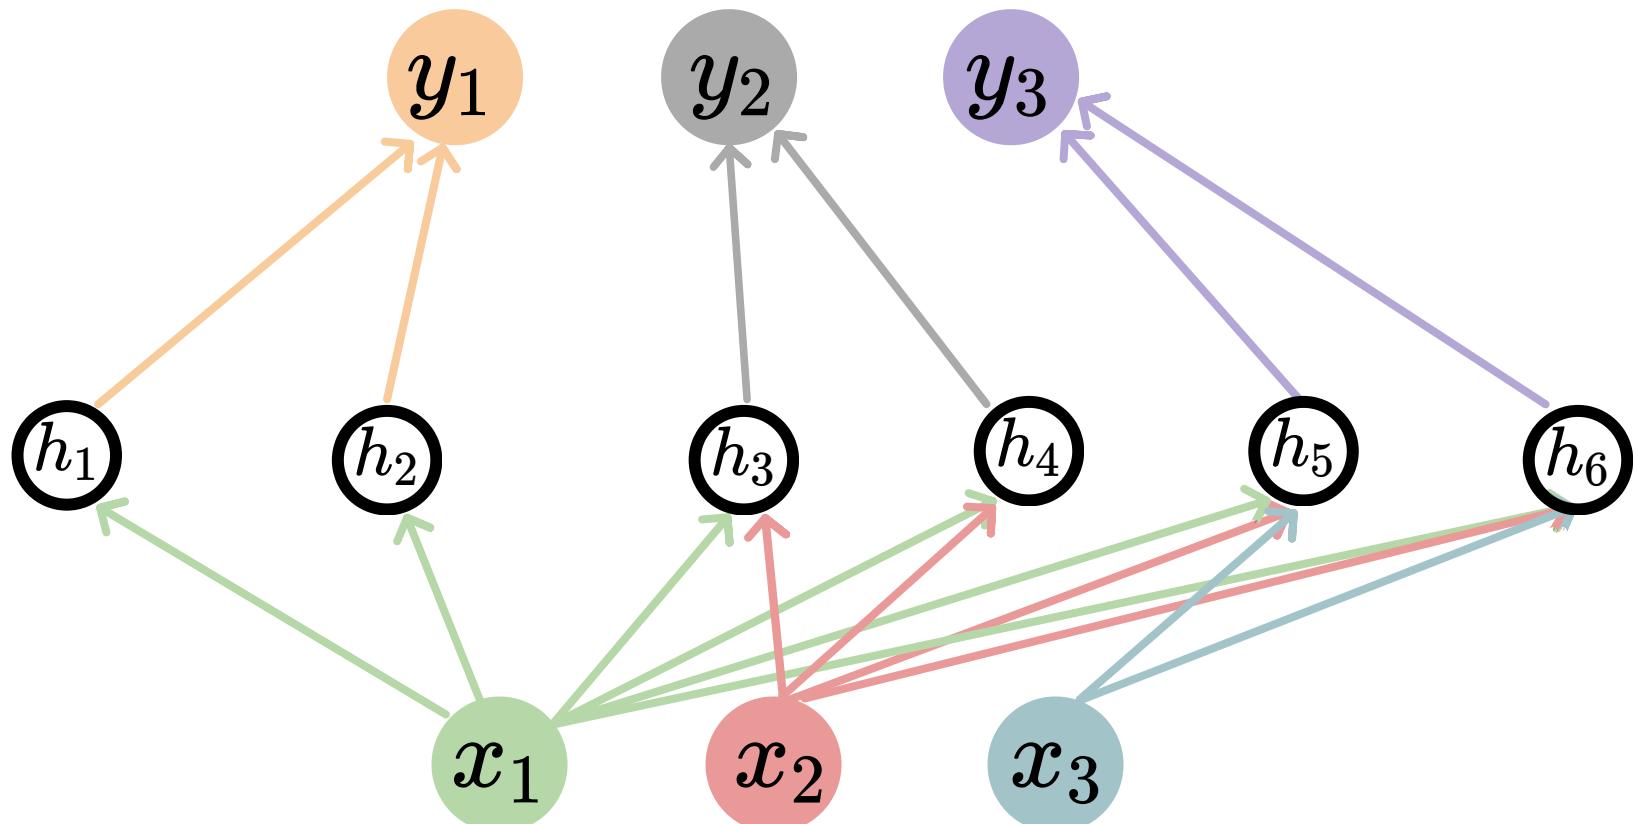
Masked MLP

$$y_i = \mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$



Masked MLP

$$y_i = \mu_i, \alpha_i = f_{\phi_i}(x_{1:i-1})$$



$$\log p(x) = \log p(z_0) - \sum_{k=1}^K \log |\det J_{f_k}(z_{k-1})|$$

Computational Complexity

$\mathcal{O}(N^3)$

$$\log p(x) = \log p(z_0) - \sum_{k=1}^K \log |\det J_{f_k}(z_{k-1})|$$

Computational Complexity

$\mathcal{O}(N^3)$

Autoregressive = Triangular matrix

$$\frac{\partial u_k^{(d)}}{\partial u_{k-1}^{(>d)}} == 0 \implies \begin{pmatrix} \frac{\partial u_k^{(1)}}{\partial u_{k-1}^{(1)}} & 0 & \dots & 0 \\ \frac{\partial u_k^{(2)}}{\partial u_{k-1}^{(1)}} & \frac{\partial u_k^{(2)}}{\partial u_{k-1}^{(2)}} & \dots & 0 \\ \vdots & \dots & \dots & \vdots \\ \frac{\partial u_k^{(D)}}{\partial u_{k-1}^{(1)}} & \dots & \dots & \frac{\partial u_k^{(D)}}{\partial u_{k-1}^{(D)}} \end{pmatrix}.$$

$$\log p(x) = \log p(z_0) - \sum_{k=1}^K \log |\det J_{f_k}(z_{k-1})|$$

Computational Complexity

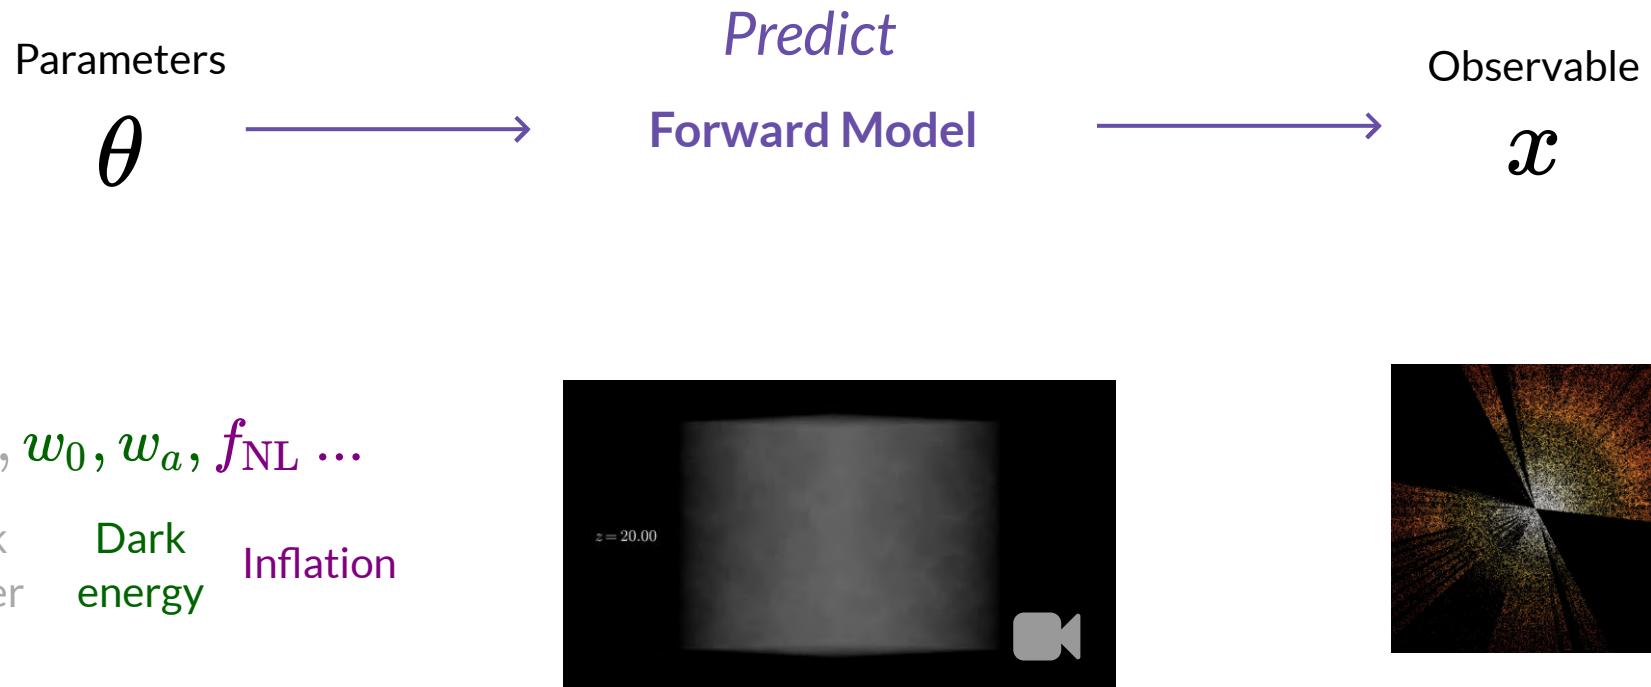
$$\mathcal{O}(N^3)$$

Autoregressive = Triangular matrix

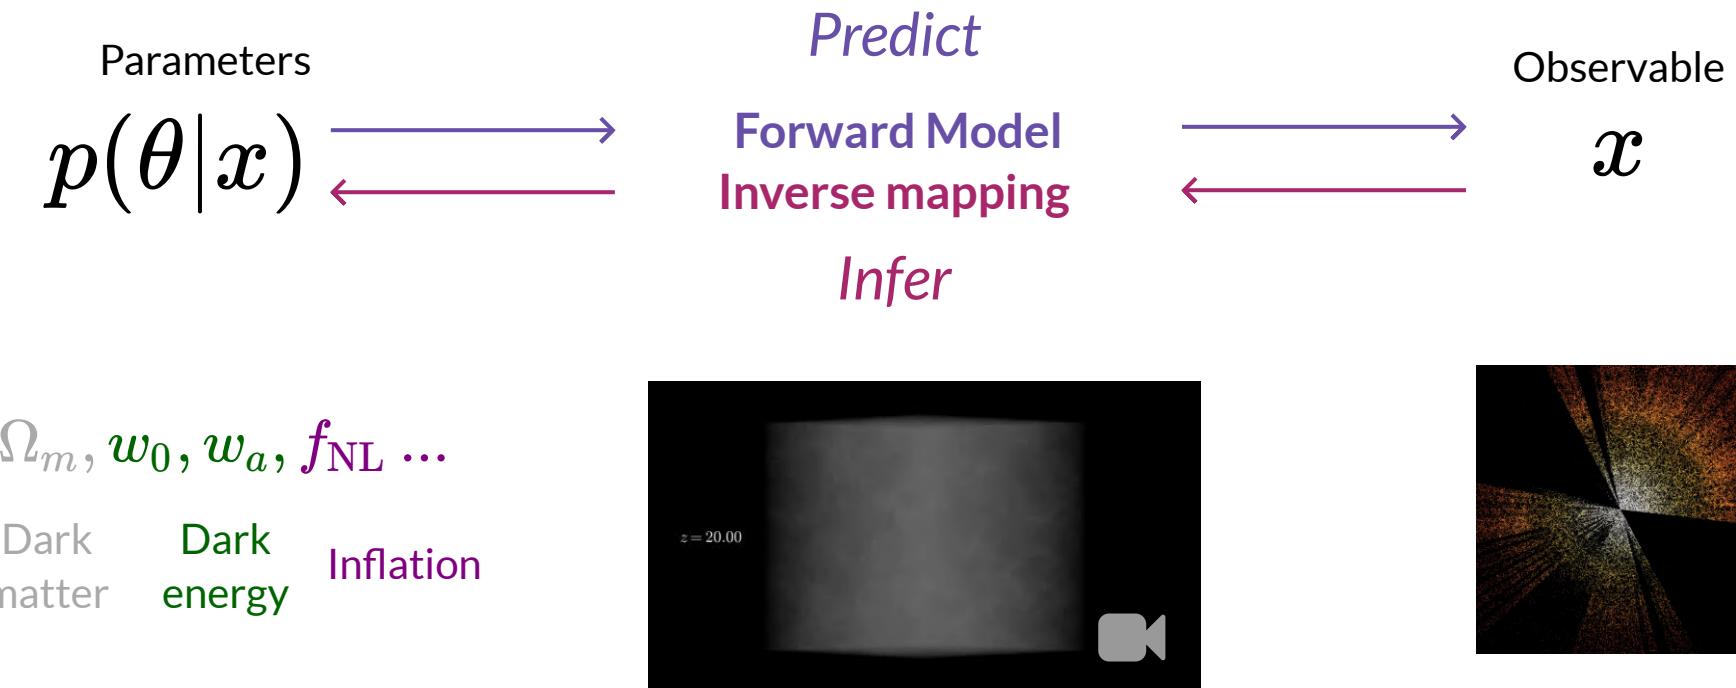
$$\mathcal{O}(N)$$

$$\frac{\partial u_k^{(d)}}{\partial u_{k-1}^{(>d)}} == 0 \implies \begin{pmatrix} \frac{\partial u_k^{(1)}}{\partial u_{k-1}^{(1)}} & 0 & \dots & 0 \\ \frac{\partial u_k^{(2)}}{\partial u_{k-1}^{(1)}} & \frac{\partial u_k^{(2)}}{\partial u_{k-1}^{(2)}} & \dots & 0 \\ \vdots & \dots & \dots & \vdots \\ \frac{\partial u_k^{(D)}}{\partial u_{k-1}^{(1)}} & \dots & \dots & \frac{\partial u_k^{(D)}}{\partial u_{k-1}^{(D)}} \end{pmatrix}.$$

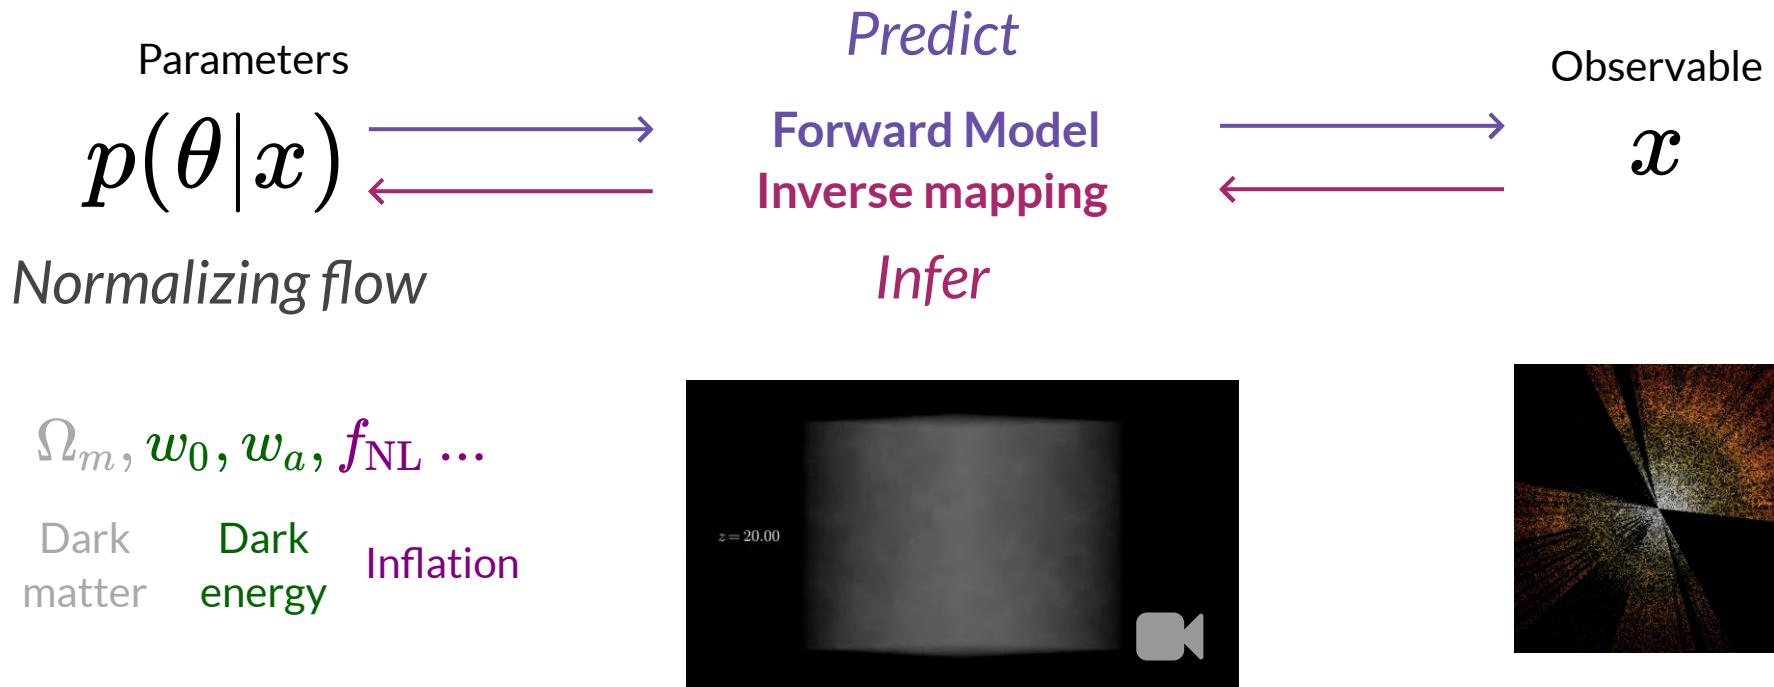
Simulation-based Inference



Simulation-based Inference



Simulation-based Inference



In continuous time

$$\frac{dx_t}{dt} = v_t^\phi(x_t)$$

[Image Credit: "Understanding Deep Learning" Simon J.D. Prince]

In continuous time

$$\frac{dx_t}{dt} = v_t^\phi(x_t)$$

$$x_1 = x_0 + \int_0^1 v_t^\phi(x_t) dt$$

$$x_0 = x_1 + \int_1^0 v_t^\phi(x_t) dt$$

[Image Credit: "Understanding Deep Learning" Simon J.D. Prince]

In continuous time

$$\frac{dx_t}{dt} = v_t^\phi(x_t)$$

$$x_1 = x_0 + \int_0^1 v_t^\phi(x_t) dt$$

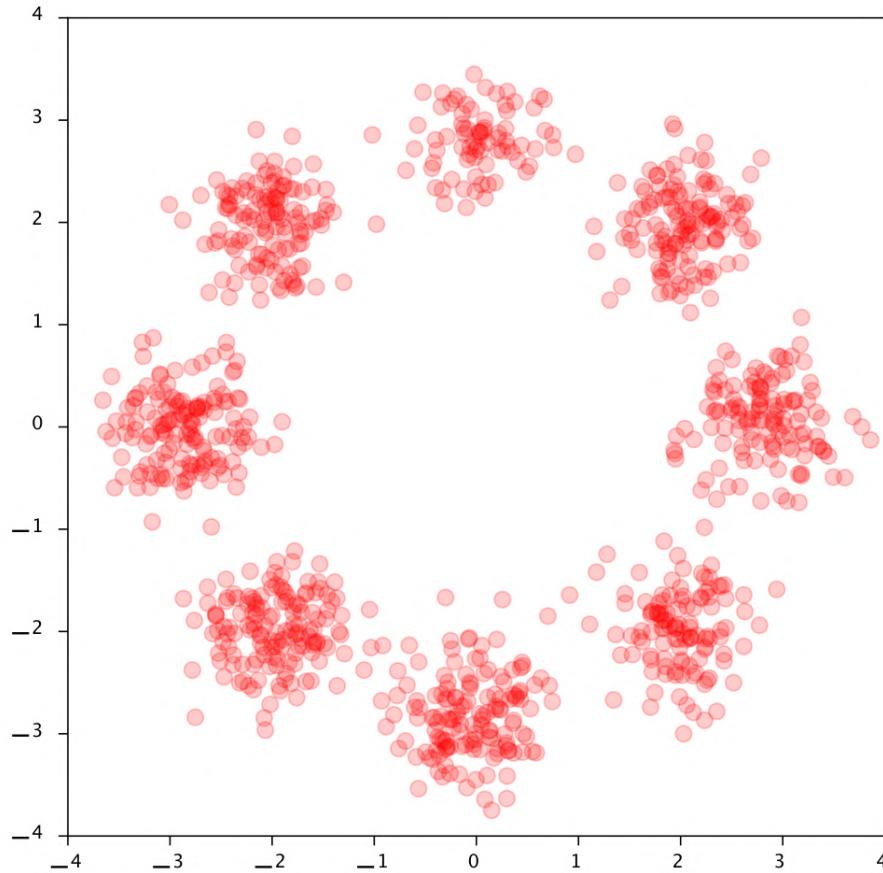
$$x_0 = x_1 + \int_1^0 v_t^\phi(x_t) dt$$

$$\frac{dp(x_t)}{dt} = -\nabla \left(v_t^\phi(x_t)p(x_t) \right)$$

Continuity Equation

[Image Credit: "Understanding Deep Learning" Simon J.D. Prince]

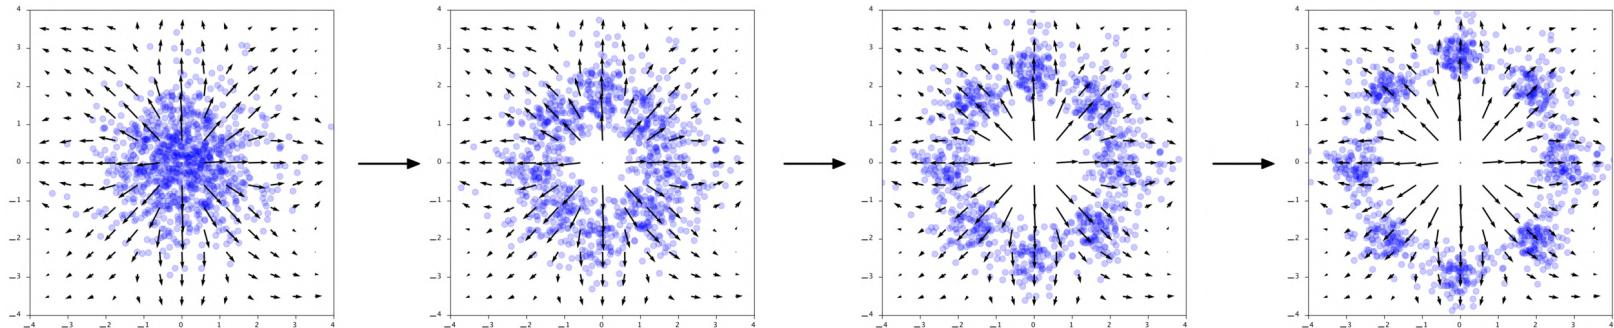
$$x_1 = x_0 + \int_0^1 v_\theta(x(t), t) dt$$



Chen et al. (2018), Grathwohl et al. (2018)

$$x_1 = x_0 + \int_0^1 v_\theta(x(t), t) dt$$

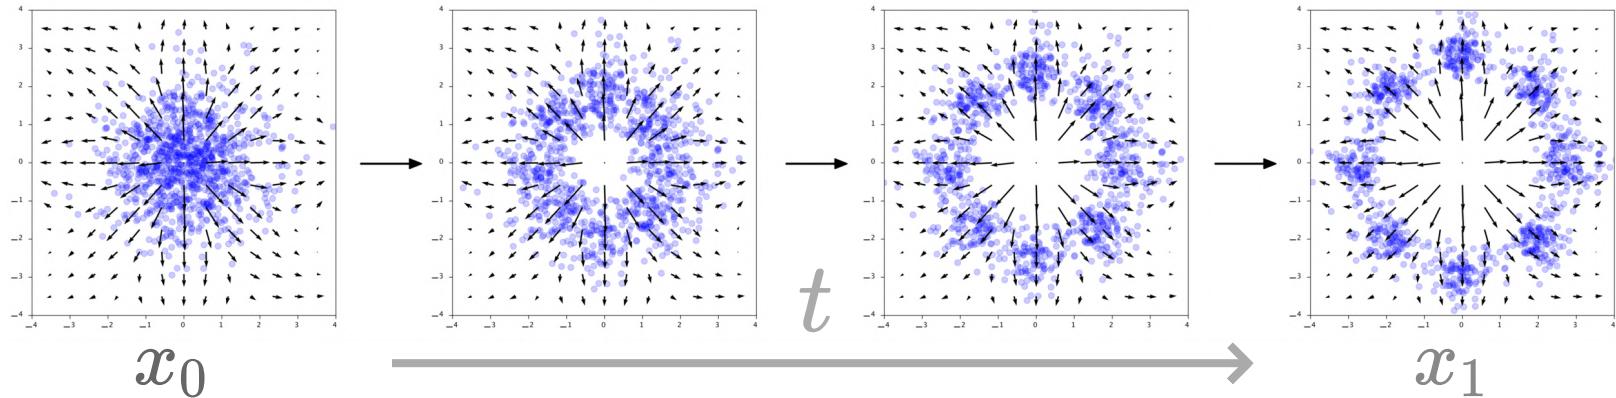
Generate



Chen et al. (2018), Grathwohl et al. (2018)

$$x_1 = x_0 + \int_0^1 v_\theta(x(t), t) dt$$

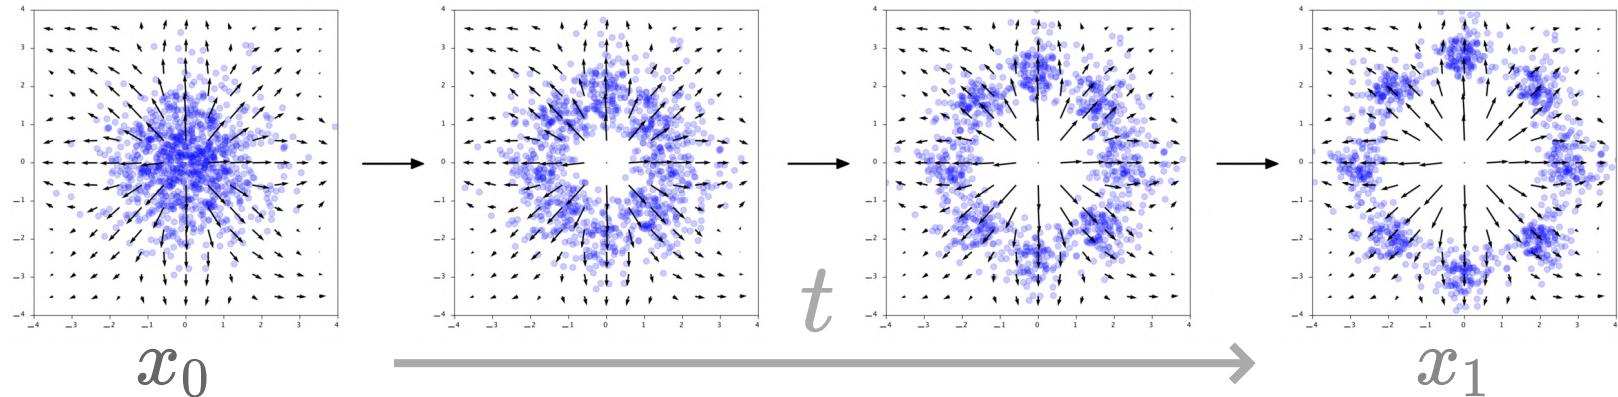
Generate



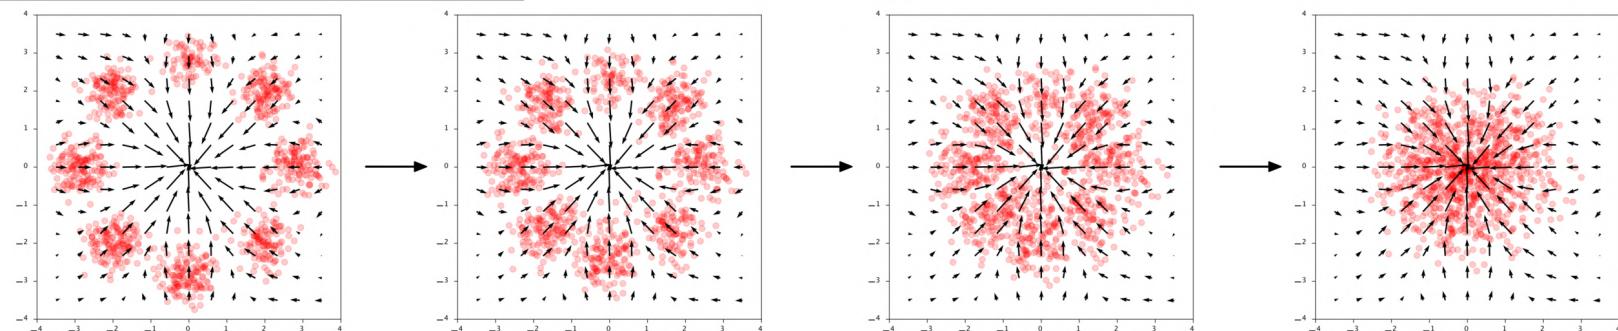
Chen et al. (2018), Grathwohl et al. (2018)

$$x_1 = x_0 + \int_0^1 v_\theta(x(t), t) dt$$

Generate

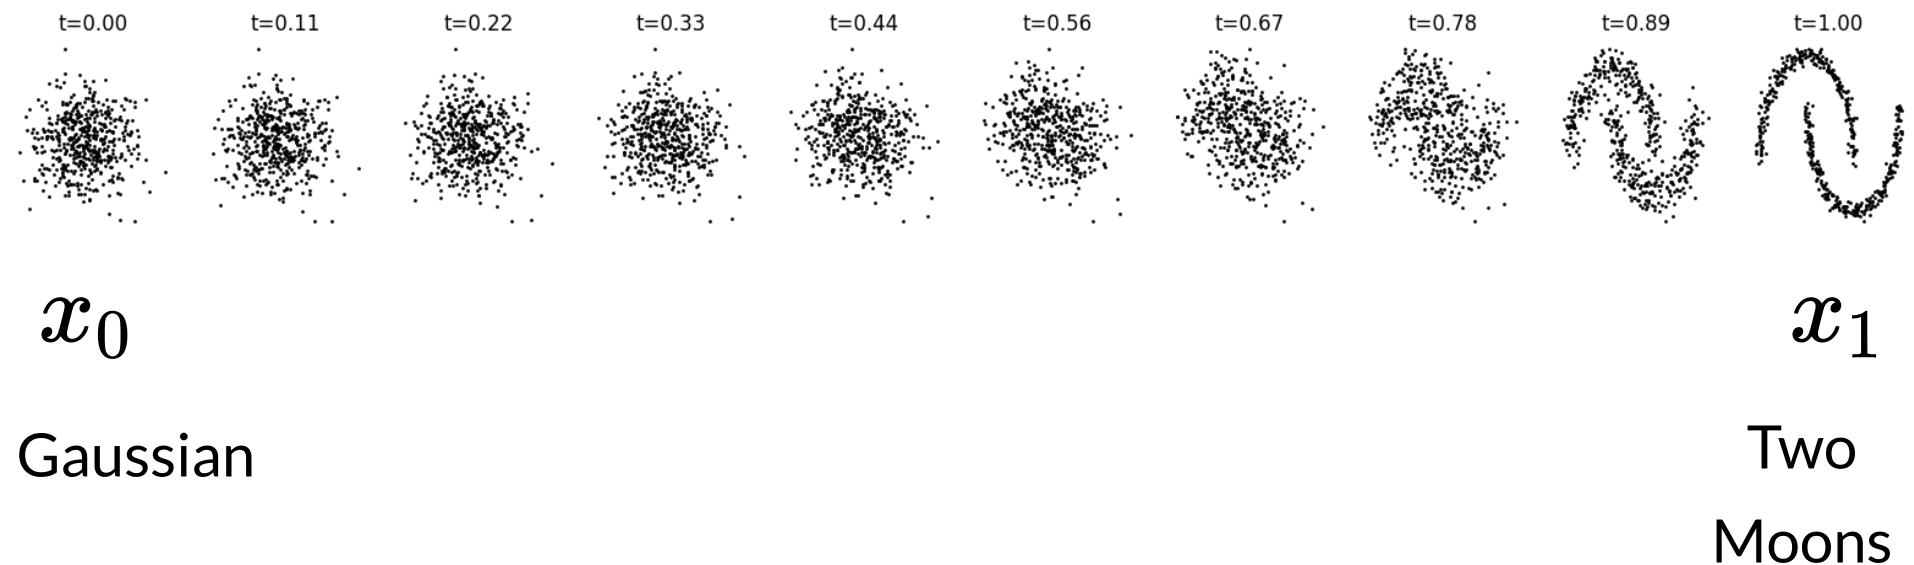


Evaluate Probability



Chen et al. (2018), Grathwohl et al. (2018)

Tutorial 2: Continuous time normalizing flows



$$\log p_X(x) = \log p_Z(z) + \int_0^1 \text{Tr} J_v(x(t)) dt$$

$$\log p_X(x) = \log p_Z(z) + \int_0^1 \text{Tr} J_v(x(t)) dt$$

Loss requires solving an ODE!

$$\log p_X(x) = \log p_Z(z) + \int_0^1 \text{Tr} J_v(x(t)) dt$$

Loss requires solving an ODE!

*Diffusion, Flow matching, Interpolants...
All ways to avoid this **at training time***

Conditional Flow matching

["Flow Matching for Generative Modeling" Lipman et al]

["Stochastic Interpolants: A Unifying framework for Flows and Diffusions" Albergo et al]

Conditional Flow matching

Assume a conditional vector field (known at training time)

$$x_t = (1 - t)x_0 + tx_1$$

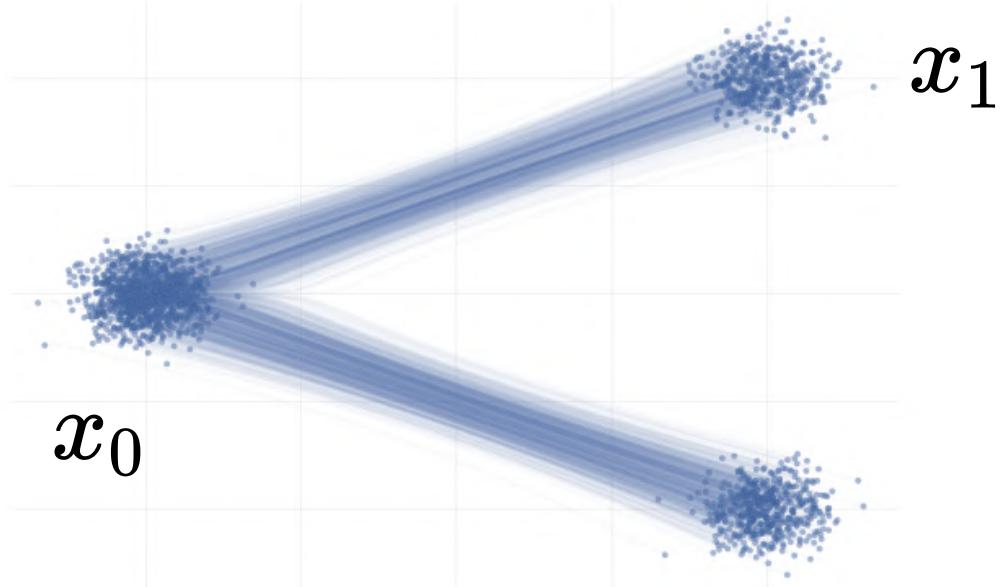
["Flow Matching for Generative Modeling" Lipman et al]

["Stochastic Interpolants: A Unifying framework for Flows and Diffusions" Albergo et al]

Conditional Flow matching

Assume a conditional vector field (known at training time)

$$x_t = (1 - t)x_0 + tx_1$$



["Flow Matching for Generative Modeling" Lipman et al]

["Stochastic Interpolants: A Unifying framework for Flows and Diffusions" Albergo et al]

Conditional Flow matching

Assume a conditional vector field (known at training time)

$$x_t = (1 - t)x_0 + tx_1$$

$$p_t(x) = \int p_t(x|x_1)q(x_1) dx_1$$
$$u_t^{\text{target}}(x) = \int u_t(x|x_1) \frac{p_t(x|x_1)p_1(x_1)}{p_t(x)} dx_1$$

["Flow Matching for Generative Modeling" Lipman et al]

["Stochastic Interpolants: A Unifying framework for Flows and Diffusions" Albergo et al]

Conditional Flow matching

Assume a conditional vector field (known at training time)

$$x_t = (1 - t)x_0 + tx_1$$

$$p_t(x) = \int p_t(x|x_1)q(x_1) dx_1$$

$$u_t^{\text{target}}(x) = \int u_t(x|x_1) \frac{p_t(x|x_1)p_1(x_1)}{p_t(x)} dx_1$$

$$\mathcal{L}_{\text{FM}} = \mathbb{E}_{t \sim U, x \sim p_t} \left[\|u_t^\phi(x) - u_t^{\text{target}}(x)\|^2 \right]$$

["Flow Matching for Generative Modeling" Lipman et al]

["Stochastic Interpolants: A Unifying framework for Flows and Diffusions" Albergo et al]

Conditional Flow matching

Assume a conditional vector field (known at training time)

$$x_t = (1 - t)x_0 + tx_1$$

Intractable

$$p_t(x) = \int p_t(x|x_1)q(x_1) dx_1$$

$$u_t^{\text{target}}(x) = \int u_t(x|x_1) \frac{p_t(x|x_1)p_1(x_1)}{p_t(x)} dx_1$$

$$\mathcal{L}_{\text{FM}} = \mathbb{E}_{t \sim U, x \sim p_t} \left[\|u_t^\phi(x) - u_t^{\text{target}}(x)\|^2 \right]$$

["Flow Matching for Generative Modeling" Lipman et al]

["Stochastic Interpolants: A Unifying framework for Flows and Diffusions" Albergo et al]

Conditional Flow matching

Assume a conditional vector field (known at training time)

$$x_t = (1 - t)x_0 + tx_1$$

Intractable

$$p_t(x) = \int p_t(x|x_1)q(x_1) dx_1$$

$$u_t^{\text{target}}(x) = \int u_t(x|x_1) \frac{p_t(x|x_1)p_1(x_1)}{p_t(x)} dx_1$$

$$\mathcal{L}_{\text{FM}} = \mathbb{E}_{t \sim U, x \sim p_t} \left[\|u_t^\phi(x) - u_t^{\text{target}}(x)\|^2 \right]$$

The loss that we can compute

$$\mathcal{L}_{\text{CFM}} = \mathbb{E}_{t \sim U, x_1 \sim p_1, x_t \sim p_t} \left[\|u_t^\phi(x_t) - u_t^{\text{target}}(x|x_1)\|^2 \right]$$

["Flow Matching for Generative Modeling" Lipman et al]

["Stochastic Interpolants: A Unifying framework for Flows and Diffusions" Albergo et al]

Conditional Flow matching

Assume a conditional vector field (known at training time)

$$x_t = (1 - t)x_0 + tx_1$$

$$p_t(x) = \int p_t(x|x_1)q(x_1) dx_1$$

$$u_t^{\text{target}}(x) = \int u_t(x|x_1) \frac{p_t(x|x_1)p_1(x_1)}{p_t(x)} dx_1$$

$$\mathcal{L}_{\text{FM}} = \mathbb{E}_{t \sim U, x \sim p_t} \left[\|u_t^\phi(x) - u_t^{\text{target}}(x)\|^2 \right]$$

The loss that we can compute

$$\mathcal{L}_{\text{CFM}} = \mathbb{E}_{t \sim U, x_1 \sim p_1, x_t \sim p_t} \left[\|u_t^\phi(x_t) - u_t^{\text{target}}(x|x_1)\|^2 \right]$$

The gradients of the losses are the same!

$$\nabla_\phi \mathcal{L}_{\text{CFM}} = \nabla_\phi \mathcal{L}_{\text{FM}}$$

["Flow Matching for Generative Modeling" Lipman et al]

["Stochastic Interpolants: A Unifying framework for Flows and Diffusions" Albergo et al]

Flow Matching

$$\frac{dz_t}{dt} = u_t^\phi(z_t)$$

[Image Credit: "Understanding Deep Learning" Simon J.D. Prince]

Flow Matching

$$\frac{dz_t}{dt} = u_t^\phi(z_t)$$

$$x = z_0 + \int_0^1 u_t^\phi(z_t) dt$$

[Image Credit: "Understanding Deep Learning" Simon J.D. Prince]

Flow Matching

$$\frac{dz_t}{dt} = u_t^\phi(z_t)$$

$$x = z_0 + \int_0^1 u_t^\phi(z_t) dt$$

$$\frac{dp(z_t)}{dt} = -\nabla \left(u_t^\phi(z_t) p(z_t) \right)$$

Continuity equation

[Image Credit: "Understanding Deep Learning" Simon J.D. Prince]

Flow Matching

$$\frac{dz_t}{dt} = u_t^\phi(z_t)$$

✓ *Sample*

$$x = z_0 + \int_0^1 u_t^\phi(z_t) dt$$

✓ *Evaluate probabilities*

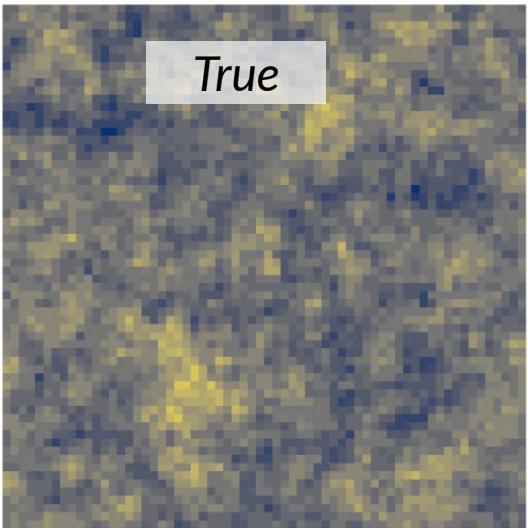
$$\frac{dp(z_t)}{dt} = -\nabla \left(u_t^\phi(z_t) p(z_t) \right)$$

Continuity equation

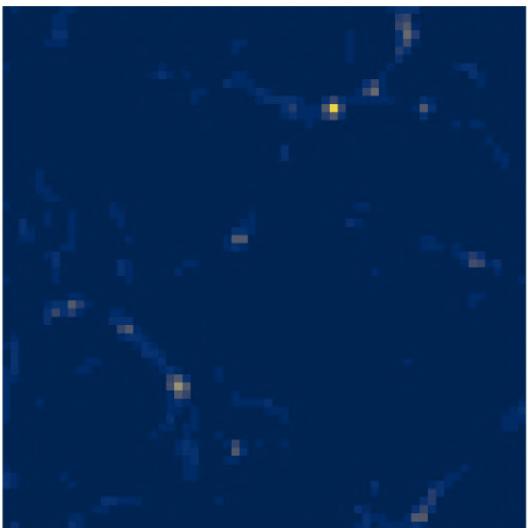
[Image Credit: "Understanding Deep Learning" Simon J.D. Prince]

Stochastic Interpolants

$$p(\delta_{\text{ICs}}, \theta | \delta_{\text{Obs}})$$



δ_{ICs}



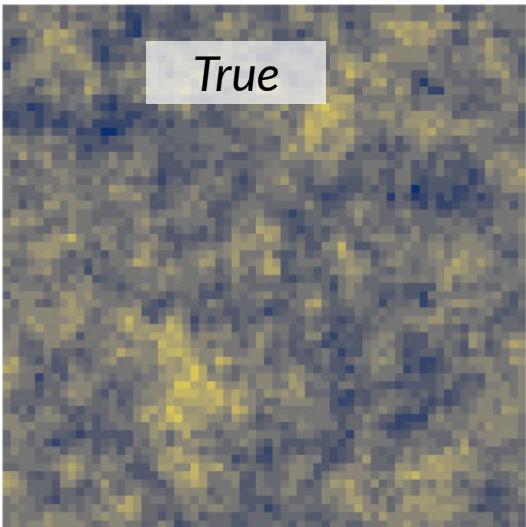
δ_{Obs}

"Joint cosmological
parameter inference and
initial condition
reconstruction with
Stochastic Interpolants"
Cuesta-Lazaro, Bayer,
Albergo et al
NeurIPS ML4PS 2024 Spotlight
talk

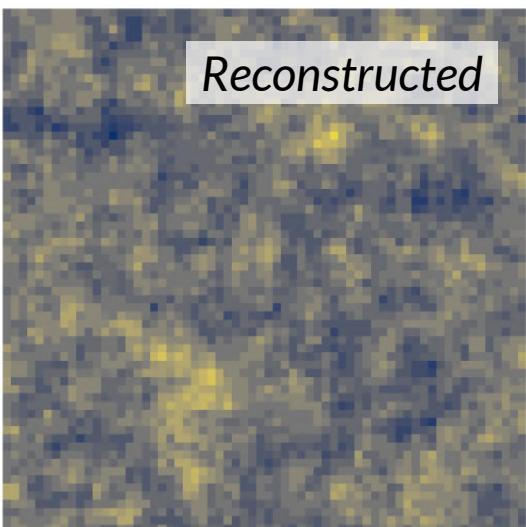
Stochastic Interpolants

$$p(\delta_{\text{ICs}}, \theta | \delta_{\text{Obs}})$$

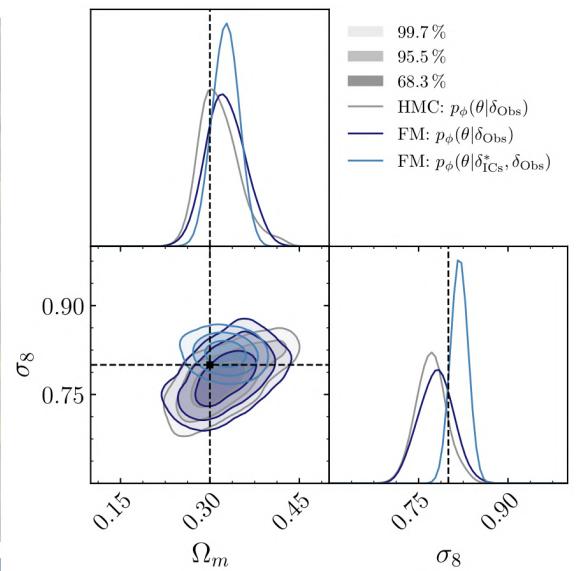
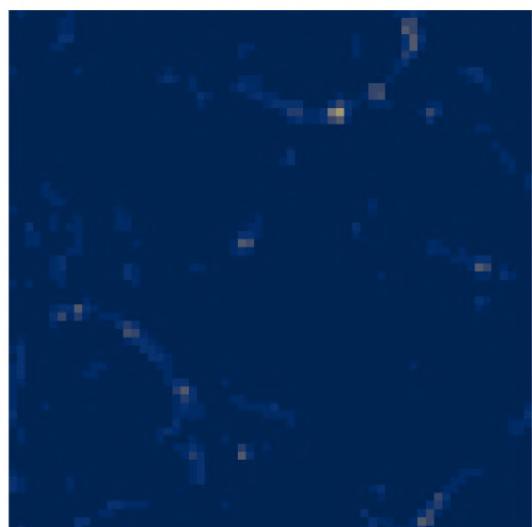
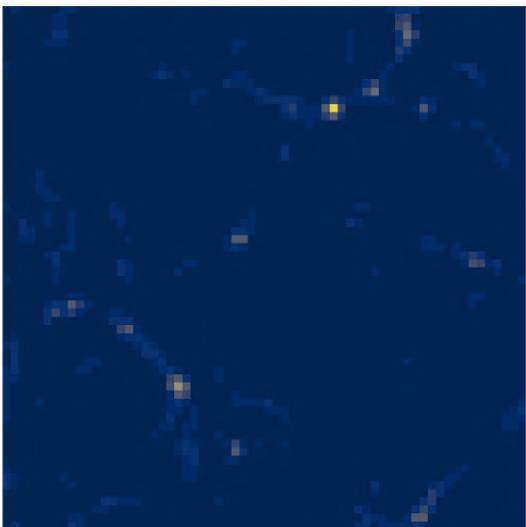
δ_{ICs}



Reconstructed



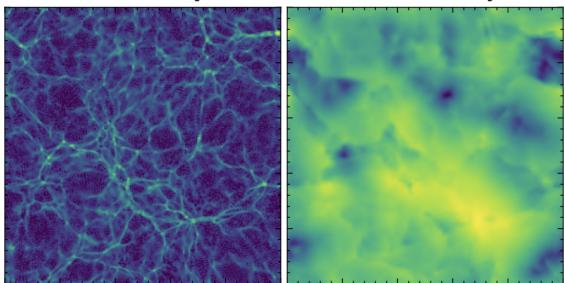
δ_{Obs}



"Joint cosmological parameter inference and initial condition reconstruction with Stochastic Interpolants"
Cuesta-Lazaro, Bayer, Albergo et al
NeurIPS ML4PS 2024 Spotlight talk

Hydro Simulations at scale

Particle Mesh for Gravity



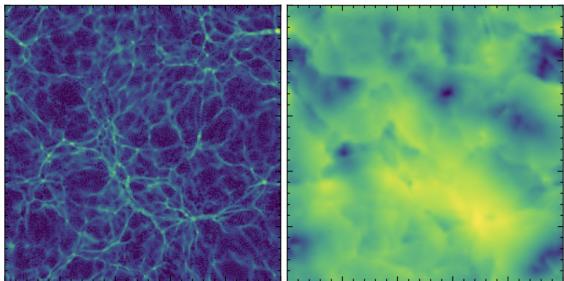
CAMELS Volumes

$25h^{-1}\text{Mpc}$

["BaryonBridge: Interpolants
models for fast
hydrodynamical simulations"
Horowitz, Cuesta-Lazaro,
Yehia ML4Astro workshop 2025]

Hydro Simulations at scale

Particle Mesh for Gravity



CAMELS Volumes
 $25h^{-1}\text{Mpc}$

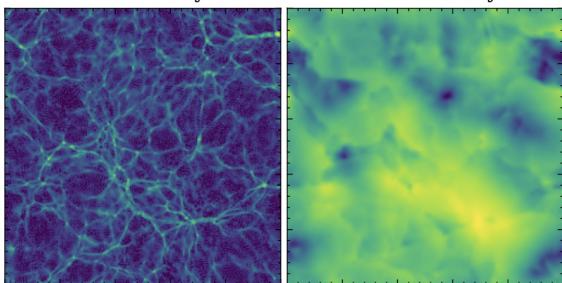
$+\mathcal{C}, \mathcal{A}$

1000 boxes with varying
cosmology and *feedback* models

["BaryonBridge: Interpolants
models for fast
hydrodynamical simulations"
Horowitz, Cuesta-Lazaro,
Yehia ML4Astro workshop 2025]

Hydro Simulations at scale

Particle Mesh for Gravity



$$p(\text{Baryons} | \text{DM}, \mathcal{C}, \mathcal{A})$$



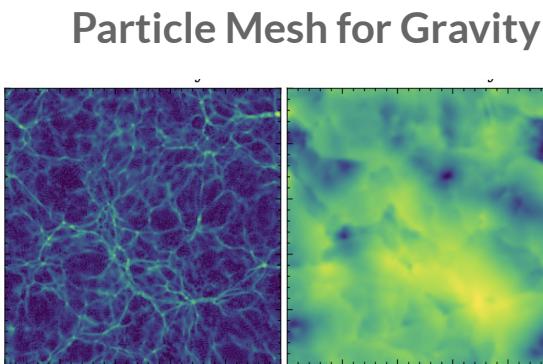
CAMELS Volumes
 $25h^{-1}\text{Mpc}$

$$+\mathcal{C}, \mathcal{A}$$

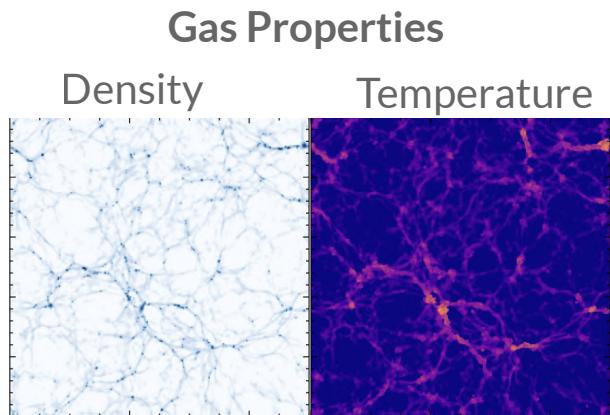
1000 boxes with varying
cosmology and *feedback* models

["BaryonBridge: Interpolants
models for fast
hydrodynamical simulations"
Horowitz, Cuesta-Lazaro,
Yehia ML4Astro workshop 2025]

Hydro Simulations at scale



$p(\text{Baryons} | \text{DM}, \mathcal{C}, \mathcal{A})$



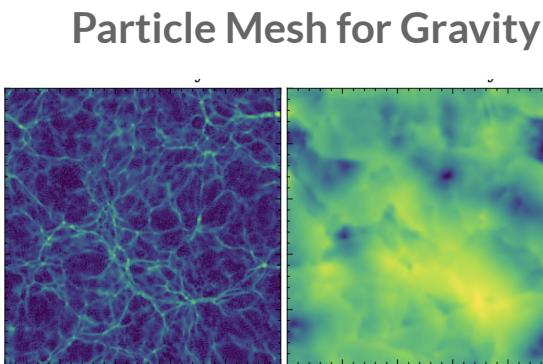
CAMELS Volumes
 $25h^{-1}\text{Mpc}$

1000 boxes with varying
cosmology and *feedback* models

$+\mathcal{C}, \mathcal{A}$

```
[ "BaryonBridge: Interpolants
  models for fast
  hydrodynamical simulations"
  Horowitz, Cuesta-Lazaro,
  Yehia ML4Astro workshop 2025]
```

Hydro Simulations at scale



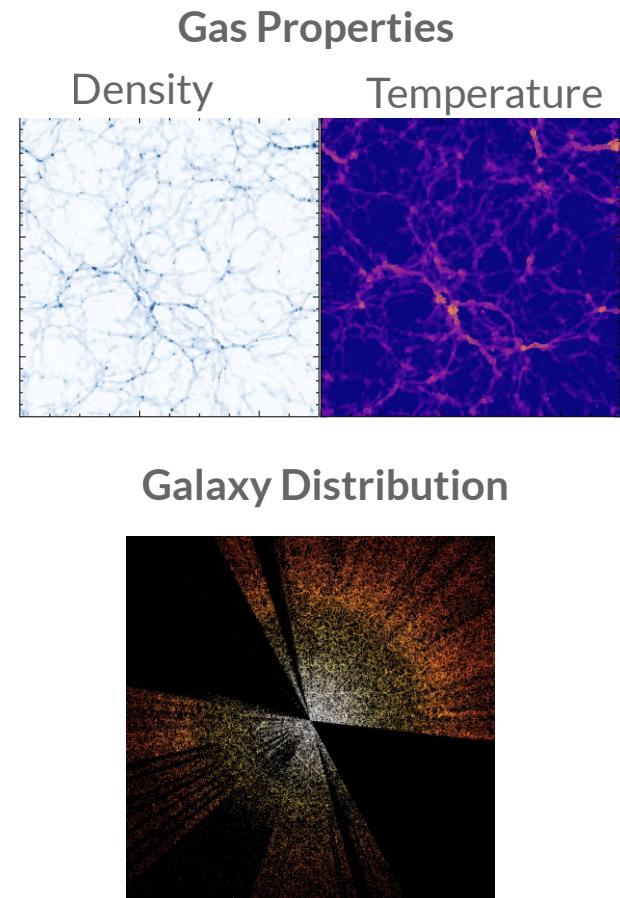
CAMELS Volumes
 $25h^{-1}\text{Mpc}$

1000 boxes with varying
cosmology and *feedback* models

$$p(\text{Baryons} | \text{DM}, \mathcal{C}, \mathcal{A})$$

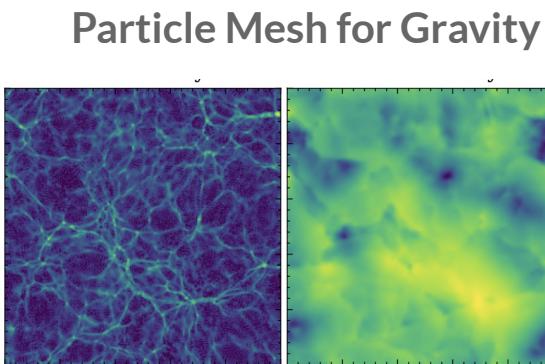


$$+\mathcal{C}, \mathcal{A}$$



```
[ "BaryonBridge: Interpolants  
models for fast  
hydrodynamical simulations"  
Horowitz, Cuesta-Lazaro,  
Yehia ML4Astro workshop 2025]
```

Hydro Simulations at scale



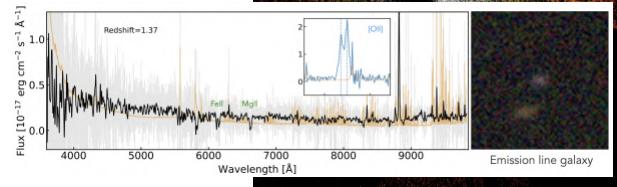
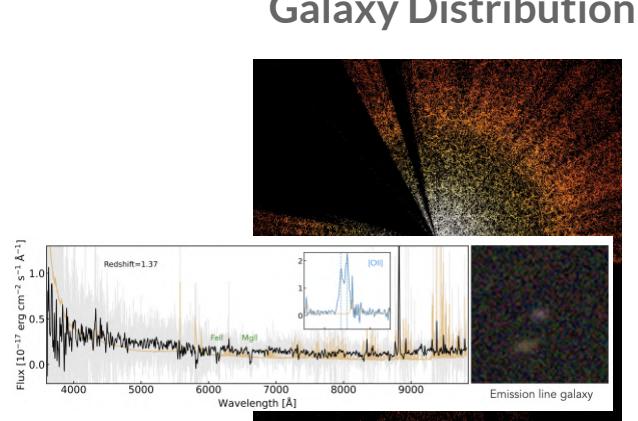
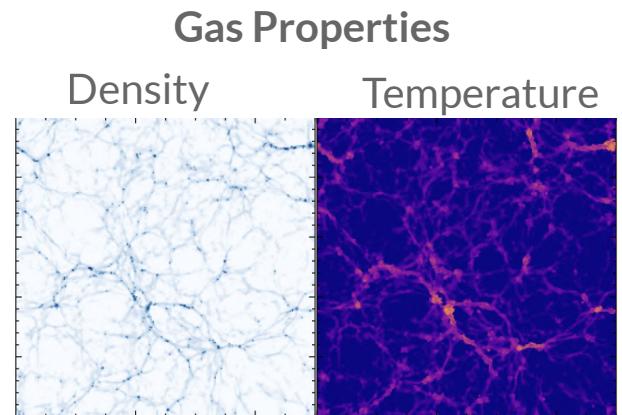
CAMELS Volumes
 $25h^{-1}\text{Mpc}$

1000 boxes with varying
cosmology and *feedback* models

$$p(\text{Baryons} | \text{DM}, \mathcal{C}, \mathcal{A})$$

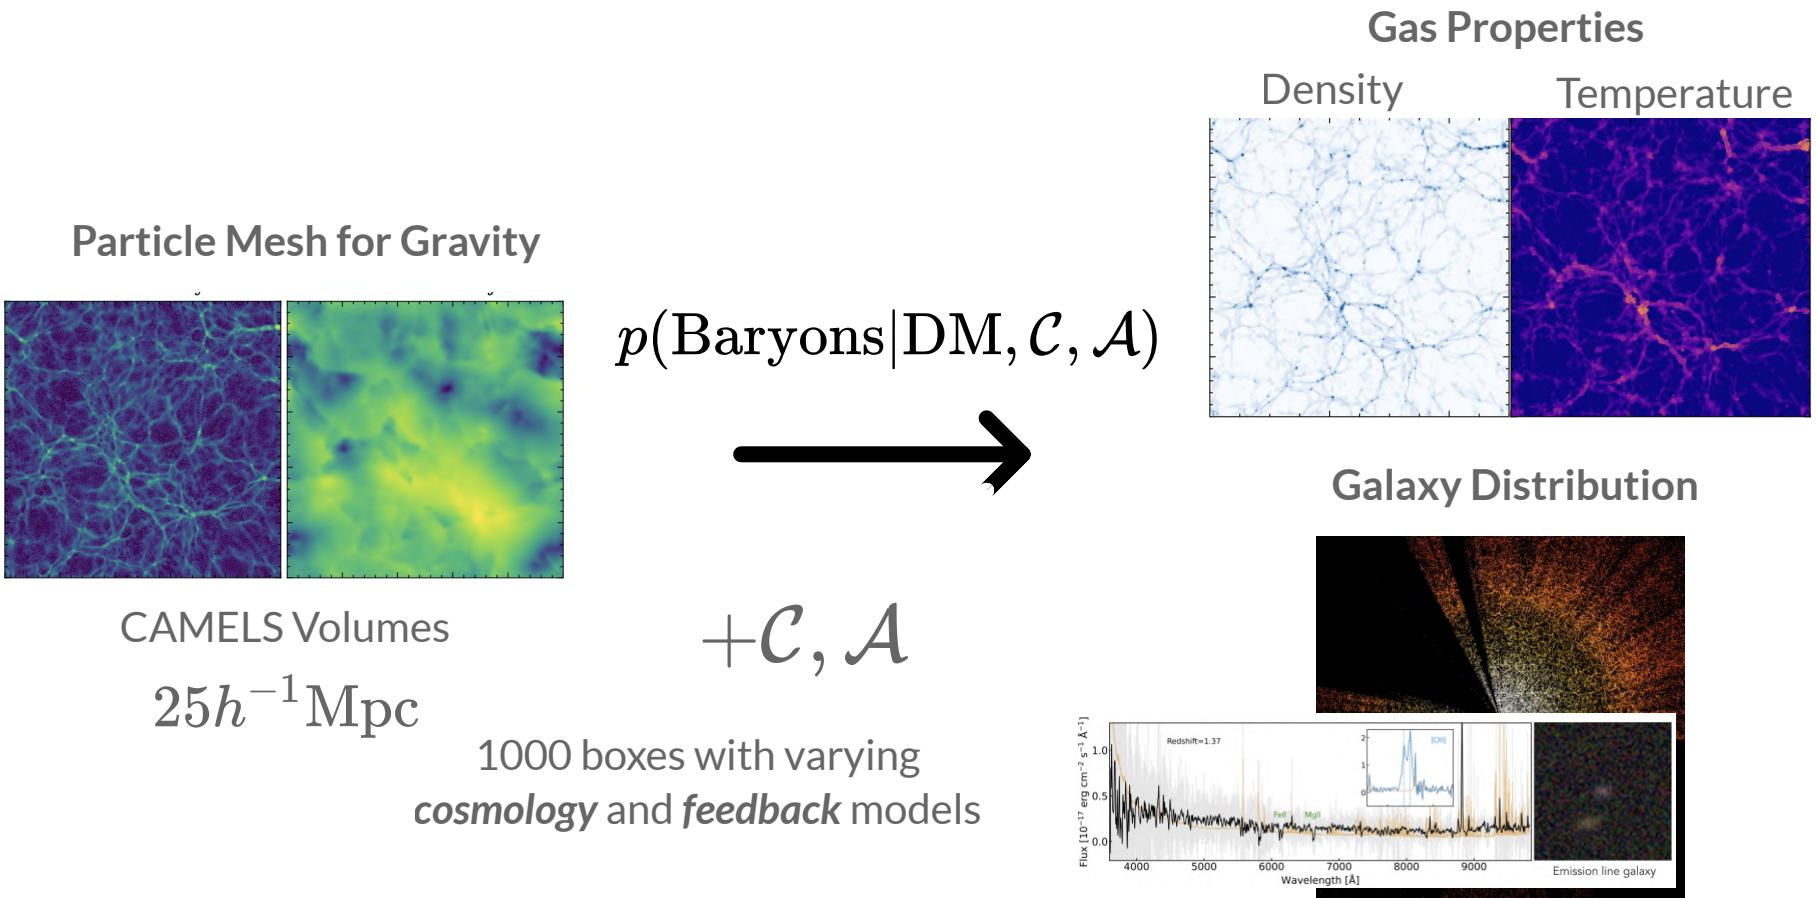


$$+\mathcal{C}, \mathcal{A}$$



["BaryonBridge: Interpolants
models for fast
hydrodynamical simulations"
Horowitz, Cuesta-Lazaro,
Yehia ML4Astro workshop 2025]

Hydro Simulations at scale



Current model optimised for Lyman Alpha forest
7 GPU minutes for a 50 Mpc simulation
130 million CPU core hours for TNG50

["BaryonBridge: Interpolants
models for fast
hydrodynamical simulations"
Horowitz, Cuesta-Lazaro,
Yehia ML4Astro workshop 2025]

Find straighter paths with OT

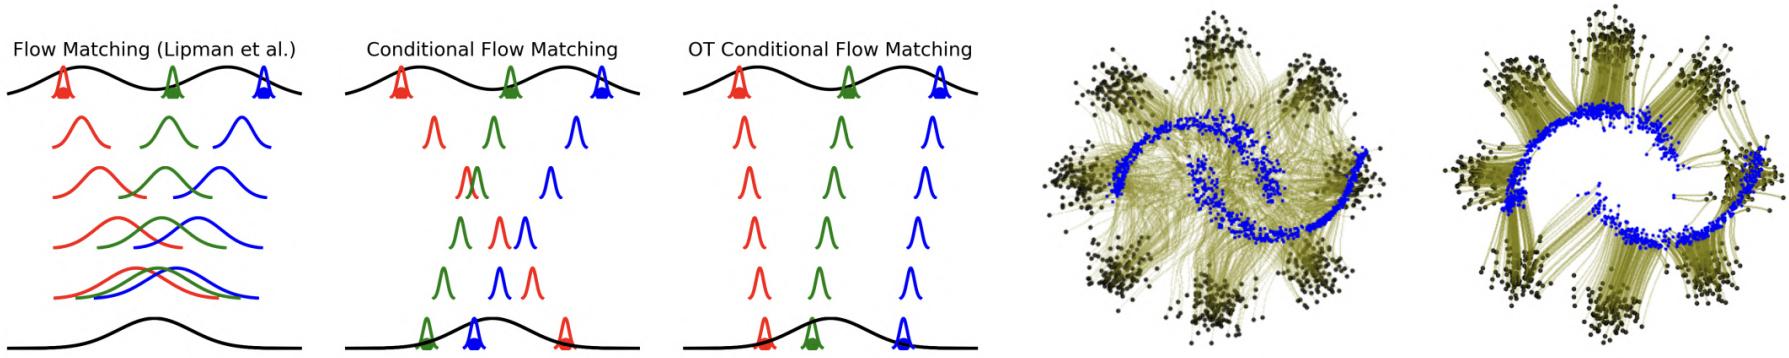
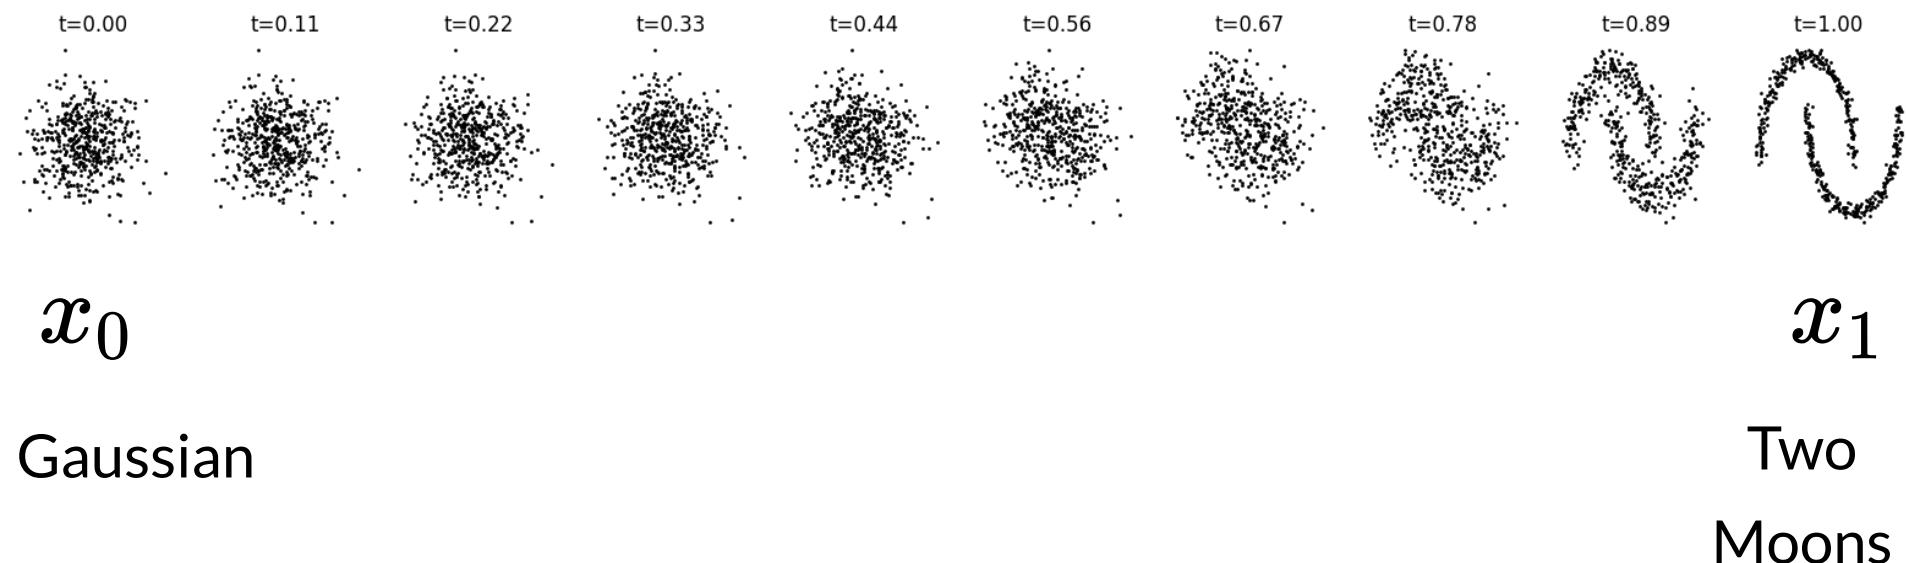


Figure 1: **Left:** Conditional flows from FM (Lipman et al., 2023), I-CFM (§3.2.2), and OT-CFM (§3.2.3). **Right:** Learned flows (green) from moons (blue) to 8gaussians (black) using I-CFM (centre-right) and OT-CFM (far right).

["Improving and generalizing flow-based generative models with minibatch optimal transport" Tong et al]

Tutorial 3: Flow Matching from scratch



Diffusion Models



Diffusion Models



←----- Forward diffusion: Add Gaussian noise (fixed)

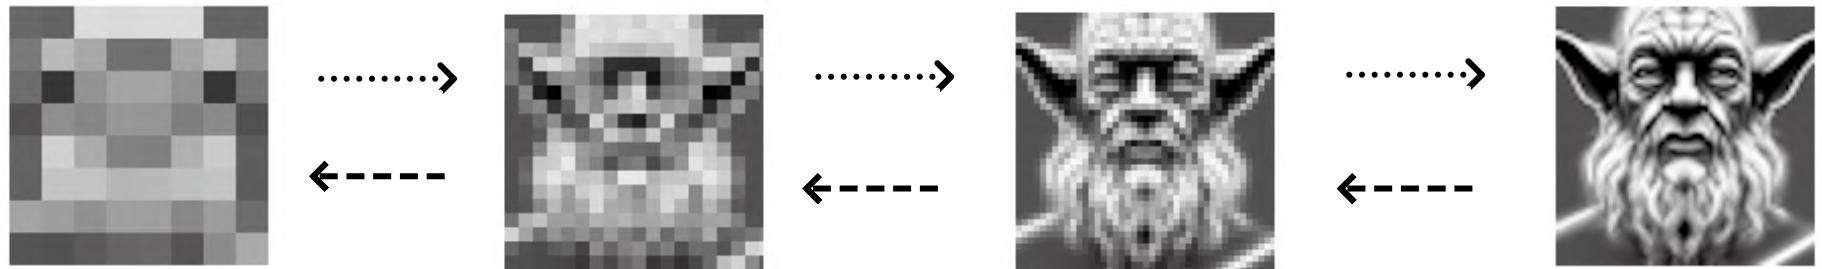
Diffusion Models



Fixed base
distribution:
Gaussian

Forward diffusion: Add Gaussian noise (fixed)

Diffusion Models



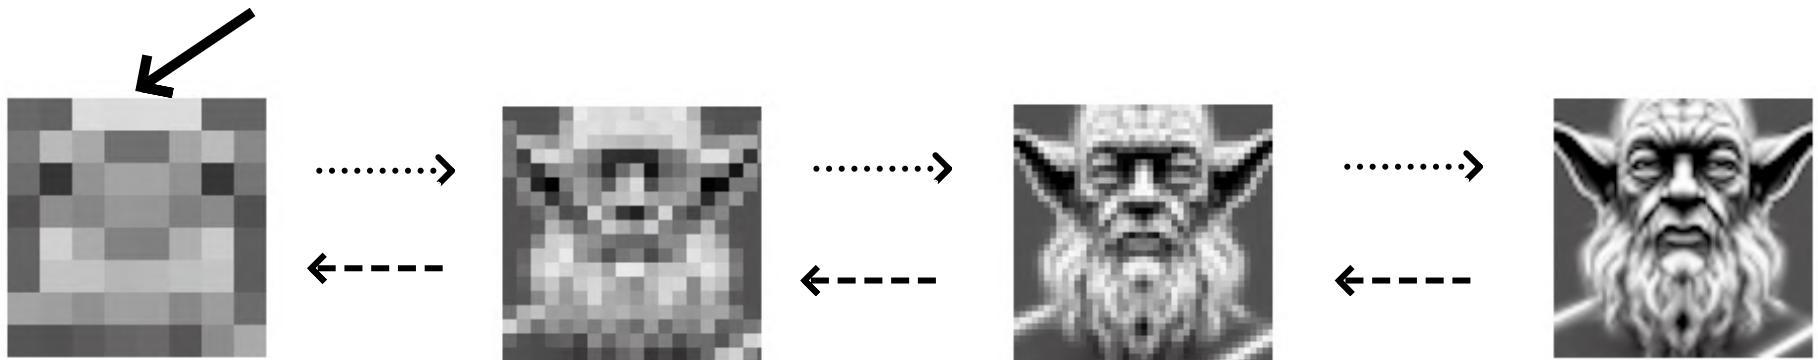
Fixed base
distribution:
Gaussian

- ←---- Forward diffusion: Add Gaussian noise (fixed)
-→ Reverse diffusion: Denoise previous step

Diffusion Models

Prompt

A person half Yoda half Gandalf



Fixed base
distribution:

Gaussian

- ←---- Forward diffusion: Add Gaussian noise (fixed)
-→ Reverse diffusion: Denoise previous step

Diffusion Models

Prompt

A person half Yoda half Gandalf



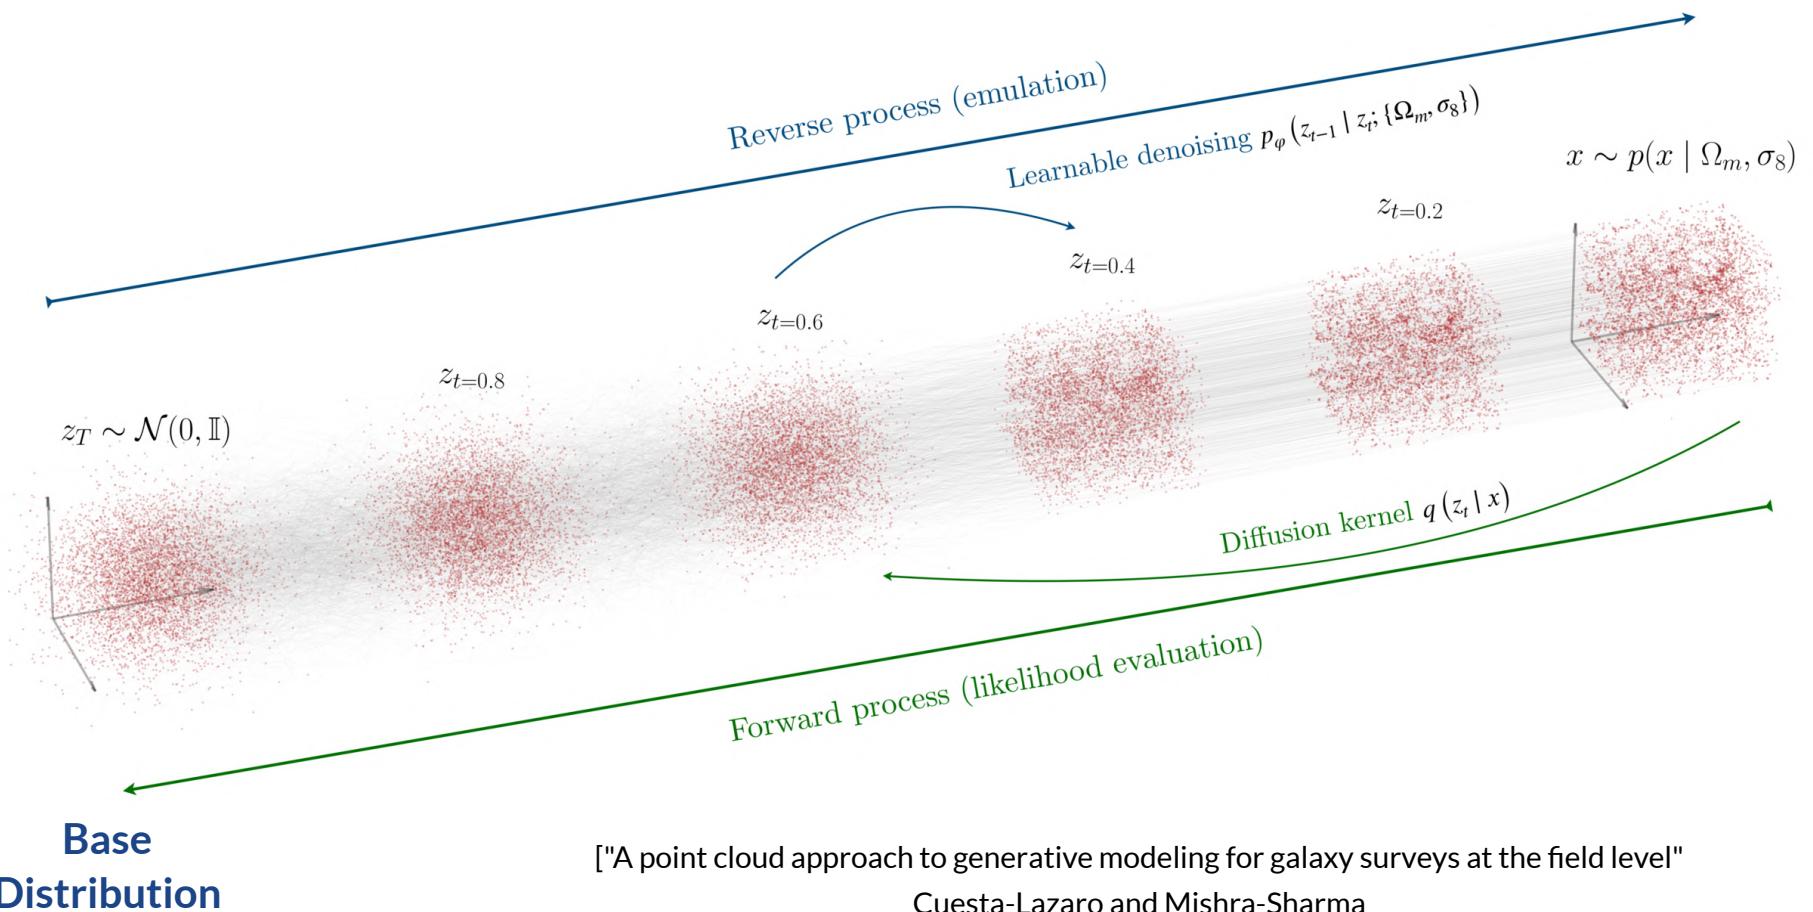
Fixed base
distribution:
Gaussian

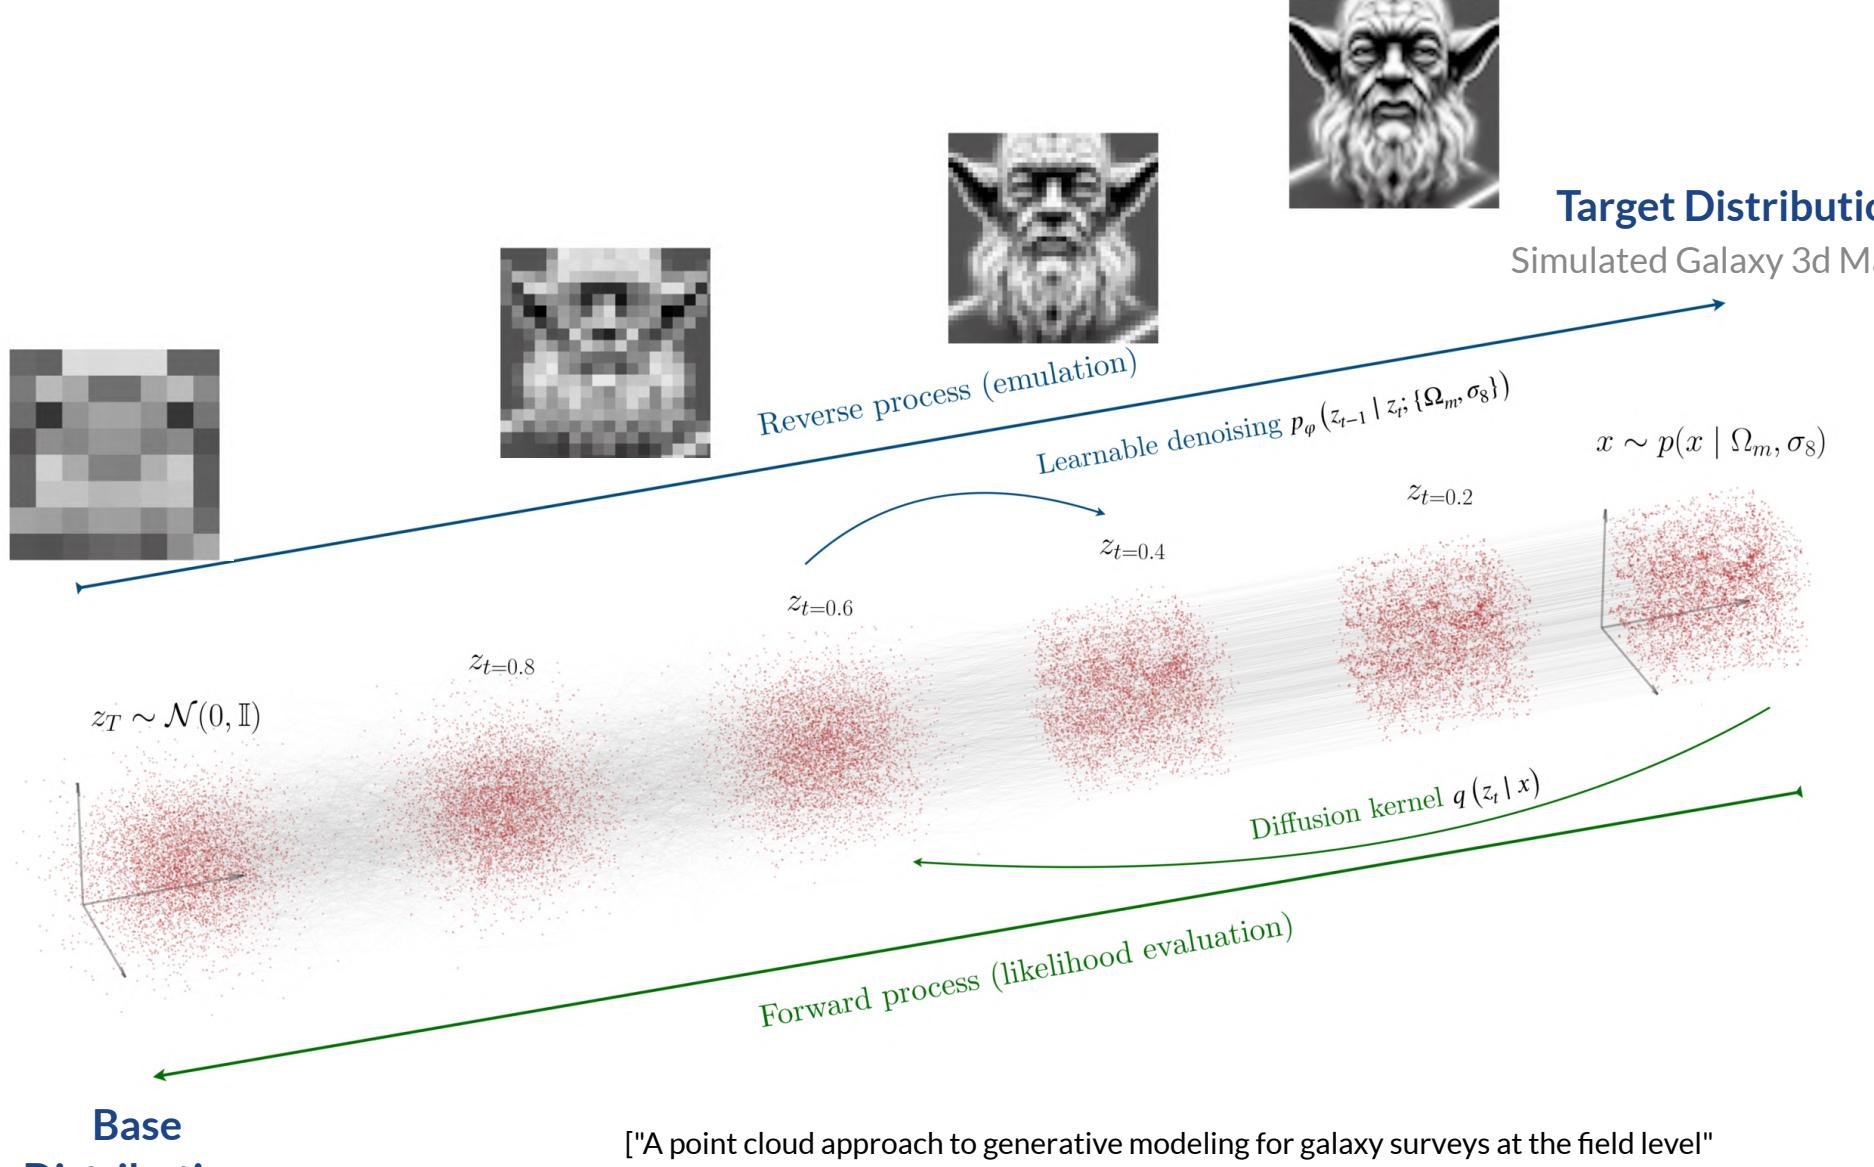
Denoising = Regression

- > Forward diffusion: Add Gaussian noise (fixed)
-> Reverse diffusion: Denoise previous step

Target Distribution

Simulated Galaxy 3d M



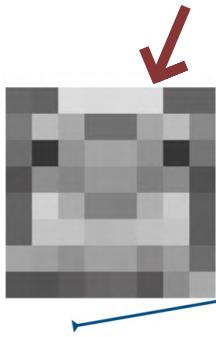


["A point cloud approach to generative modeling for galaxy surveys at the field level"

Cuesta-Lazaro and Mishra-Sharma

International Conference on Machine Learning ICML AI4Astro 2023, Spotlight talk, arXiv:2311.17141]

Prompt: A person half
Yoda half Gandalf



Target Distribution
Simulated Galaxy 3d Ma

$z_T \sim \mathcal{N}(0, \mathbb{I})$

$z_{t=0.8}$

Reverse process (emulation)

Learnable denoising $p_\varphi(z_{t-1} | z_t; \{\Omega_m, \sigma_8\})$

$z_{t=0.4}$

$z_{t=0.2}$

$x \sim p(x | \Omega_m, \sigma_8)$

$z_{t=0.6}$

Diffusion kernel $q(z_t | x)$

Forward process (likelihood evaluation)

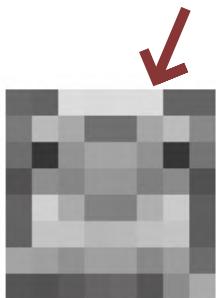
**Base
Distribution**

["A point cloud approach to generative modeling for galaxy surveys at the field level"

Cuesta-Lazaro and Mishra-Sharma

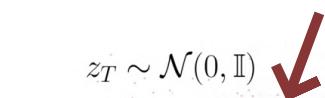
International Conference on Machine Learning ICML AI4Astro 2023, Spotlight talk, arXiv:2311.17141]

Prompt: A person half
Yoda half Gandalf



Prompt: Ω_m, σ_8

$z_T \sim \mathcal{N}(0, \mathbb{I})$



Base
Distribution



$z_{t=0.8}$

Reverse process (emulation)

$z_{t=0.6}$



$z_{t=0.4}$

Learnable denoising $p_\varphi(z_{t-1} | z_t; \{\Omega_m, \sigma_8\})$

$z_{t=0.2}$



Target Distribution

Simulated Galaxy 3d Ma

$x \sim p(x | \Omega_m, \sigma_8)$

Diffusion kernel $q(z_t | x)$

Forward process (likelihood evaluation)

["A point cloud approach to generative modeling for galaxy surveys at the field level"

Cuesta-Lazaro and Mishra-Sharma

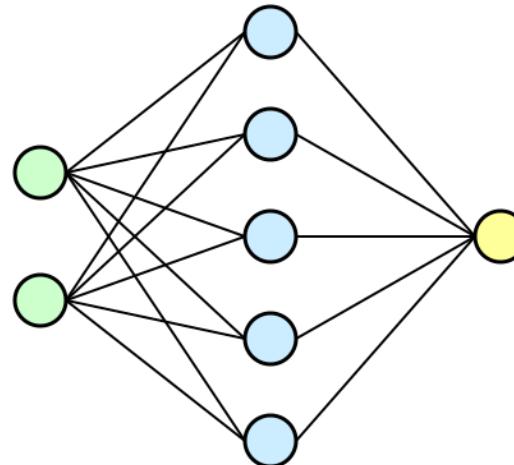
International Conference on Machine Learning ICML AI4Astro 2023, Spotlight talk, arXiv:2311.17141]

How good is my generative model?

$R \sim p(\theta|x)$



$F \sim \hat{p}(\theta|x)$



Real or Fake?

["A Practical Guide to Sample-based Statistical Distances for Evaluating Generative Models in Science"

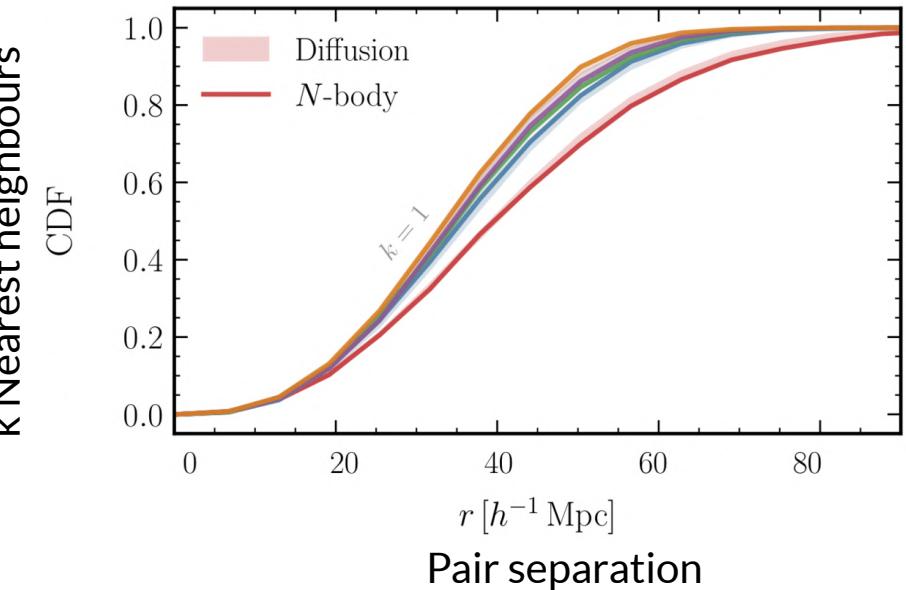
Bischoff et al 2024

arXiv:2403.12636]

Reproducing summary statistics

Varying cosmological parameters

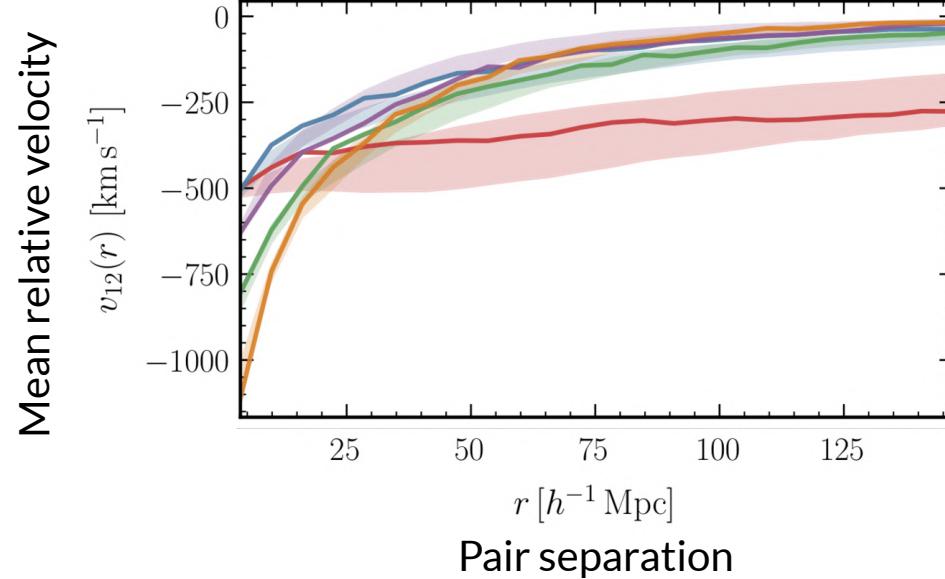
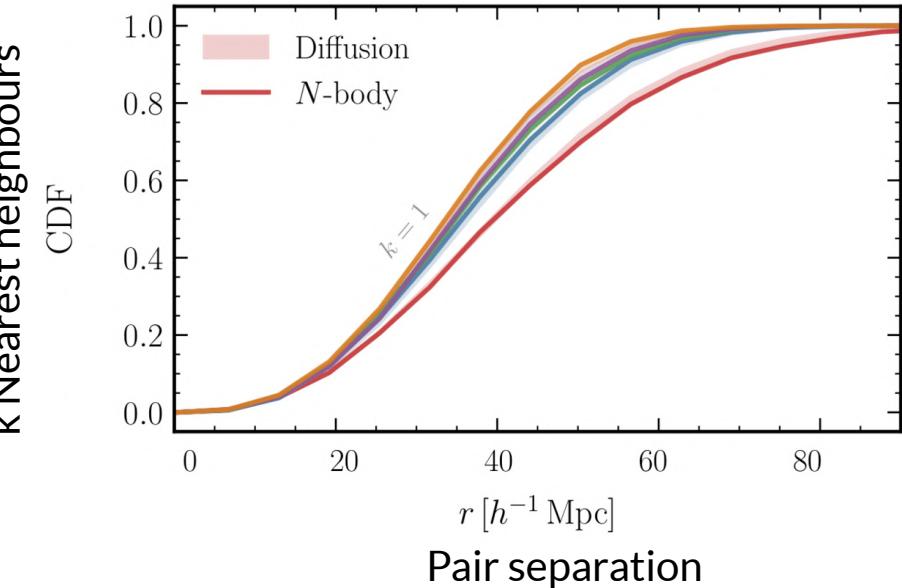
— $\Omega_m = 0.10 \ \sigma_8 = 0.92$ — $\Omega_m = 0.18 \ \sigma_8 = 0.68$ — $\Omega_m = 0.27 \ \sigma_8 = 0.92$
— $\Omega_m = 0.35 \ \sigma_8 = 0.62$ — $\Omega_m = 0.44 \ \sigma_8 = 0.95$



Reproducing summary statistics

Varying cosmological parameters

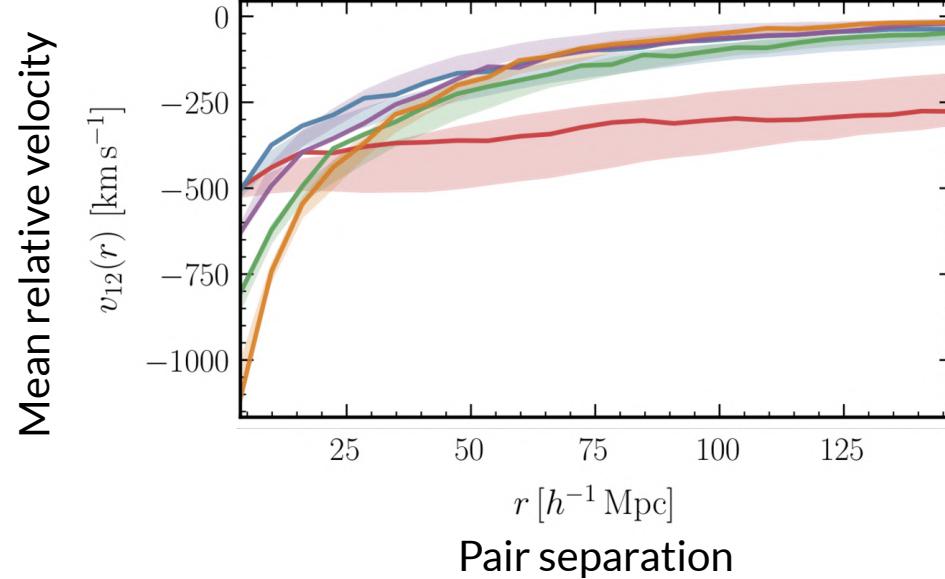
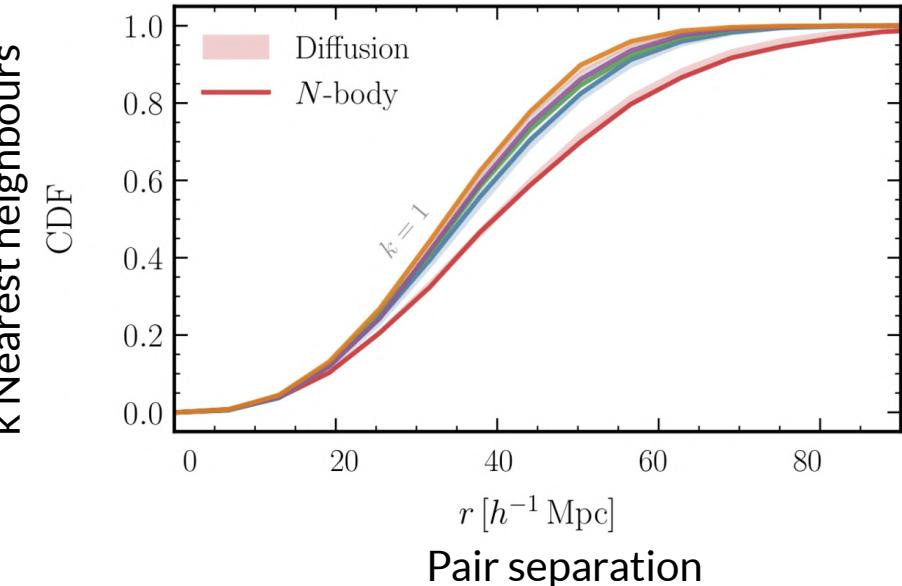
— $\Omega_m = 0.10 \quad \sigma_8 = 0.92$ — $\Omega_m = 0.18 \quad \sigma_8 = 0.68$ — $\Omega_m = 0.27 \quad \sigma_8 = 0.92$
— $\Omega_m = 0.35 \quad \sigma_8 = 0.62$ — $\Omega_m = 0.44 \quad \sigma_8 = 0.95$



Reproducing summary statistics

Varying cosmological parameters

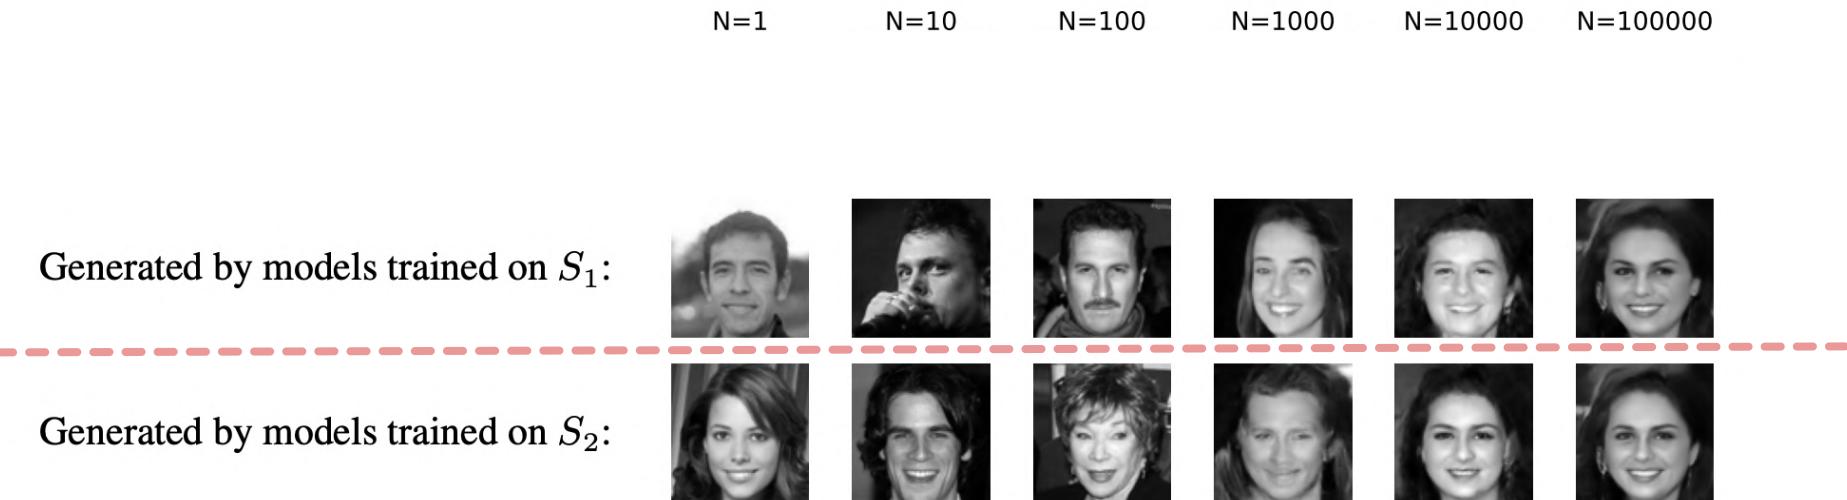
— $\Omega_m = 0.10 \ \sigma_8 = 0.92$ — $\Omega_m = 0.18 \ \sigma_8 = 0.68$ — $\Omega_m = 0.27 \ \sigma_8 = 0.92$
— $\Omega_m = 0.35 \ \sigma_8 = 0.62$ — $\Omega_m = 0.44 \ \sigma_8 = 0.95$



Physics as a testing ground: Well-understood summary statistics enable rigorous validation of generative models

Has my model learned the underlying density?

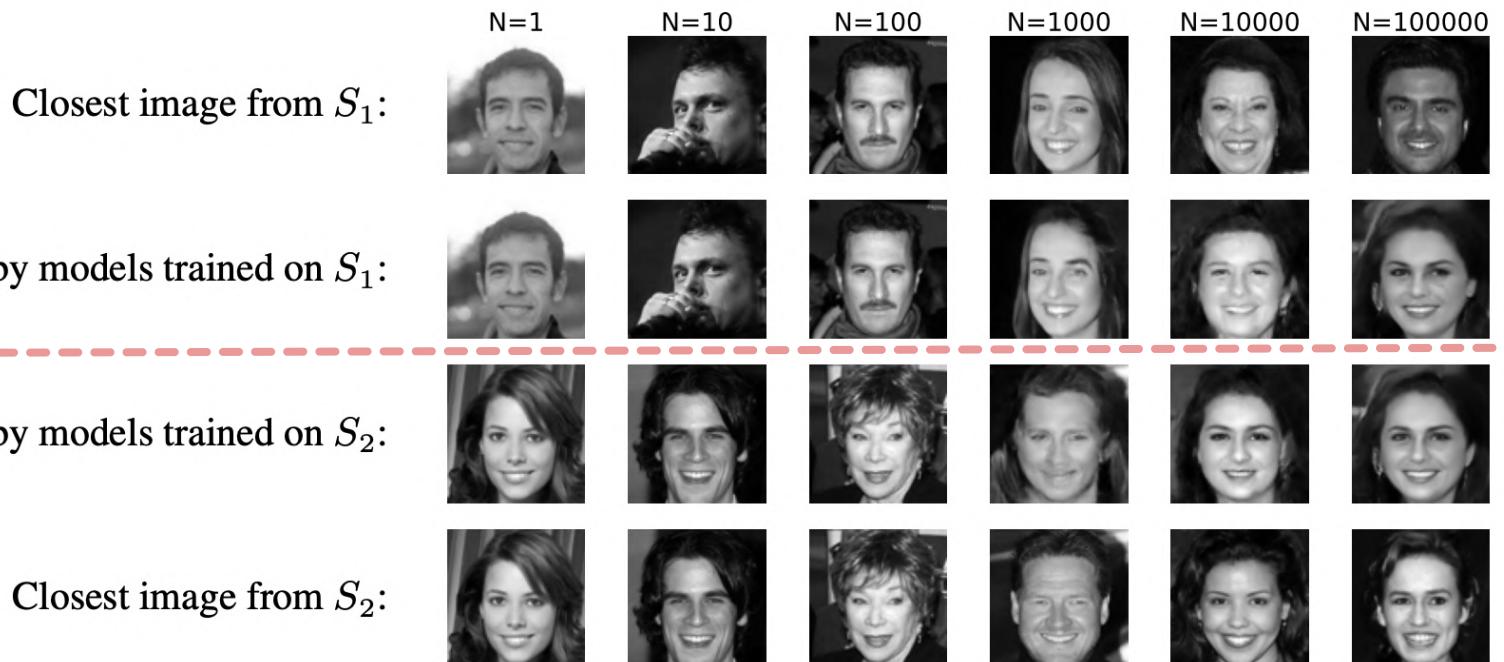
Split training set into non-overlapping S_1, S_2



["Generalization in diffusion models arises from geometry-adaptive harmonic representations" Kadkhodaie et al (2024)]

Has my model learned the underlying density?

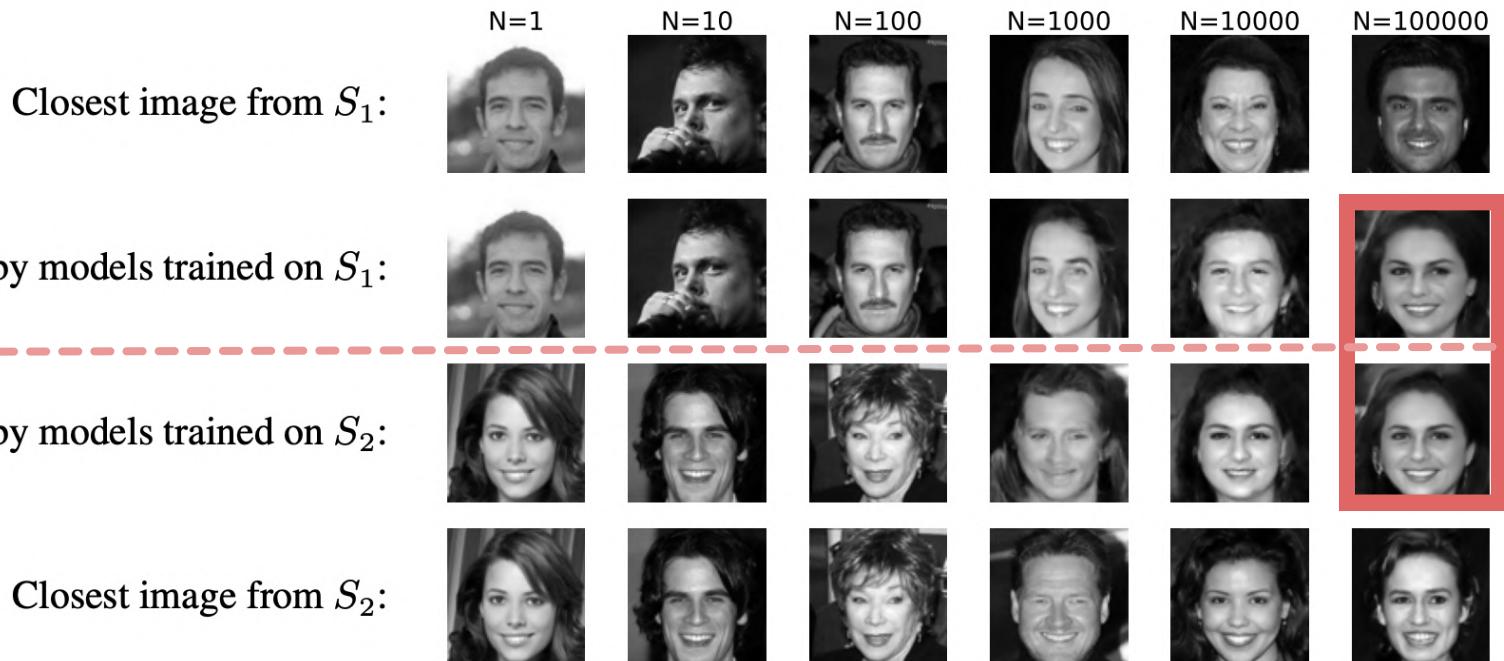
Split training set into non-overlapping S_1, S_2



["Generalization in diffusion models arises from geometry-adaptive harmonic representations" Kadkhodaie et al (2024)]

Has my model learned the underlying density?

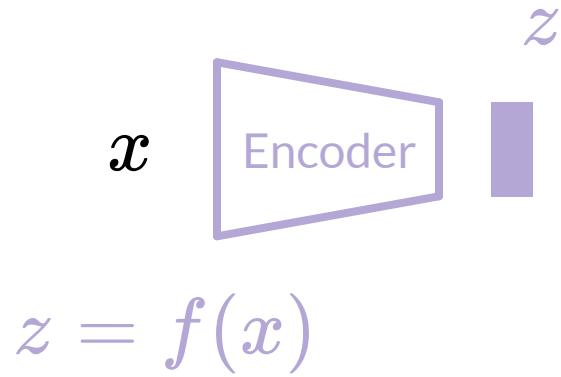
Split training set into non-overlapping S_1, S_2



["Generalization in diffusion models arises from geometry-adaptive harmonic representations" Kadkhodaie et al (2024)]

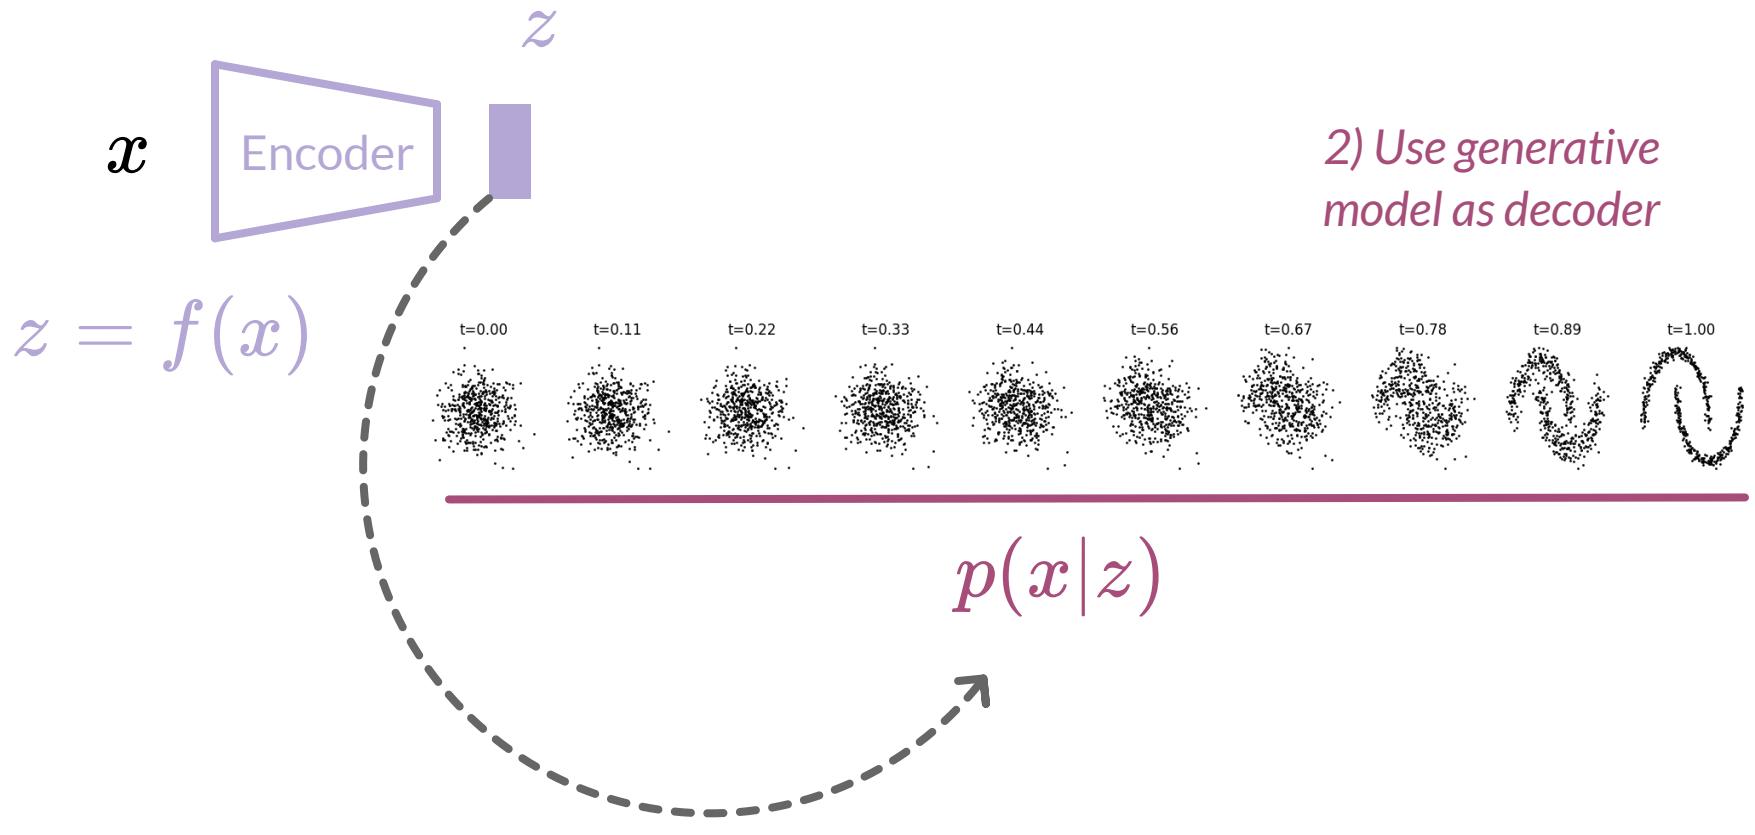
Generate = Understand?

1) Compress into low dimensional latent



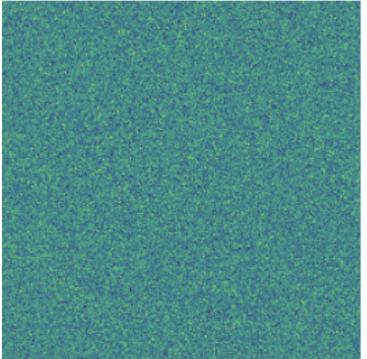
Generate = Understand?

1) Compress into low dimensional latent

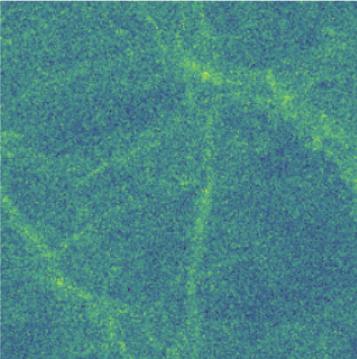


Autoregressive in Frequency

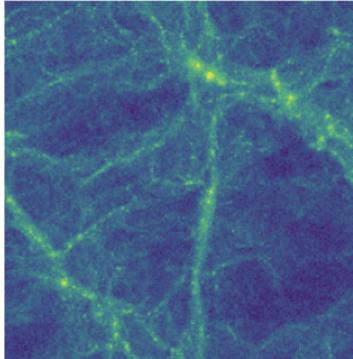
Reconstruction
 $t = 0.03$



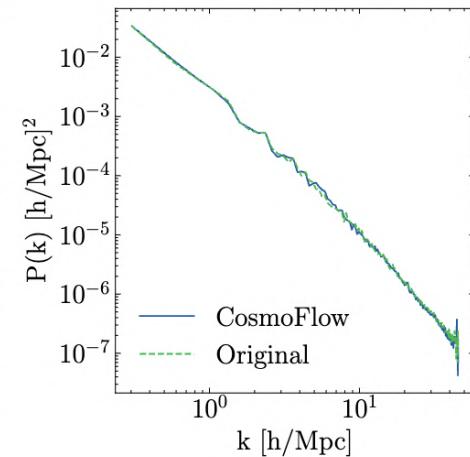
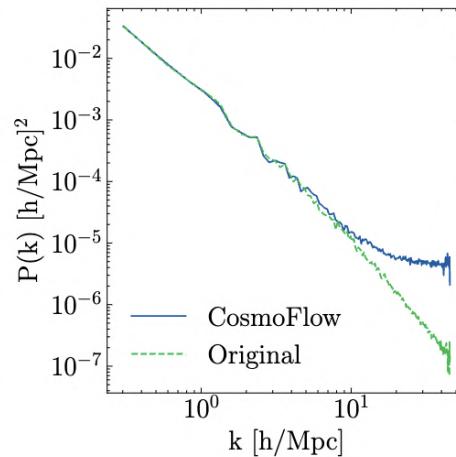
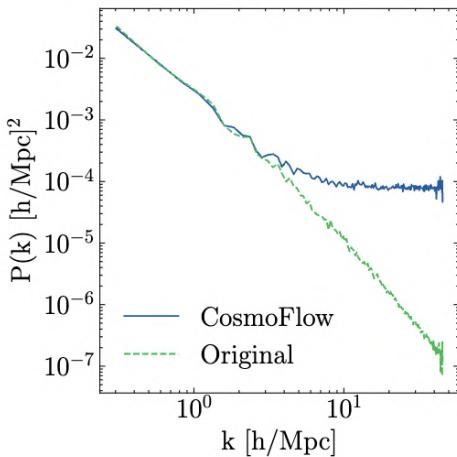
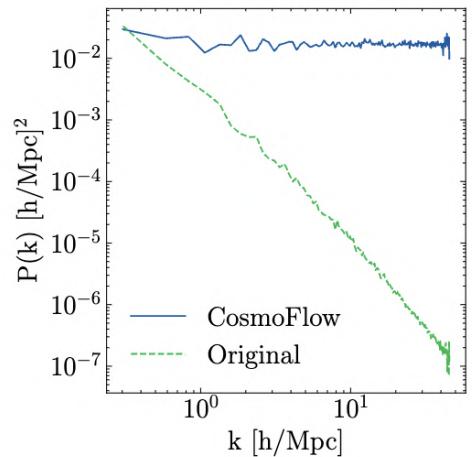
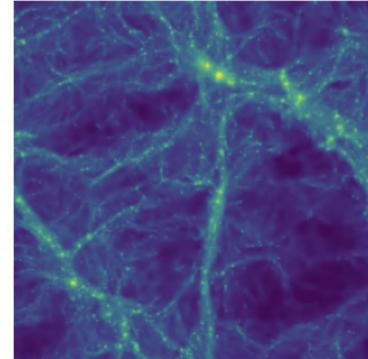
Reconstruction
 $t = 0.34$



Reconstruction
 $t = 0.69$

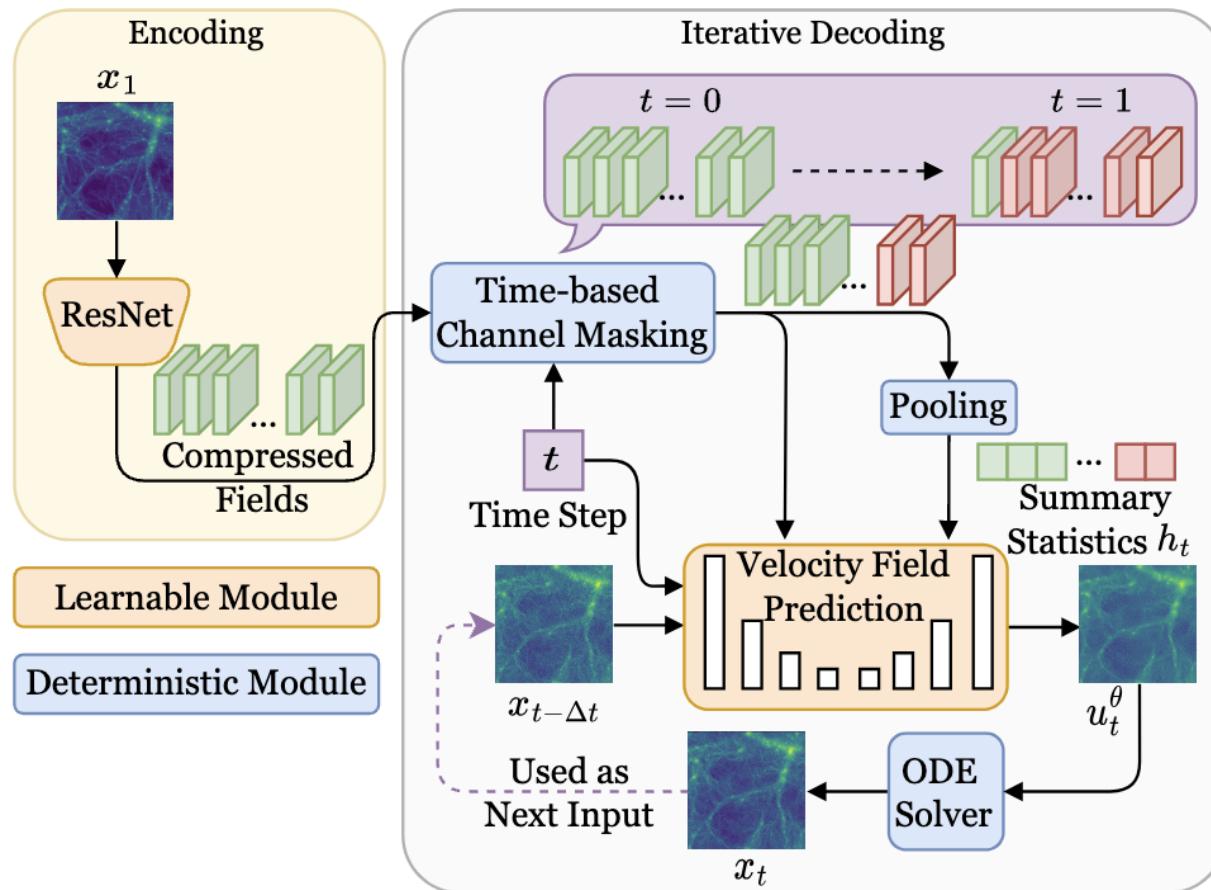


Reconstruction
 $t = 1.00$



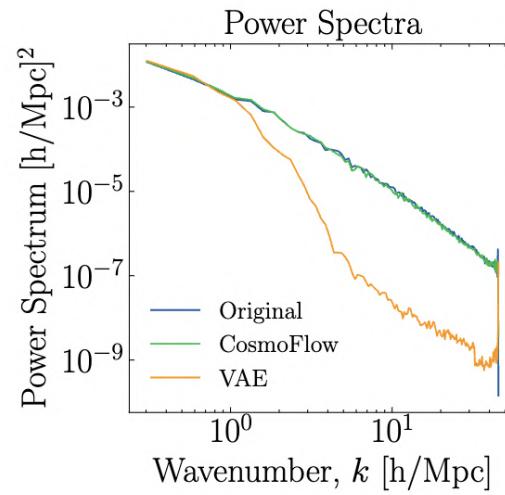
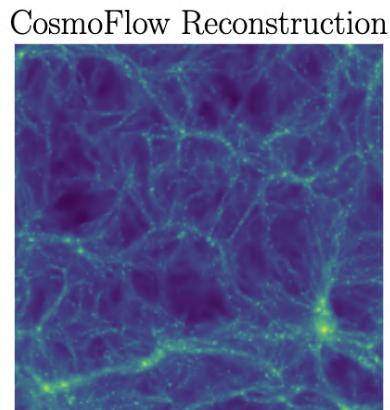
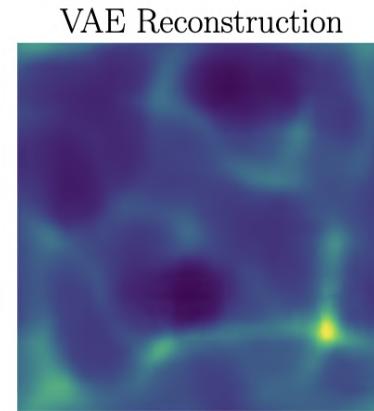
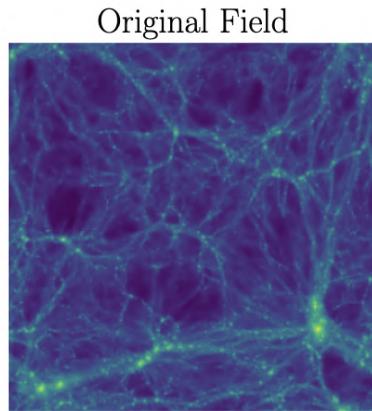
["CosmoFlow: Scale-Aware Representation Learning for Cosmology with Flow Matching" Kannan, Qiu, Cuesta-Lazaro, Jeong (in prep)]

Generate = Understand?



["CosmoFlow: Scale-Aware Representation Learning for Cosmology with Flow Matching" Kannan, Qiu, Cuesta-Lazaro, Jeong (in prep)]

Generate = Understand?

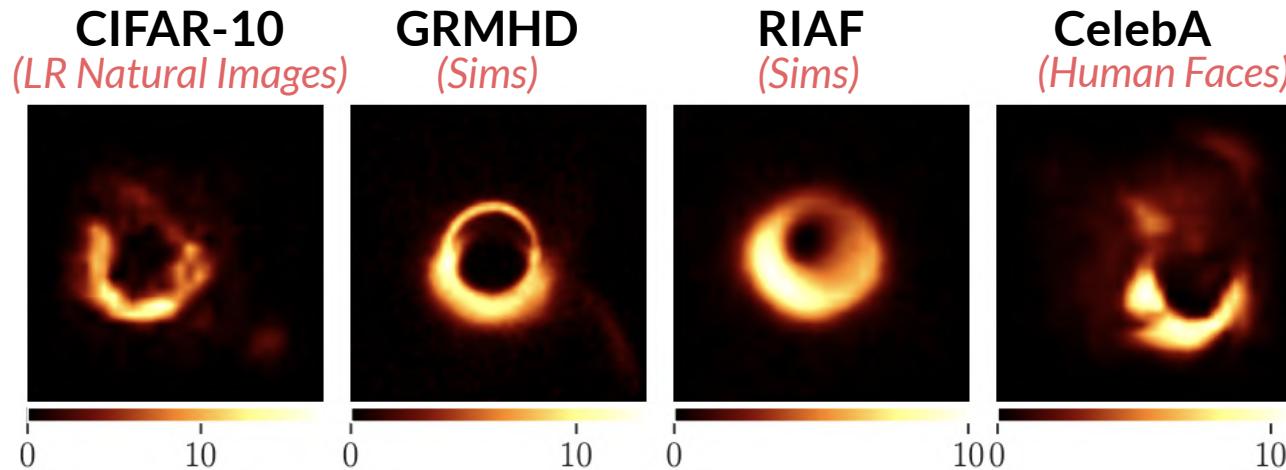


["CosmoFlow: Scale-Aware Representation Learning for Cosmology with Flow Matching" Kannan, Qiu, Cuesta-Lazaro, Jeong (in prep)]

Generative Priors: Plug and Play

$$\log p(x|y) = \log p(y|x) + \underbrace{\log p(x)}_{\text{Prior}} + \text{constant}$$

EHT posterior samples with different priors



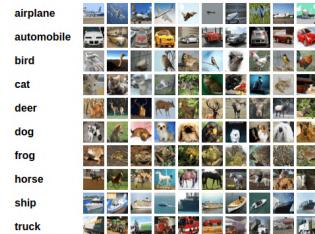
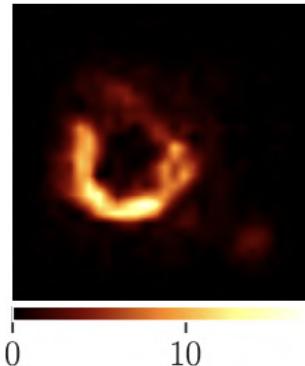
[*"Event-horizon-scale Imaging of M87* under Different Assumptions via Deep Generative Image Priors"* Feng et al]

Generative Priors: Plug and Play

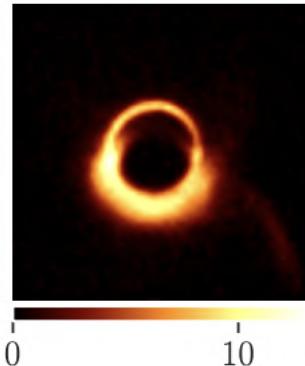
$$\log p(x|y) = \log p(y|x) + \underbrace{\log p(x)}_{\text{Prior}} + \text{constant}$$

EHT posterior samples with different priors

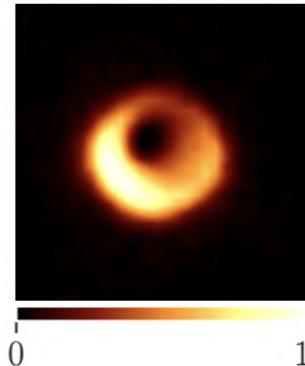
CIFAR-10
(LR Natural Images)



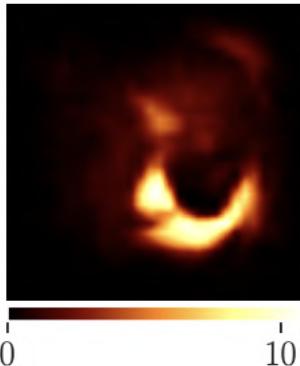
GRMHD
(Sims)



RIAF
(Sims)



CelebA
(Human Faces)



["Event-horizon-scale Imaging of M87* under Different Assumptions via Deep Generative Image Priors" Feng et al]

Generative priors: learn directly?

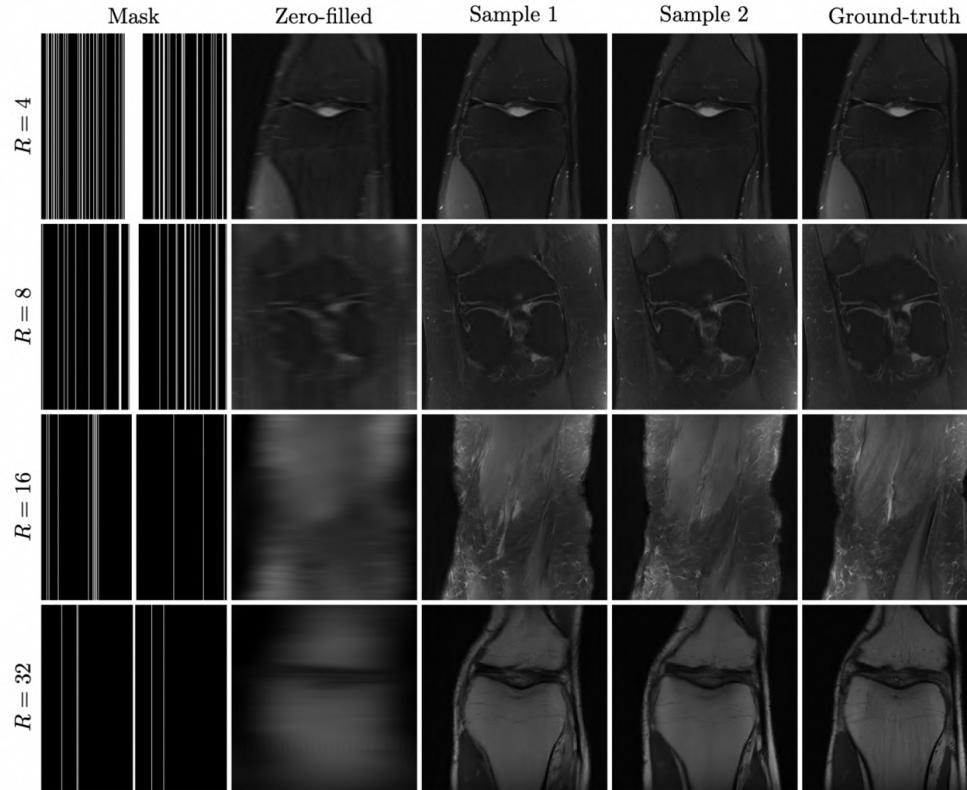


Figure 6. Examples of posterior samples for accelerated MRI using a diffusion prior trained from k -space observations only. Posterior samples are detailed and present plausible variations, while remaining consistent with the observation. We provide the zero-filled inverse, where missing frequencies are set to zero, as baseline.

["Learning Diffusion Priors from Observations by Expectation Maximization"

Rozet et al]

Bridging two distributions

How is the bridge constrained?

Normalizing flows: Reverse = Forward inverse

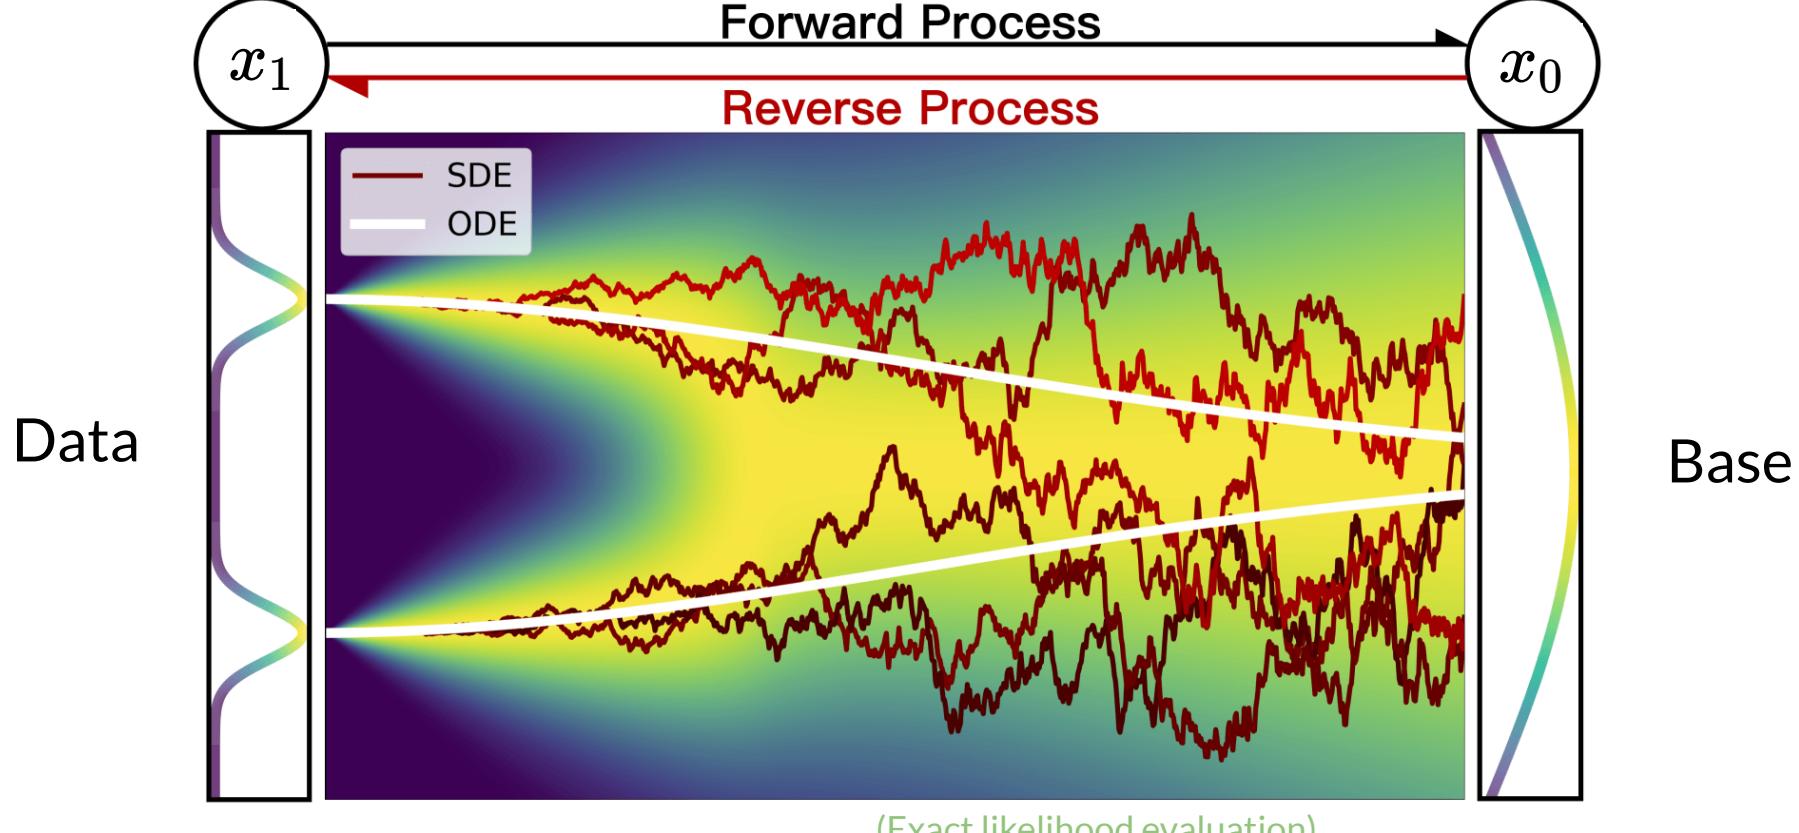
Diffusion: Forward = Gaussian noising

Flow Matching: Forward = Interpolant

is $p(x_0)$ restricted?

Diffusion: $p(x_0)$ is Gaussian

Normalising flows: $p(x_0)$ can
be evaluated

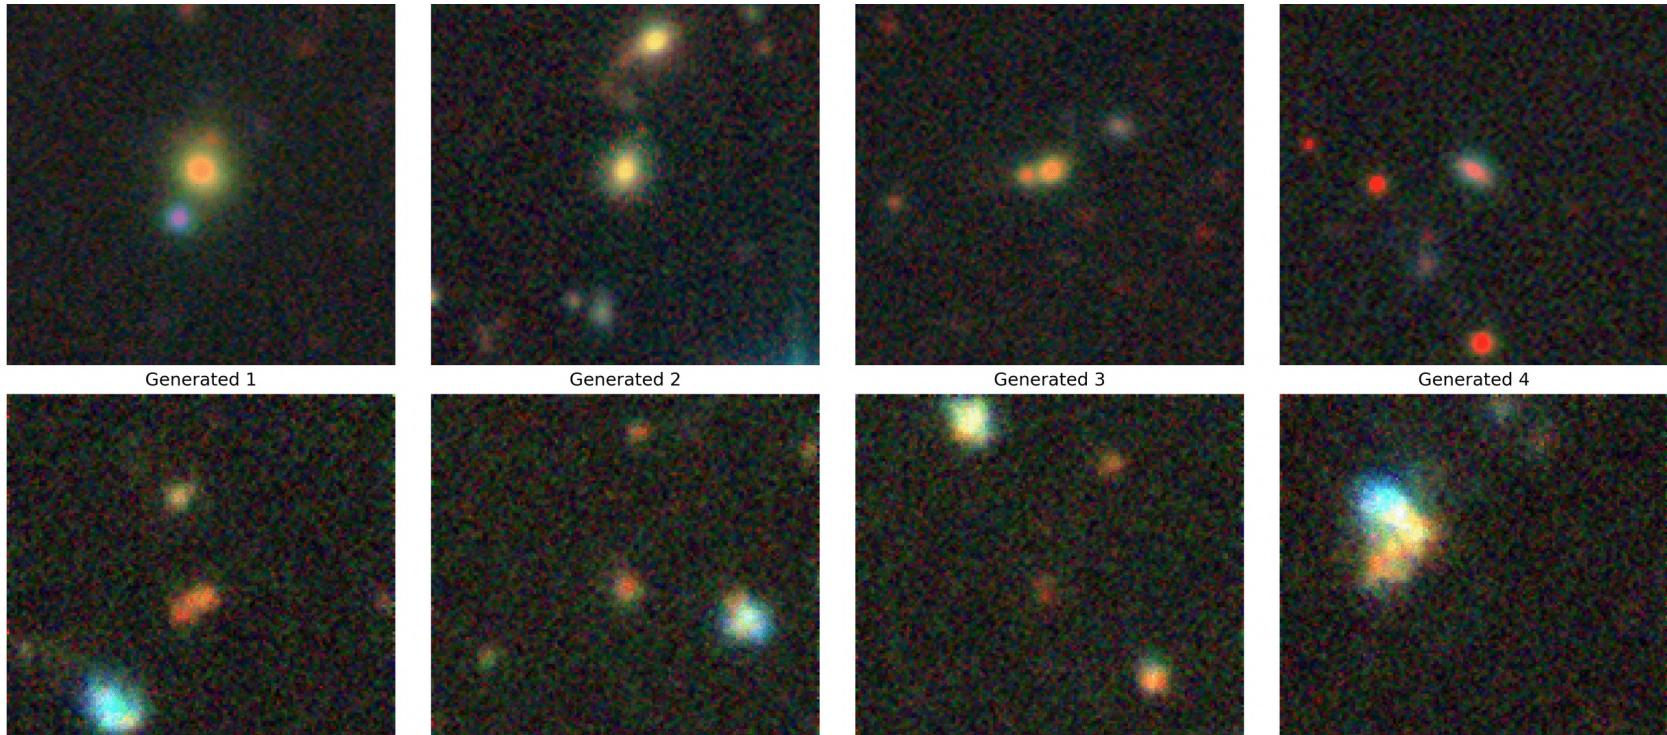


Is bridge stochastic (SDE) or deterministic (ODE)?

Diffusion: Stochastic (SDE)

Normalising flows: Deterministic (ODE)

Tutorial 4: Generating Galaxy Images



References



- Books by Kevin P. Murphy carolina.clzr@gmail.com
 - *Machine learning, a probabilistic perspective*
 - *Probabilistic Machine Learning: advanced topics*
- ML4Astro workshop <https://ml4astro.github.io/icml2023/>
- ProbAI summer school <https://github.com/probabilicai/probai-2023>
- IAIFI Summer school
 - <https://github.com/iaifi/summer-school-2023>
 - Tutorials https://github.com/florpi/summer_school_generative
- Blogposts
 - <https://lilianweng.github.io/posts/2018-10-13-flow-models/>
 - <https://sander.ai/2023/07/20/perspectives.html>
 - <https://mlg.eng.cam.ac.uk/blog/2024/01/20/flow-matching.html>